

CMS Pixel Online Software and Calibrations

Anders Ryd (*anders.ryd@cornell.edu*)
Steve Stroiney (*srs63@lepp.cornell.edu*)
Souvik Das (*sd259@cornell.edu*)
Karl Ecklund (*Karl.Ecklund@cern.ch*)
Josh Thompson (*joshua.thompson@cern.ch*)
Ben Kreis (*benjaminkreis@gmail.com*)

April 16, 2019

Abstract

Pixel Online Software (POS) is the online data acquisition software for the CMS Pixel Detector installed in 2008. The POS orchestrates its operation, calibration and monitoring. This Note describes the structure and installation of this software. It also documents the algorithms used to calibrate the detector. This document currently serves both as a Reference Manual for and a User's Guide to POS.

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 7 |
| 2 | Pixel DAQ System | 7 |
| 2.1 | Overview of DAQ components | 7 |
| 2.1.1 | The Read Out Chip | 8 |
| 2.1.2 | The Token Bit Manager | 9 |
| 2.1.3 | The Pixel FED | 10 |
| 2.1.4 | The Pixel FEC | 11 |
| 2.2 | Installation at P5 | 11 |
| 3 | Pixel Online Software Overview | 14 |
| 4 | Package structure | 15 |
| 4.1 | Pixel Function Manager | 16 |
| 4.1.1 | FSM Implementation Status | 16 |
| 4.1.2 | Control of the L1 Function Manager | 18 |
| 4.1.3 | Outline of L1FM implementation | 18 |
| 4.2 | PixelSupervisor | 19 |
| 4.2.1 | Functions | 19 |
| 4.2.2 | Interface | 19 |
| 4.3 | PixelFECSupervisor | 20 |
| 4.3.1 | Halted State | 20 |
| 4.3.2 | Configured State | 21 |
| 4.3.3 | Running State | 21 |
| 4.3.4 | Paused State | 21 |
| 4.4 | PixelFEDSupervisor | 22 |
| 4.4.1 | Halted State | 23 |
| 4.4.2 | Configured State | 23 |
| 4.4.3 | Running State | 23 |
| 4.4.4 | Paused State | 23 |
| 5 | Coding practices | 25 |
| 5.1 | Makefile | 25 |
| 5.2 | Include files | 25 |
| 5.3 | CVS tags | 25 |
| 5.4 | Building RPMs | 25 |
| 6 | Configuration Data Management | 26 |

| | |
|---|-----------|
| 7 Configuration Database Interface | 26 |
| 7.1 C++ Configuration Data Access API | 26 |
| 7.1.1 Retrieving data from the database | 26 |
| 7.1.2 Storing data in the database | 28 |
| 7.1.3 Configuration keys | 29 |
| 7.1.4 Alias manipulation | 29 |
| 7.1.5 Support of polymorphism | 30 |
| 7.2 Command line interface to configuration data | 30 |
| 7.2.1 Inserting new data | 30 |
| 7.2.2 Retrieving data | 31 |
| 7.2.3 Creating new configurations | 31 |
| 7.2.4 Alias manipulation | 32 |
| 7.2.5 Use cases | 32 |
| 7.3 Managing Configurations | 33 |
| 8 Configuration Objects | 34 |
| 8.1 Introduction | 34 |
| 8.2 Trim and mask bits | 37 |
| 8.3 ROC DACs | 41 |
| 8.4 PixelDetectorConfig | 42 |
| 8.5 PixelROCStatus | 43 |
| 8.6 PixelNameTranslation | 44 |
| 8.7 PixelFECCConfig | 44 |
| 8.8 PixelTKFECCConfig | 44 |
| 8.9 PixelFEDConfig | 45 |
| 8.10 PixelCalibConfiguration | 45 |
| 8.11 PixelFedCard | 49 |
| 8.12 PixelTBMSettings | 49 |
| 8.13 PixelPortcardMap | 49 |
| 8.14 PixelPortCardConfig | 49 |
| 8.15 PixelDelay25Calibration | 51 |
| 8.16 PixelGlobalDelay25 | 51 |
| 9 Usage of configuration data in XDAQ applications | 51 |
| 9.1 Global Delay 25 usage | 54 |
| 10 Configuration | 55 |
| 10.1 Configuration Sequence | 55 |
| 10.1.1 Possible modifications to the configuration Sequence | 56 |
| 10.2 Configuration steps of the underlying supervisors | 57 |
| 10.2.1 PixelDCSFSMInterface and PixelDCSToTrkFECDpInterface | 57 |
| 10.2.2 PixelTKFECSupervisor | 57 |
| 10.3 Quick reconfiguration for the fine delay scan | 57 |

| | | |
|-----------|--|-----------|
| 10.3.1 | Relevant settings | 57 |
| 10.3.2 | Implementation | 58 |
| 10.3.3 | Limitations | 58 |
| 11 | Event builder configurations | 58 |
| 12 | Directory Structure | 59 |
| 12.1 | Tree structure | 59 |
| 13 | Calibration Algorithms | 60 |
| 13.1 | AOH and FED channel mapping test | 60 |
| 13.1.1 | Mapping test steps | 60 |
| 13.1.2 | Parameters | 61 |
| 13.2 | FED phase and delay scan | 62 |
| 13.2.1 | Output | 62 |
| 13.2.2 | Example configuration | 63 |
| 13.3 | Delay25 settings for send data and return data | 63 |
| 13.3.1 | Output | 63 |
| 13.4 | Delay25 trigger setting | 66 |
| 13.5 | FED baseline calibration | 66 |
| 13.6 | AOH bias settings | 66 |
| 13.6.1 | Introduction and discussion | 66 |
| 13.6.2 | AOH bias calibration steps | 68 |
| 13.6.3 | Parameters | 70 |
| 13.7 | AOH gain calibration | 70 |
| 13.7.1 | Introduction and discussion | 70 |
| 13.7.2 | AOH gain calibration steps | 71 |
| 13.7.3 | Parameters | 72 |
| 13.8 | TBM UB calibration | 72 |
| 13.8.1 | Introduction and discussion | 72 |
| 13.8.2 | TBM UB calibration steps | 74 |
| 13.8.3 | Parameters | 75 |
| 13.9 | ROC UB equalization calibration | 75 |
| 13.9.1 | Introduction and discussion | 75 |
| 13.9.2 | ROC UB equalization calibration steps | 75 |
| 13.9.3 | Parameters | 77 |
| 13.10 | Address level determination | 78 |
| 13.10.1 | Data volume and time estimate | 78 |
| 13.11 | Pixel alive, Scurve, and Gain calibration | 78 |
| 13.11.1 | Practical issues | 79 |
| 13.12 | Iana vs. Vana Calibration | 79 |
| 13.13 | Settings of CalDelay and VcThreshold | 81 |
| 13.13.1 | Output | 81 |

| | |
|---|-----------|
| 13.13.2 Example | 81 |
| 13.14 CalDel calibration | 83 |
| 13.14.1 Output | 83 |
| 13.14.2 Example | 83 |
| 13.15 Idigi vs. Vsf | 85 |
| 13.15.1 Output | 85 |
| 13.15.2 Example | 85 |
| 13.16 Linearity vs. Vsf | 87 |
| 13.16.1 Introduction and discussion | 87 |
| 13.16.2 Linearity vs. Vsf calibration steps | 88 |
| 13.16.3 Parameters | 89 |
| 13.17 Vsf and VHldDel | 90 |
| 13.17.1 Introduction and discussion | 90 |
| 13.17.2 Vsf and VHldDel calibration steps | 90 |
| 13.17.3 Parameters | 92 |
| 13.18 Pulse height range calibration | 93 |
| 13.18.1 Introduction and discussion | 93 |
| 13.18.2 Pulse height calibration procedure | 93 |
| 13.18.3 Parameters | 94 |
| 13.19 Vana and time walk | 94 |
| 13.20 Trim bit determination | 95 |
| 13.21 Trim bit analysis | 96 |
| 13.21.1 Running the threshold analysis | 96 |
| 13.21.2 Quality of trimming | 97 |
| 13.21.3 Determining new VcThr settings | 97 |
| 13.21.4 Make simple plots of thresholds | 98 |
| 13.21.5 Full trim bit determination | 98 |
| 13.21.6 Determining all trimbites | 99 |
| 13.22 Emulated physics | 99 |
| 14 Calibration Procedure | 99 |
| 14.1 Useful Tools, Background Knowledge | 99 |
| 14.2 Starting Point | 100 |
| 14.3 AOH and FED Channel Mapping | 100 |
| 14.4 Achieving Successful Programming | 101 |
| 14.5 Basic Signal Properties | 101 |
| 14.5.1 FEDBaseline | 101 |
| 14.5.2 AOH Settings | 103 |
| 14.5.3 Ultrablacks | 104 |
| 14.5.4 Clock Phase | 105 |
| 14.6 Getting Hits | 106 |
| 14.7 Address Levels | 107 |

| | |
|--|------------|
| 14.8 Initial Trimming | 108 |
| 14.9 Vana Calibration | 108 |
| 14.10 Final Trimming | 109 |
| 14.11 Threshold Minimization | 110 |
| 14.12 Pulse Height Optimizations | 113 |
| 14.12.1 Vsf | 113 |
| 14.12.2 VHldDel | 114 |
| 14.12.3 PHRange | 114 |
| 14.12.4 Gain Calibration | 115 |
| 15 XDAQ-DCS Interaction | 116 |
| 15.1 PSX Server | 117 |
| 15.2 PixelDCSFSMInterface | 117 |
| 15.3 Use of DCS information by the supervisors | 118 |
| 15.4 Configuration | 118 |
| 15.5 HV-dependent DAC programming | 122 |
| 15.6 Start and Resume transitions | 123 |
| 15.7 Pause, Stop, and Halt transitions | 123 |
| 15.8 Expected behavior | 123 |
| 15.9 Notes on hardcoding of information | 124 |
| 16 Procedures to follow when hardware components are replaced | 124 |
| 16.1 FED board | 124 |
| 16.2 Pixel FEC including mFECs | 125 |
| 16.3 TKFEC | 125 |
| 17 Low-level commands | 125 |
| 17.1 Resets | 125 |
| A FED phase and delay calibration | 127 |
| B Calibration Organization | 132 |
| B.1 PixelSupervisor calibration code | 132 |
| B.2 PixelFEDSupervisor calibration code | 135 |
| B.3 Implementation of algorithms | 136 |

1 Introduction

This note describes the design and implementation of Pixel Online Software (POS). POS is used for controlling and calibrating the CMS pixel detector that was installed in 2008. In particular, it performs the following functions:

- Configuring the detector
- Performing online calibrations
- Analyzing calibration data in online farm (CMSSW)
- Monitoring the detector during data taking

Other information about Pixel Online Software is available at the pixel online software wiki pages [1].

2 Pixel DAQ System

The intention of this section is to give an overview of the CMS pixel DAQ system. For more details the reader can follow the references given in the text. But the introduction given here should be sufficient to understand the main goals of POS.

2.1 Overview of DAQ components

The CMS pixel DAQ system consists of a number of components as illustrated in Fig. 1. This figure is specific to the FPix detector, but the modification for the BPix are minor and not really relevant to the discussion here unless otherwise mentioned. Starting from the detector itself we have the Read Out Chip (ROC) [2] and the Token Bit Manager (TBM) [3]. Some properties of the ROC are described below as needed to understand the calibrations. The ROC reads out 4,160 pixels and the TBM coordinates the communication with a group of 8 to 24 ROCs. The TBM is electrically connected to the portcard via extension cables ≈ 2 feet in length.

The portcard receives and sends optical signals to the Front-End Controller (FEC) and Front-End Driver (FED) respectively. The FEC is used to program the TBM and ROCs and the data read out from the detector is sent to the FED to be digitized. The portcard is home to several discrete components. We have the Digital Opto-Hybrid (DOH) that receives data from the FEC and the Analog Opto-Hybrid (AOH) that transmits data to the FED. On the portcard we have in addition the (t)PLL, Delay25, and gatekeeper chips.

Of particular interest is the Delay25 chip. The communication to the portcard via the DOH is done at 40MHz on a serial line. We have a clock line and the data line (and return clock and return data). In order to have this communication working the timing between the clock and data lines must be right. The purpose of the Delay25 chip is to adjust delays to make this communication work. The setting of the parameters on the Delay25 chip as

well as the other components on the portcard is done using the I2C protocol from the CCU. The CCU is again controlled from a Tracker FEC (TkFEC).

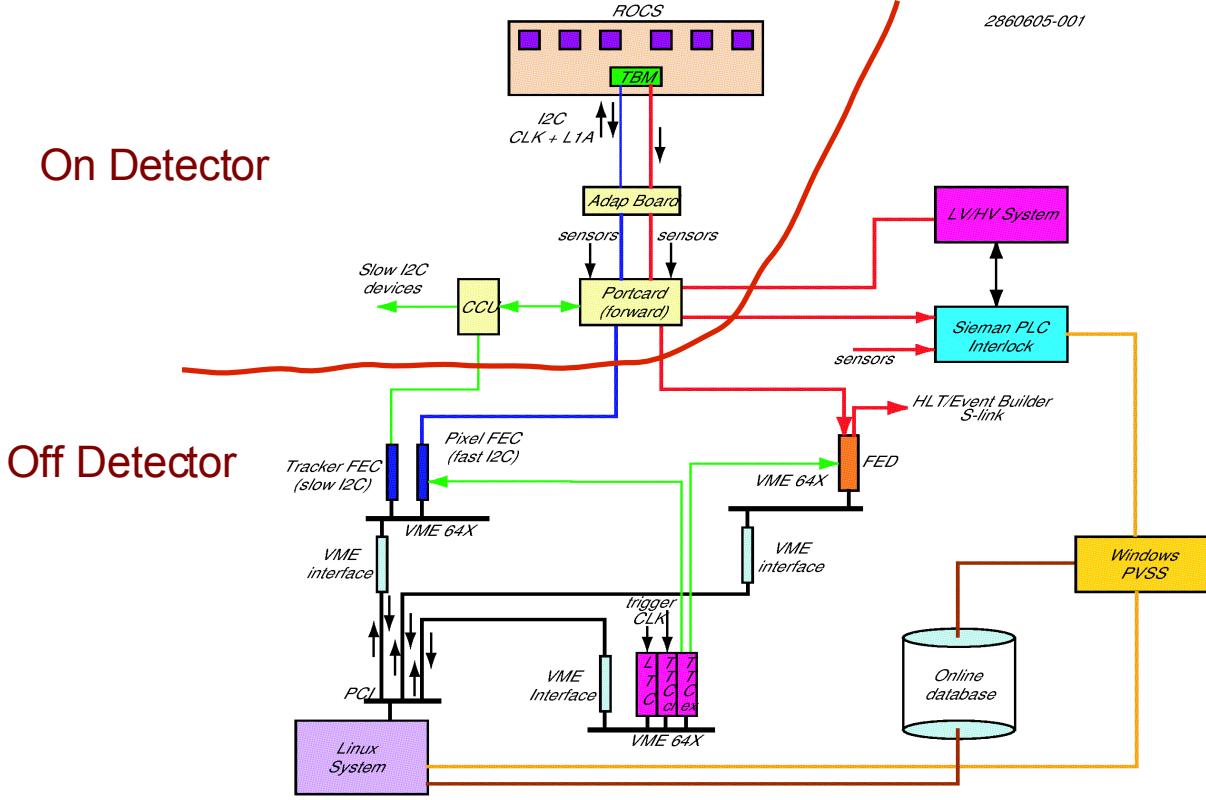


Figure 1: The main components in the CMS pixel DAQ system.

2.1.1 The Read Out Chip

Each Read Out Chip (ROC) [2], illustrated in Fig. 2, collates data from $52 \times 80 = 4,160$ pixels and contains about 1.3 million transistors. It amplifies and zero suppresses data using 4 trim bits that specify a threshold for every pixel. The chip buffers hit data until the trigger decision arrives. It is developed at PSI and manufactured by IBM [4].

For pixel hits to be read out, the following consecutive actions must take place.

- The accumulated charge in any of the pixels of the double column must exceed a (programmable) threshold (trim-input to comparator). Then the corresponding time-stamp (bunch crossing number) is written into the time-stamp buffer presently pointed at, and the analog signals and the pixel addresses of all hit pixels are written into the next free data buffers. This hit-recording into the time stamp and data buffers runs autonomously and asynchronously in each double column of the chip, independently

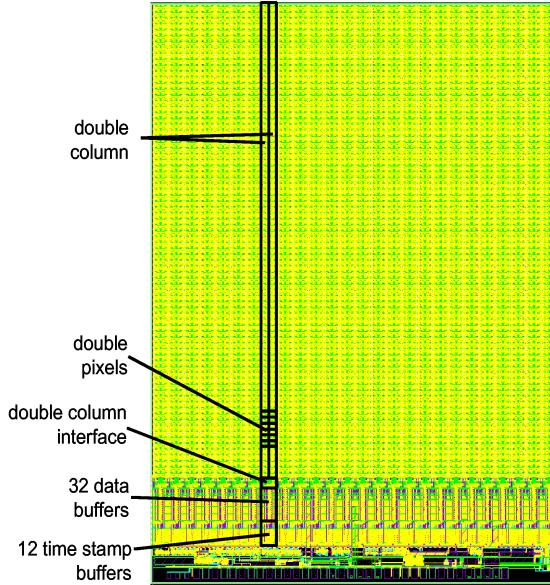


Figure 2: The Read Out Chip of the CMS Pixel Detector.

of the bunch crossing clock. The recorded hit information must be kept in the buffers during the latency time of the L1-trigger.

- Hits in the column are validated by an external L1-trigger, by comparing their corresponding time stamp with a counter running behind the bunch crossing counter by the trigger delay, else the hits are cleared. If the hits in the column are validated, the value of a 4-bit trigger counter is latched into the column periphery. The column is frozen and cannot record further until reset after the readout of the triggered hits. Other (untriggered) columns remain active.
- All frozen columns with the latched value of the trigger counter equal to the present value of the token counter are set to the readout mode. Directly afterwards, the trigger counter is incremented. The readout process starts when the token bit enters the chip. After a three bit chip header is sent, each consecutive double column which is in the readout mode is read out and reset. Just before the token leaves the chip, the token counter is incremented. A new token is sent for each trigger and only data belonging to that trigger will be sent onto the readout bus, even if more triggers arrive between a trigger and the corresponding token.

2.1.2 The Token Bit Manager

A Token Bit Manager (TBM) controls groups of 8 to 24 ROCs. It distributes the clock and trigger signals. It serializes analog readout using a token bit passed from ROC to ROC. It is physically mounted next to the ROCs. The TBM is developed at Rutgers.

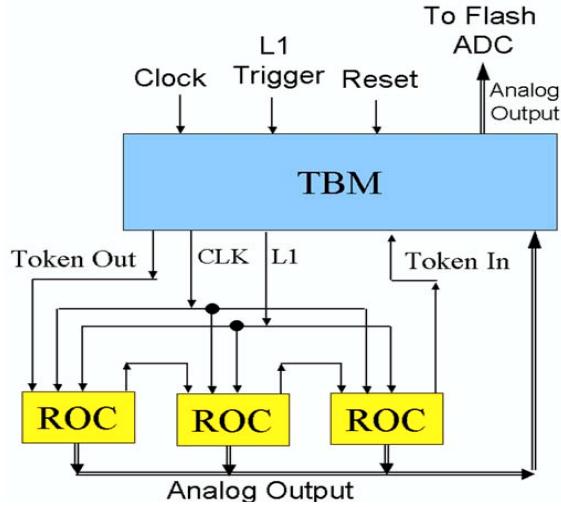


Figure 3: A schematic of how the Token Bit Manager reads out data from the ROCs.

The principal functions of the TBM are the following.

- It controls the readout of the ROCs by initiating a token pass for each incoming L1-trigger.
- On each token pass, it writes a header and a trailer word to the data stream.
- The header contains an 8-bit event number. The trailer contains an 8-bit error status. These are transmitted in 2-bit analog encoded digital signals.
- It distributes the L1-trigger and LHC Clock to the ROCs.
- Each arriving L1-trigger is placed on a 32-deep stack awaiting its associated token pass. Normally the stack is empty but is needed to accommodate high burst rates due to either noise, high track density or trigger bursts. Only the first 16 triggers placed on the stack are passed to the ROCs, and have an associated token pass. All others are marked in the trailer as a “No Token Pass” event. This is done to keep the DAQ synchronized, while speeding up the clearing of the stack.

The TBM receives 4 control signals: A 40 MHz clock synchronized to the CMS beam crossing, an active low reset signal (which resets all control registers to an initial state), a serial line for configuration settings, and a serial encoded trigger signal.

2.1.3 The Pixel FED

The Pixel Front End Driver (FED) is a VME module with 36 optical inputs. It sits in the counting room and receives the analog output of the ROCs via an optical fiber from the

center of the CMS detector. In order for the FED to decode and digitize this data, it has to have address levels well calibrated, and this is one of the several tasks performed by the Pixel Online Software. Internal timings of the FED also have to be calibrated within errors of a few nanoseconds, and this is also accomplished using the POS. The output of the FED is dispatched, in a binary format readable to the entire CMS experiment, via an S-Link.

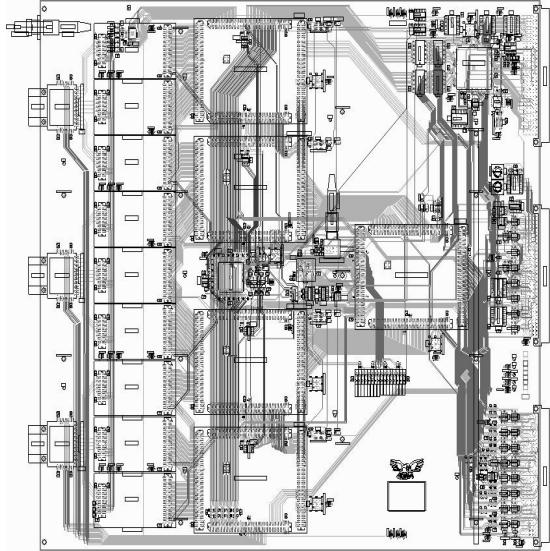


Figure 4: A schematic of the pixel FED.

2.1.4 The Pixel FEC

The Pixel Front End Controller (FEC) sends triggers, clocks and programming data to the ROCs. For the pixel detector, we use a one-way ‘fast’ I2C protocol to send this data. We need to send 1 byte of data per pixel, or 66 MB of data for the entire detector. For this, we use a 40 MHz serial line, and this calls for timing calibrations in order to enable downloads.

2.2 Installation at P5

Table 1 lists the online PCs that we have at P5 for Pixel Online Software. Table 2 lists the FED boards and their locations in the VME crate within the CMS Counting Room. The slots and the VME addresses are also provided. Similarly, we list the FEC boards and their locations in Table 3. The CCU connections in the Counting Room are shown in Table 4.

Table 1: Pixel online PCs at P5

| Slot in rack | Node name | Front | Label | Size | Function | Comments | OS |
|--------------|-----------------|-------|-------|------|----------------|--------------|------------|
| - | vmepcs2b18-17 | Pixel | - | ?U | Histoviewer | | SLC4 |
| 19 | vmepcs2b18-16 | Pixel | 1/10 | 1U | VME/S1G01i | BPix FECs | SLC4 |
| 17 | vmepcs2b18-15 | Pixel | 10/10 | 1U | VME/S1G01e | FPix FECs | SLC4 |
| 16 | vmepcs2b18-14 | Pixel | 9/10 | 1U | VME/S1G04e | BPix FEDs do | SLC4 |
| 14/15 | vmepcs2b18-13 | Pixel | 8/10 | 2U | VME/S1G04i | BPix FEDs up | SLC4 |
| 13 | vmepcs2b18-12 | Pixel | 7/10 | 1U | VME/S1G03i | FPix FEDs | SLC4 |
| 12 | vmepcs2b18-11 | Pixel | 6/10 | 1U | PixelSuperFPix | FPix Supv. | SLC4 |
| 11 | vmepcs2b18-10 | Pixel | 5/10 | 1U | PixelSuperBPix | BPix Supv. | SLC4 |
| 10 | vmepcs2b18-09 | Pixel | 4/10 | 1U | DCS | Fpix daq | Win |
| 9 | vmepcs2b18-08 | Pixel | 3/10 | 1U | DCS/Siemens | | Win |
| 8 | vmepcs2b18-07 | Pixel | 2/10 | 1U | DCS/CAEN | | Win |
| | vmepcs2b16-10 | | | | TTC+LTC | Supv. | SLC4 |
| | cmsrc-pixel | | | | L1 FM | | Pixel FM |
| | fmmppc-s1d12-08 | | | | FMM PC | | FMM PC |
| | cmspsx | | | | PSX server | | PSX server |
| | srv-c2c02-05 | | | | DB server | | SLC4 |
| | srv-c2c02-06 | | | | DB server | | SLC4 |

Table 2: FED connections for FPIX

| Rack | Crate | Slot | VME addr. | Channels | FED id | Name(Official) | Name(Construction) |
|-------|-------|------|------------|----------|--------|--------------------------|--------------------|
| S1G03 | upper | 6 | 0x13000000 | 13-24 | 33 | BpO_D(1,2)_BLD(1,2,3) | HC+Z1 1.1 2.1 |
| S1G03 | upper | 6 | 0x13000000 | 1-12 | 33 | BpO_D(1,2)_BLD(4,5,6) | HC+Z1 1.2 2.2 |
| S1G03 | upper | 7 | 0x14000000 | 13-24 | 34 | BpO_D(1,2)_BLD(7,8,9) | HC+Z1 1.3 2.3 |
| S1G03 | upper | 7 | 0x14000000 | 1-12 | 34 | BpO_D(1,2)_BLD(10,11,12) | HC+Z1 1.4 2.4 |
| S1G03 | upper | 5 | 0x12000000 | 1-12 | 32 | BpI_D(1,2)_BLD(1,2,3) | HC+Z2 1.4 2.4 |
| S1G03 | upper | 5 | 0x12000000 | 13-24 | 32 | BpI_D(1,2)_BLD(4,5,6) | HC+Z2 1.3 2.3 |
| S1G03 | upper | 8 | 0x15000000 | 1-12 | 35 | BpI_D(1,2)_BLD(7,8,9) | HC+Z2 1.2 2.2 |
| S1G03 | upper | 8 | 0x15000000 | 13-24 | 35 | BpI_D(1,2)_BLD(10,11,12) | HC+Z2 1.1 2.1 |
| S1G03 | upper | 11 | 0x17000000 | 13-24 | 37 | BmI.D(1,2)_BLD(1,2,3) | HC-Z1 1.1 2.1 |
| S1G03 | upper | 11 | 0x17000000 | 1-12 | 37 | BmI.D(1,2)_BLD(4,5,6) | HC-Z1 1.2 2.2 |
| S1G03 | upper | 12 | 0x18000000 | 13-24 | 38 | BmI.D(1,2)_BLD(7,8,9) | HC-Z1 1.3 2.3 |
| S1G03 | upper | 12 | 0x18000000 | 1-12 | 38 | BmI.D(1,2)_BLD(10,11,12) | HC-Z1 1.4 2.4 |
| S1G03 | upper | 10 | 0x16000000 | 1-12 | 36 | BmO.D(1,2)_BLD(1,2,3) | HC-Z2 1.4 2.4 |
| S1G03 | upper | 10 | 0x16000000 | 13-24 | 36 | BmO.D(1,2)_BLD(4,5,6) | HC-Z2 1.3 2.3 |
| S1G03 | upper | 13 | 0x19000000 | 1-12 | 39 | BmO.D(1,2)_BLD(7,8,9) | HC-Z2 1.2 2.2 |
| S1G03 | upper | 13 | 0x19000000 | 13-24 | 39 | BmO.D(1,2)_BLD(10,11,12) | HC-Z2 1.1 2.1 |

Table 3: FEC connections for FPIX

| Rack | Crate | Slot | VME addr. | mFEC | Name(Official) | Name(Construction) |
|-------|--------|------|------------|------|--------------------------|--------------------|
| S1G01 | middle | 5 | 0x28000000 | 3 | BpO_D(1,2)_BLD(1,2,3) | HC+Z1 1.1 2.1 |
| S1G01 | middle | 5 | 0x28000000 | 4 | BpO_D(1,2)_BLD(4,5,6) | HC+Z1 1.2 2.2 |
| S1G01 | middle | 5 | 0x28000000 | 5 | BpO_D(1,2)_BLD(7,8,9) | HC+Z1 1.3 2.3 |
| S1G01 | middle | 5 | 0x28000000 | 6 | BpO_D(1,2)_BLD(10,11,12) | HC+Z1 1.4 2.4 |
| S1G01 | middle | 5 | 0x28000000 | 8 | BpI_D(1,2)_BLD(1,2,3) | HC+Z2 1.4 2.4 |
| S1G01 | middle | 5 | 0x28000000 | 7 | BpI_D(1,2)_BLD(4,5,6) | HC+Z2 1.3 2.3 |
| S1G01 | middle | 5 | 0x28000000 | 2 | BpI_D(1,2)_BLD(7,8,9) | HC+Z2 1.2 2.2 |
| S1G01 | middle | 5 | 0x28000000 | 1 | BpI_D(1,2)_BLD(10,11,12) | HC+Z2 1.1 2.1 |
| S1G01 | middle | 10 | 0x50000000 | 3 | BmI_D(1,2)_BLD(1,2,3) | HC-Z1 1.1 2.1 |
| S1G01 | middle | 10 | 0x50000000 | 4 | BmI_D(1,2)_BLD(4,5,6) | HC-Z1 1.2 2.2 |
| S1G01 | middle | 10 | 0x50000000 | 5 | BmI_D(1,2)_BLD(7,8,9) | HC-Z1 1.3 2.3 |
| S1G01 | middle | 10 | 0x50000000 | 6 | BmI_D(1,2)_BLD(10,11,12) | HC-Z1 1.4 2.4 |
| S1G01 | middle | 10 | 0x50000000 | 8 | BmO_D(1,2)_BLD(1,2,3) | HC-Z2 1.4 2.4 |
| S1G01 | middle | 10 | 0x50000000 | 7 | BmO_D(1,2)_BLD(4,5,6) | HC-Z2 1.3 2.3 |
| S1G01 | middle | 10 | 0x50000000 | 2 | BmO_D(1,2)_BLD(7,8,9) | HC-Z2 1.2 2.2 |
| S1G01 | middle | 10 | 0x50000000 | 1 | BmO_D(1,2)_BLD(10,11,12) | HC-Z2 1.1 2.1 |

Table 4: CCU connections for FPIX

| Rack | Crate | Slot | mFEC | Name(Official) | Name(Construction) |
|-------|--------|------|------|----------------|--------------------|
| S1G01 | middle | 18 | 8 | BpI | HC+Z2 |
| S1G01 | middle | 18 | 7 | BpO | HC+Z1 |
| S1G01 | middle | 18 | 6 | Bm0 | HC-Z2 |
| S1G01 | middle | 18 | 5 | BmI | HC-Z1 |

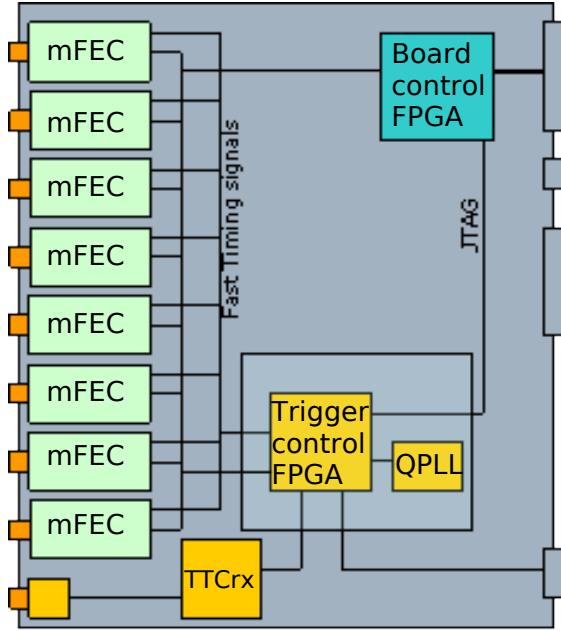


Figure 5: A schematic of the pixel FEC.

3 Pixel Online Software Overview

The Pixel Online Software is based on the XDAQ toolkit and is a suite of applications. The different XDAQ based components are shown in green in Fig. 6. The top level application is the PixelSupervisor. This application is responsible for the overall coordination of the pixel DAQ. The PixelSupervisor talks to the supervisors that directly control the hardware. For example we have the PixelFECSupervisor that provides the interface to the pixel FECs. Similarly the PixelFEDSupervisor controls FEDs. In production at P5, there are multiple instances of the PixelFECSupervisor and PixelFEDSupervisor; one per VME crate¹.

The PixelTKFECSupervisor controls the tracker FEC hardware. The pixel system uses the tracker FEC hardware slow I2C to initialize the fast I2C used for the download of most configuration data.

The PixelTTCSupervisor controls the pixel TTC module used for trigger and timing. Among other things the TTC module is used during calibrations to generate triggers. In modern releases of the software, the PixelTTCSupervisor has been deprecated in favor of the TTCCiControl, which is a standard package maintained by the TTC group.

The PixelLTCSupervisor is used for the local trigger control.

The various supervisors run as independent processes, or even on different computers. Therefore, in order to communicate with each other they must exchange messages on the network. This is done using the SOAP protocol.

¹The strip tracker uses a design where there is one supervisor per VME board.

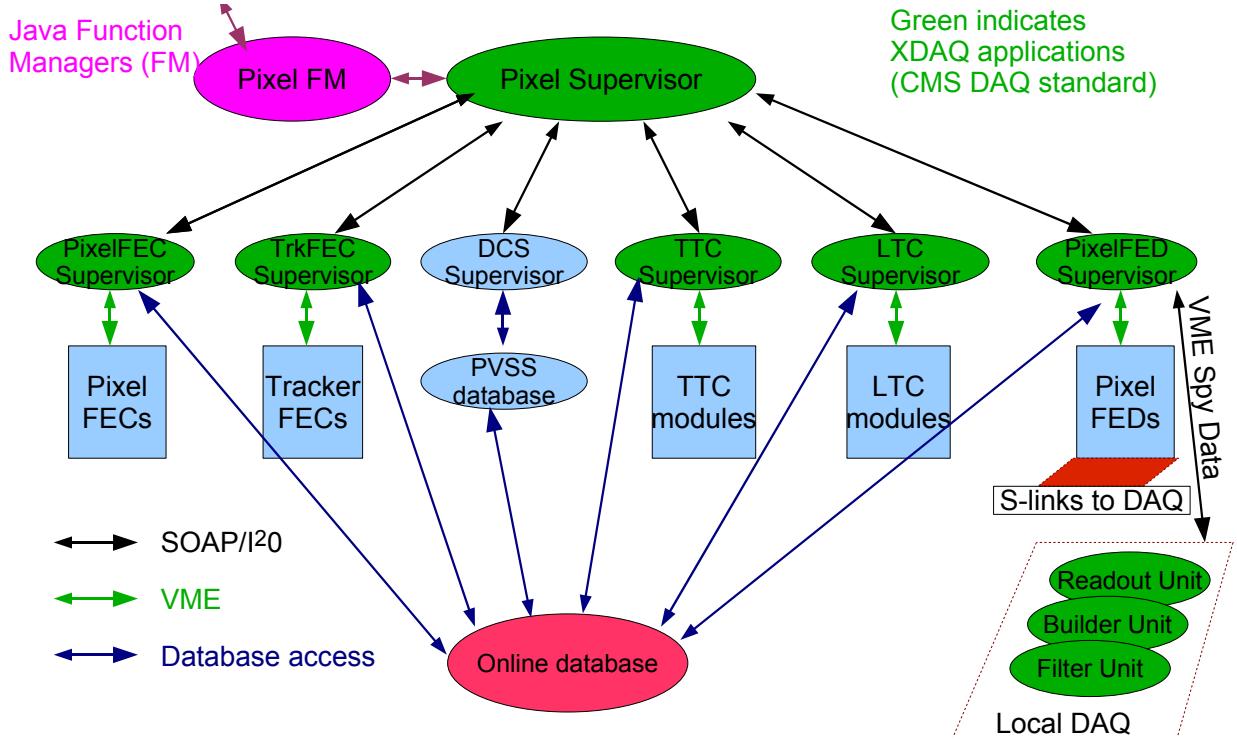


Figure 6: The different applications that compose the Pixel Online Software.

The Level 1 function manager (L1FM or FM) for the pixels is the interface the pixel system has to the global run control (RCMS for Run Control and Monitoring System). The FM is a java application. It responds to requests from RCMS to change states in the state diagram that describe the state of the DAQ system. This state diagram is shown in Fig. 8. The pixel FM is a relatively thin layer that basically just passes the state changes on to the PixelSupervisor.

4 Package structure

In the Pixel Online Software the code is distributed among a number of packages. These packages are listed here.

- PixelCalibrations
- CalibFormats/SiPixelObjects
- PixelConfigDBInterface
- PixelDCSInterface

- PixelFECInterface
- PixelFECSupervisor
- PixelFEDInterface
- PixelFEDSupervisor
- PixelFunctionManager
- PixelLTCSupervisor
- PixelSupervisor
- PixelTKFECSupervisor
- PixelTTCSupervisor²
- PixelUtilities

The package dependency tree is shown in Fig 7. The supervisor applications are at the top and depend on the packages below. We should make sure that the dependencies form a tree and not contain loops.

4.1 Pixel Function Manager

The Pixel Function Manager (the Level 1 Function Manager) acts as an interface between Run Control (the Level 0 Function Manager) and POS. The Pixel Function Manager is a Java application. It implements the state machine of CMS [5]. The Function Manager interacts with the PixelSupervisor to carry out the different tasks needed in state transitions of the run control.

4.1.1 FSM Implementation Status

The CMS Run Control finite state machine (FSM) is shown in Fig. 8. This model is implemented in the Pixel Function Manager, and also in the PixelSupervisor. The other Supervisors implement portions of this model as required. The FSM depicted in the figure is completely implemented in the PixelSupervisor. The PixelSupervisor does include transitions from any state into the "Error" state, and then allows a "Recover" transition that returns the FSM to the "Halted" state.

²In use through tag `POS_3_1_2`; deprecated starting in `POS_3_2_0`.

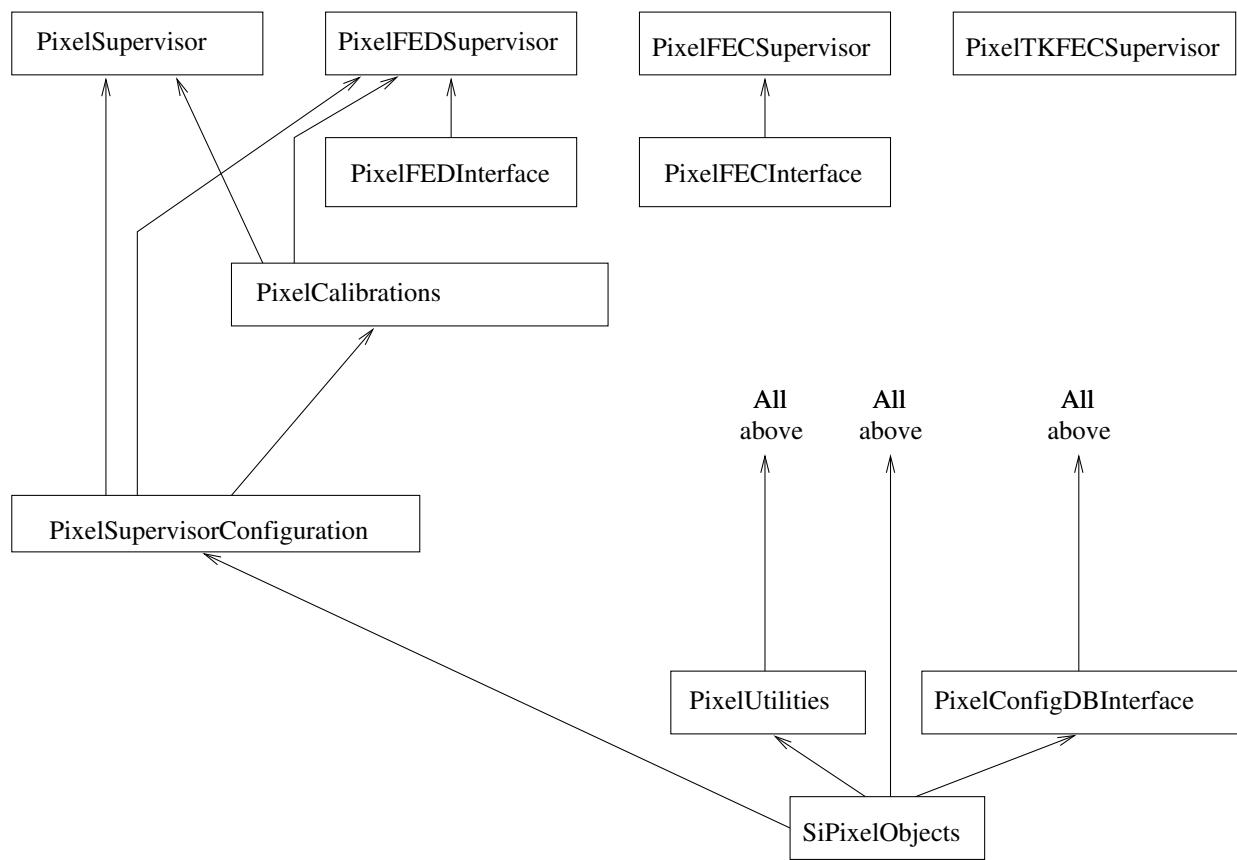


Figure 7: The dependencies among the packages are indicated here. At the top are Supervisor applications.

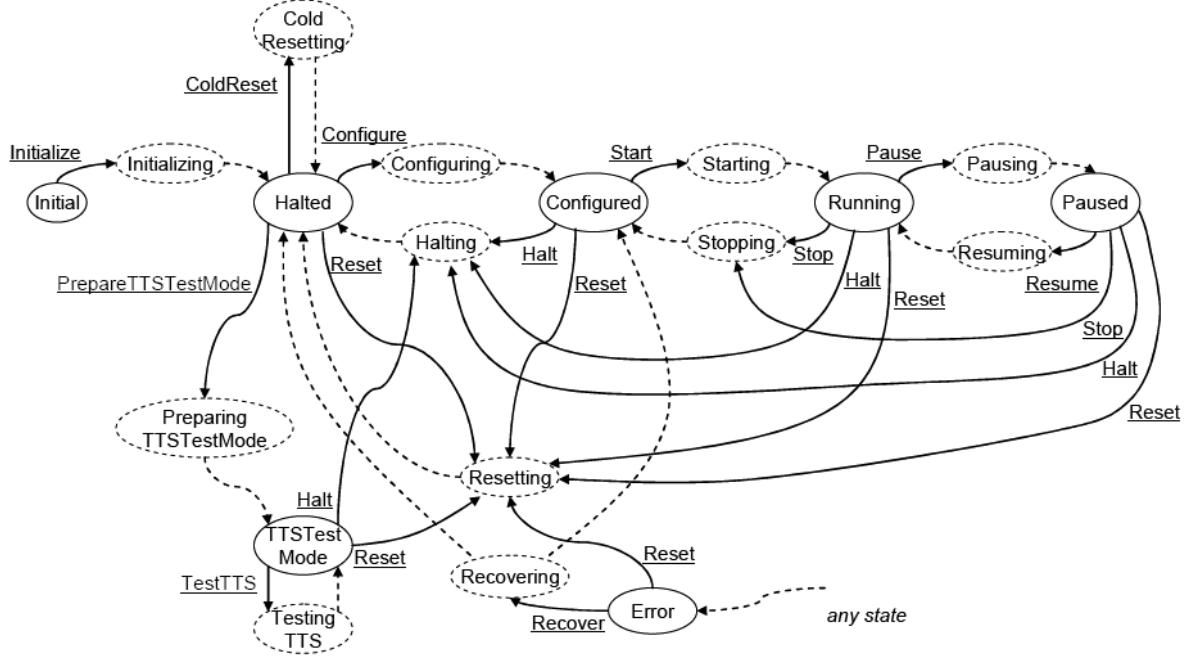


Figure 8: The CMS finite state machine definition. Figure taken from Ref. [5].

4.1.2 Control of the L1 Function Manager

During global running, the L1 Function Manager (L1FM) is driven by the L0FM, which is operated by central DAQ. The pixel user should not intervene via the L1FM GUI, except to check its status.

During local running, the L1FM must be created via its GUI. State transitions of the FSM can then be driven via the L1FM GUI. In general, we drive the “Initialize” state transition via the GUI, then drive subsequent state transitions directly from the PixelSupervisor. However, in principle all transitions can be initiated from the L1FM GUI³.

4.1.3 Outline of L1FM implementation

The L1FM is created using the Create button in RCMS. This calls the method of `PixelFunctionManager.java` called `createAction()`.

After creation, transitions of the FM FSM are handled by methods in `PixelEventHandler.java`. Which method is triggered depends on the transition, as defined by a table near the top of this class. The names are fairly logical (for example, “Initialize” corresponds to the `initAction` method and “Configure” corresponds to the `configureAction` method). Inside each `fooAction` method are two blocks: one to handle objects of type

³In practice, this is only useful for simple tests of the L1FM to PixelSupervisor communication.

`StateEnteredEvent` and one to handle objects of type `StateNotification`.

When an FSM transition “foo” is initiated by the L0FM or L1FM GUI, we enter the `StateEnteredEvent` block of `fooAction`. This block then contains the code to send a message to PixelSupervisor, telling it to proceed with transition “foo”. The L1FM then does nothing while PixelSupervisor coordinates the necessary activities in POS. When the transition is completed by the PixelSupervisor, it passes a message back to the L1FM⁴. This message triggers entry into the `StateNotification` block of `fooAction`. If PixelSupervisor reports that the transition was successful, this block moves the L1FM FSM into the state “foo”. If PixelSupervisor reports an error, this block moves the L1FM FSM into state “Error”. In this way, the L0FM (central DAQ) learns whether the transition was successful.

4.2 PixelSupervisor

The PixelSupervisor is the top level Pixel Online Software application in the Pixel Online Software. As described above it takes commands from the function manager. There is one pixel supervisor for the pixel online system.

4.2.1 Functions

The main function of the PixelSupervisor is to coordinate the activities of the other supervisors, particularly during configuration (see Sec. 10) and calibration (see Sec. 13). It is responsible for updating the configuration database with new settings obtained by calibrations.

The PixelSupervisor also communicates the state of the pixel XDAQ software (POS) to the Level 1 Function Manager (Sec. 4.1).

4.2.2 Interface

The PixelSupervisor web GUI is an html page, which by default refreshes every few seconds. It displays information about the current configuration, or if it is not configured it allows the user to select a possible configuration from a list and configure the detector using that configuration.

The PixelSupervisor runs the JobControl Monitor, which is a utility that periodically sends “heartbeat” SOAP messages to the JobControl processes running on the various machines at P5. The PixelSupervisor GUI uses the replies from these SOAP messages to display whether any of the POS XDAQ processes has crashed, or whether any of the JobControl processes themselves are unresponsive. (Note that we typically only run JobControl at P5, so this feature is not available elsewhere.)

⁴This is implemented in PixelSupervisor using the `stateChanged` method of the `rcmsStateNotifier_-` object.

The figure consists of four separate screenshots of a web-based control interface for the Pixel Front End Controller Supervisor. Each screenshot shows a header with the XDAQ logo, the title 'Pixel Front End Controller Supervisor', the version number (3.0), and the date ('Wed, 31 Jan 2007').

- Halted State:** The 'Current State' is 'Halted'. A 'Global Key' input field is present. Buttons for 'Configure', 'Halt', 'Pause', 'Resume', and 'Start' are available.
- Configured State:** The 'Current State' is 'Configured'. A 'Global Key' input field is present. Buttons for 'Configure', 'Halt', 'Pause', 'Resume', and 'Start' are available.
- Running State:** The 'Current State' is 'Running'. A 'Global Key' input field is present. Buttons for 'Configure', 'Halt', 'Pause', 'Resume', and 'Start' are available.
- Paused State:** The 'Current State' is 'Paused'. A 'Global Key' input field is present. Buttons for 'Configure', 'Halt', 'Pause', 'Resume', and 'Start' are available.

Below each state section is a 'Finite State Machine' and a 'Low Level Commands' section. The 'Low Level Commands' section contains a link to 'FEC with Base Address 0x30000000'.

Figure 9: The operational states of the PixelFECSupervisor: Halted, Configured, Running and Paused.

4.3 PixelFECSupervisor

The PixelFECSupervisor supervises one crate of pixel-FEC boards. PixelFECSupervisor may be run by PixelSupervisor via SOAP messages, or independently through its own web interface called PixelFECSupervisorGUI. The PixelFECSupervisor is used to load configuration parameters for the ROCs from the configuration database and program them into the ROCs. Its web interface offers two levels of interactivity: A finite state machine based GUI that issues commands to all pixel-FECs in its crate as the user steps through the states of the FSM, and a low level GUI that gives the user control down to individual pixels.

The PixelFECSupervisor has a state machine identical to that of the PixelSupervisor and follows that of the PixelSupervisor. If its state machine is triggered using its web interface, care must be taken to return it to the state of the PixelSupervisor before PixelSupervisor is allowed to regain control of it.

4.3.1 Halted State

This is the initial state of the PixelFECSupervisor. As shown in Fig. 9, the user is allowed to enter a Global Key and press the “Configure” button. The Global Key is then used to retrieve configuration data and download them to the FEC hardware while transitioning

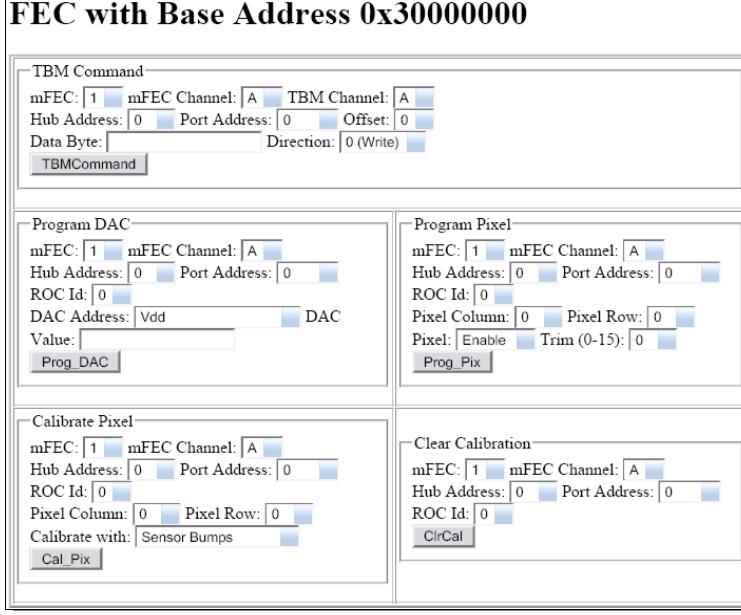


Figure 10: Low Level GUI for the FEC board at VME base address 0x30000000 in the crate controlled by the PixelFECSupervisor

the state machine to the Configured state. The Low Level GUI is not accessible yet as the PixelFECSupervisor has not retrieved the hardware addresses of the FECs it controls.

4.3.2 Configured State

“Halt” and “Start” are the two state machine input keys that are enabled in this state, as shown in Fig. 9. The Low Level GUI displays a list of FECs with their VME base addresses. Clicking on any opens a new browser window with the Low Level GUI corresponding to that FEC board.

4.3.3 Running State

On pressing the “Start” button, the PixelFECSupervisor transitions to its Running state. It is shown in Fig. 9. “Pause” and “Halt” buttons are active. The Low Level GUI is accessible.

4.3.4 Paused State

On pressing the “Pause” button in the Running state, the PixelFECSupervisor transitions to its Paused state. It is shown in Fig. 9. “Resume” and “Halt” state input buttons are active. “Resume” returns the state machine to its Running state, and hence executes the FEC’s part of the last calibration it was last configured for. “Halt” takes the state machine back to its “Halted State”.

The figure consists of four separate browser windows, each representing a different operational state of the PixelFEDSupervisor. Each window has a header with the XDAQ logo, the title 'Pixel Front End Driver Supervisor', a version number (3.0), and a date (Wed, 31 Jan 2007). The footer of each window shows a URL like 'FED with Base Address 0x1c000000'.

- Halted:** The 'Current State' is 'Halted'. Below it is a table with three rows: 'Global Key', 'FED Control Register', and 'FED Mode Register'. At the bottom are buttons for 'Configure', 'Halt', 'Pause', 'Resume', and 'Start'.
- Configured:** The 'Current State' is 'Configured'. Below it is a table with two rows: 'Configure' and a group of five buttons: 'Halt', 'Pause', 'Resume', 'Start', and 'Stop'.
- Running:** The 'Current State' is 'Running'. Below it is a table with two rows: 'Configure' and a group of five buttons: 'Halt', 'Pause', 'Resume', 'Start', and 'Stop'.
- Paused:** The 'Current State' is 'Paused'. Below it is a table with two rows: 'Configure' and a group of five buttons: 'Halt', 'Pause', 'Resume', 'Start', and 'Stop'.

Figure 11: The operational states of the PixelFEDSupervisor: Halted, Configured, Running and Paused.

Clicking any hyperlink for the FEC boards under the title “Low Level Commands” opens a new browser window for the FEC’s Low Level GUI. An example of such a GUI is shown in Fig. 10. It allows the user to send TBM commands to individual TBMs, program-DAC and clear-calibration commands to each Readout Chip, and program-pixel and calibrate-pixel commands to any pixel of any Readout Chip within its control.

4.4 PixelFEDSupervisor

A PixelFEDSupervisor supervises one crate of Front End Driver (FED) boards. PixelFEDSupervisor may be run by PixelSupervisor via SOAP messages, or independently through its own web interface. Its web interface offers two levels of interactivity: A state machine GUI that issues commands to all FEDs in its crate while the user steps through its state machine states, and a Low Level GUI that allows the user to adjust various settings on any FED board in the crate.

The PixelFEDSupervisor has a state machine identical to that of the PixelSupervisor and follows that of the PixelSupervisor. If its state machine is triggered using its web interface, care must be taken to return it to the state of PixelSupervisor before PixelSupervisor is allowed to regain control of it.

4.4.1 Halted State

This is the initial state of the PixelFEDSupervisor. As shown in Fig. 11, the user is allowed to enter a Global Key, a value for the FED Control Register, and a value for the FED Mode Register and press the “Configure” button. The Global Key is then used to retrieve configuration data and download them to the FED hardware along with the FED Control and Mode Register settings while transitioning the state machine to the Configured state. The Low Level GUI is not accessible yet as the PixelFEDSupervisor has not retrieved the hardware addresses of the FEDs it controls.

4.4.2 Configured State

“Halt” and “Start” are the two state machine input keys that are enabled in this state, as shown in Fig. 11. The Low Level GUI displays a list of FEDs with their VME base addresses. Clicking on any opens a new browser window with the Low Level GUI corresponding to that FED board.

4.4.3 Running State

On pressing the “Start” button, the PixelFEDSupervisor transitions to its Running state. It is shown in Fig. 11. “Pause” and “Halt” buttons are active. The Low Level GUI is accessible.

4.4.4 Paused State

On pressing the “Pause” button in the Running state, the PixelFEDSupervisor transitions to its Paused state. It is shown in Fig. 11. “Resume” and “Halt” state input buttons are active. “Resume” returns the state machine to its Running state, and hence executes the FED’s part of the last calibration it was last configured for. “Halt” takes the state machine back to its “Halted State”.

Clicking any hyperlink for the FED boards under the title “Low Level Commands” opens a new browser window for the FED’s Low Level GUI. An example of such a GUI is shown in Fig. 10. A brief description of the buttons follows.

The “Reload Firmware” button reloads the firmware on all four FPGAs of the FED and resets the FED. The “Reset” button resets the FED. The user can modify the Capacitor Adjustment, Optical Receiver Input and Output Offsets, and the Channel Offset DACs for each FED channel. One can also read data from Spy FIFO 1, 2 or 3. If FIFO 1 is chosen, a channel must be specified along with a mode – Transparent or Normal. One can also choose to ship the data to screen, in which case it will be decoded, or shipped to a file, in which case a binary dump of the Spy FIFO data will be placed in the file. It may also be shipped to the RU Builder. However, since we cannot ensure the stability of the RU Builder in the current version of POS, we do not recommend this option.

The user can modify the phase and delay settings for any channel on the FED using the “SetPhasesAndDelays” button.

FED with Base Address 0x1c000000

Reload Firmware

Reset FEDs

Channel Offset

| Channel | Capacitor Adjust | Input Offset | Output Offset | Offset DAC |
|---------|------------------|--------------|---------------|------------|
| 1 | 0 | 0 | 0 | |

Control and Mode Registers
Control Registers
Transparent Mode Disable Enable
Transparent Gate Start by L1A VME or EFT (OPTO Module)
Use simulated test-DAC Disable Enable
Event number generated by TTC VME
L1A triggers from TTCrx Disable Enable
EFT Signals from the OPTO Module Disable Enable
TTSReady Disable Enable
TTSError Disable Enable
OUTofSYN Disable Enable

Mode Registers
S-Link Disable Enable
Write Spy Memory Disable Enable
S-Link Let it be, or Reset

Reading FIFO
Read

- Spy FIFO 1
- Spy FIFO 2 Normal Mode
- Spy FIFO 3 Normal Mode

Ship Spy FIFO data to

- Screen
- RUBuilder
- File

Enable FIFO 3

Clock Phases and Delays
Channel 1 Phase 0 Delay 0

Figure 12: Low Level GUI for the FED board at VME base address 0x1c000000 in the crate controlled by the PixelFEDSupervisor

5 Coding practices

5.1 Makefile

The Pixel Online Software is built with a Makefile located in each (sub)package. The Makefile builds on the XDAQ tools. This file should ideally be as short as possible and use the functionality in the XDAQ package. You invoke the Makefile in each package by doing a ‘make’. There is also a `clean` target.

In addition to the package level makefiles there is a top level makefile in the pixel directory. This makefile allows you to build all the XDAQ packages using `make Set=pixel`. You can also clean all packages using `make Set=pixel clean`.

5.2 Include files

Include statements should include the file path starting from the project. For example you should do

```
#include "PixelCalibrations/include/PixelAOHBiasCalibration.h"  
#include "CalibFormats/SiPixelObjects/interface/PixelCalibConfiguration.h"
```

5.3 CVS tags

We create ‘official’ tags of the form ‘POS_X_Y_Z’, e.g. ‘POS_2_4_5’. For every tag created, starting with POS_2_5_0, an entry should be made in the file pixel/README that describes briefly the new features in the tag.

5.4 Building RPMs

The building of RPMs should be straightforward. The following steps are required.

- Prepare the code, update the README file, and VERSION file and commit and tag the code. The VERSION file contains the version of the RPM to build.
- Invoke ‘make Set=pixel rpm’ to build RPMs for all Pixel Online Software packages.
- In the `PixelUtilities` directory, invoke ‘buildExternalRPMs.sh’ to build RPMs for DiagSystem, TTCSoftware, and FecSoftwareV3_0. At the end of building the external RPMs this script copies all RPMs to `$BUILD_HOME/RPM_X.Y.Z-V`, where X, Y, Z, and V are the major version, minor version, patch, and build version, respectively.

The set of RPMs built can be tested (in the online environment) for consistency by invoking the command ‘`rpm -test -Uvh *.rpm`’ in the directory where all the RPMs are located.

Old RPMs can be cleaned out of the build area using the command ‘`make Set=pixel cleanrpm`’.

6 Configuration Data Management

The next two sections describe a C++ interface for configuration data management and the different classes that are used to store this information. For details of the organizational details of the database structure, please refer to the presentation:

<http://indico.cern.ch/conferenceDisplay.py?confId=8768>

7 Configuration Database Interface

Since early summer 2006, we have used an interface for accessing configuration data in the online software framework. The access was originally file based, and has now (starting in 2009) been supplemented with infrastructure to access an Oracle database. The interface used is rather generic and the purpose of this document is to write down the interface so that we can have a clear separation between the database code and the applications that use data from the configuration database.

First in Sect. 7.1 we describe the C++ API. The interface is defined in the class `PixelConfigInterface`. This interface is fully implemented in the `PixelConfigFile` implementation based on files.

In Sect. 7.2 a simple command line tool is described. This tool is implemented through the C++ interface and should work both for the file based and the final data base implementation.

Examples below use a few of the configuration data classes that are used.

7.1 C++ Configuration Data Access API

The class `PixelConfigInterface` in the package `PixelConfigDBInterface` defines the interface for access of configuration data. This interface is intended to provide type safe access methods for retrieving and storing configuration data.

7.1.1 Retrieving data from the database

The primary method for retrieving the data is

```
template <class T>
static void get(T* &pixelObject,
std::string path,
pos::PixelConfigKey key)
```

The `PixelConfigKey` is just an integer that holds the top level configuration key to be used. This interface returns a pointer to the data. The caller is assumed to take ownership of the data and delete it. The path is a 'secondary key'. It would allow us to store more than one object of the same type in a given configuration. (In the file based implementation

this label is used to build the path to where the file is stored.) Below are a few examples of using this interface.

```
PixelConfigKey theGlobalKey(5);

PixelNameTranslation *theNameTranslation=0;
PixelConfigInterface::get(theNameTranslation, "nametranslation/",
                         theGlobalKey);

PixelDetectorConfig *theDetectorConfiguration=0;
PixelConfigInterface::get(theDetectorConfiguration, "detconfig/",
                         theGlobalKey);

PixelFECCConfig *theFECCConfiguration=0;
PixelConfigInterface::get(theFECCConfiguration, "fecconfig/",
                         theGlobalKey);
```

These examples show how you extract objects for which there is only one instance of for the whole detector configuration. You pass in as the first argument a pointer. The pointer will return 0 if the object was not successfully retrieved⁵. The second argument is the label for the object. This is essentially a key that is used to look up the data. The interface is type safe, i.e., if you specify a path to an object of the wrong type you will get back a null pointer. The third argument is the global configuration key. This basically specifies the versions of all objects used in a given configuration. This key is implemented as an integer.

Besides objects like the name translation and detector configuration listed above, there are objects such as trim bits, mask bits, and DAC values that we need to access on a finer granularity than for the whole detector. To do this we use slightly modified arguments

```
PixelDACSettings *tempDACs=0;
PixelConfigInterface::get(tempDACs, 'pixel/dac/FPix_BpI_D1_BLD1_PNL1',
                         theGlobalKey);
```

where we have added to the path the module name for which we want to extract the DAC settings.

As a given application, for example the PixelFECSupervisor, will need to access DAC settings for many modules, and it is more efficient to extract the data 'in bulk' from the database, we have also added an interface that allows extraction of multiple objects at the time

⁵We have now tried to improve on this error-handling scheme by throwing a `std::exception` in case there is a failure to retrieve the configuration information. When we attempt to retrieve configuration data in the POS software, we both test for a null pointer and handle any exceptions thrown.

```

std::map<std::string, PixelDACSettings*> dacs;
dacs["pixel/dac/FPix_BpI_D1_BLD1_PNL1"]=0;
dacs["pixel/dac/FPix_BpI_D1_BLD1_PNL2"]=0;
PixelConfigInterface::get(dacs, theGlobalKey);

```

The method will add pointers to the modules; if a module is not in the configuration a null pointer is returned.

In addition to the access method that uses the configuration key, you can also retrieve data based on the specific version of the object

```

template <class T>
static void get(T* &pixelObject,
                std::string path,
                unsigned int version)

```

There is also a version that retrieves all the objects

```

template <class T>
static void get(std::map<std::string, T*> &pixelObjects,
                unsigned int version)

```

In general the access based on the configuration key should be used in the XDAQ applications.

7.1.2 Storing data in the database

To store data in the configuration database use the following method

```

template <class T>
int PixelConfigInterface::put(const T* ,std::string path);

```

This method will install the data in T* using the specified path. The new version number is returned by the method. If you instead have a set of configuration data that needs to be installed you use the interface

```

template <class T>
int PixelConfigInterface::put(std::vector<T*>,
                             std::string path);

```

Where a vector of pairs is passed, the second argument in the pair is the sub path used for the object.

7.1.3 Configuration keys

A configuration key consists of a set of objects and their versions. The simplest form is just to specify a list of the paths and the versions to use in the key

```
static unsigned int makeKey(std::vector<
                           std::pair<std::string, unsigned int> > versions)
```

The method returns the new configuration key.

7.1.4 Alias manipulation

For configurations aliases you can retrieve the list of defined aliases using

```
static std::vector<std::pair<std::string, unsigned int> > getAliases()
```

The string is the name of the alias and the unsigned int is the corresponding key.

A very similar method is

```
static std::map<std::string, unsigned int> getAliases_map()
```

That returns a map between the alias name and the corresponding key.

The method

```
static void addAlias(std::string alias, unsigned int key)
```

inserts a new alias. Note that if the alias already exists it will just be updated to point to the new key.

The method below is a lower level method. This method allows you to define an alias to a version that has key already created and link it to version aliases. This method should eventually be removed as it is probably to error prone...

```
static void addAlias(std::string alias,
                     unsigned int key,
                     std::vector<std::pair<std::string, std::string> > versionaliases)
```

To create an alias for a configuration object use the method:

```
static void addVersionAlias(std::string path,
                            unsigned int version,
                            std::string alias)
```

To add an alias for a configuration object, as supposed to a configuration key, use

```
static void addVersionAlias(std::string path,
                            unsigned int version,
                            std::string alias)
```

To get the actual version that an alias points to use the method

```
static unsigned int getVersion(std::string path,
                             std::string alias)
```

7.1.5 Support of polymorphism

The interface should support polymorphism in the sense described below. For example, for trim bits it may be convenient to have ways of storing the trim bits in different ways. We will need to be able to store trim bits so that we can set them independently for each pixel. But there are other cases where we might want to set all trim bits the same on a whole ROC, or the same in each double column.

Consider that we have a base class, `PixelTrimBase` and that there are the concrete implementations `PixelTrimAll`, `PixelTrimROC`, and `PixelTrimDCol` which implements the per pixel, per roc, and per double column respectively. The interface should support operations like

```
PixelTrimBase *tempTrims=0;  
PixelConfigInterface::get(tempTrims,"pixel/trim/FPix_BpI_D1_BLD1_PNL1",  
                         theGlobalKey);
```

where after `tempTrims` points to the data type stored in the database for the configuration.

Similarly you should be able to store data

```
PixelTrimBase *tempTrims=new PixelTrimROC; // no such constructor.  
unsigned int ver=PixelConfigInterface::put(tempTrims,  
                                         "pixel/trim/FPix_BpI_D1_BLD1_PNL1");
```

7.2 Command line interface to configuration data

Based on the interface described above in the C++ interface a simple command line tool has been written to allow manipulation of configuration data. The functionality of this interface is described below.

7.2.1 Inserting new data

Here we will start by discussing how you insert a `PixelDetectorConfig` object. There is only one such object in the configuration. Assuming that we have this in a file named `detconfig.dat`. We now want to insert this into the configuration database. This is done with

```
PixelConfigDBCmd.exe --insertData detconfig detectconfig.dat
```

This would install the content of the file `detconfig.dat` under the path `detconfig` by creating a new version. This new version is returned by the command so that it is known where it was installed. The implementation of this tool reads the file to create a `PixelDetectorConfig` object and then uses the C++ interface to store the data.

As a variation of this you might want to install, e.g., a new set of DAC settings for the ROCs. As we insist that no data can be changed in the configuration database after it has been loaded this implies that the DAC settings for all ROCs needs to be loaded at once. To insert that prepare the files that contains the ROC DAC settings. Then create the file daclist.txt that lists all the files that you want to install. To install them into the database use

```
PixelConfigDBCmd.exe --insertDataSet dac daclist.txt
```

Again a new version has been created under the path dac and all the DAC settings have been uploaded. The new version is then printed out.

Note that these interfaces are designed to not allow you to change any existing data. (We might want to consider allowing adding comments to existing versions of configuration data.)

7.2.2 Retrieving data

From the command line you can retrieve data using

```
PixelConfigDBCmd.exe --getVersion nametranslation/ 0
```

This will retrieve the data from version 0 of the nametranslation and write it out as a file.

7.2.3 Creating new configurations

So far we have discussed how to add new data to the configuration database. The next step is to combine versions of several data types into a configuration. This is done with a command like

```
PixelConfigDBCmd.exe --insertConfigAlias Physics dac 1 detconfig 0 nametranslation 0
```

where the versions of the different objects on the paths are listed. This will create a new configuration key. This key will be returned by the command.

Very often we have an existing configuration, `oldKey`, that we want to 'update'. Note that update actually means that we will create a new configuration. For example you can do

```
PixelConfigDBCommand --updateConfigAlias 5 tbm 6
```

So in the example above a copy of configuration key number 5 would be made in which the TBM settings version 6 was added.

If one wants to remove an existing object from a configuration we could imagine doing something like

```
PixelConfigDBCommand --updateConfigAlias 5 calib -1
```

In both these cases the new configuration key that was created is returned. Again this interface guarantees that no existing data has been changed. Only new data has been added.

7.2.4 Alias manipulation

The simplest use case is that we can create an alias that points to a specific key. This could e.g. be to point the 'Physics' alias to configuration key 5

```
PixelConfigDBCommand --insertConfigAlias Physics 5
```

But we need to expand on this to make it more manageable to handle the configurations.

In addition to have aliases for the top level configurations we can define aliases for the configuration data versions. For example you can define the 'Physics' alias for the TBM settings version 4 using

```
PixelConfigDBCommand --insertVersionAlias tbm 4 Default
```

Very similar to the --createconfig from above you can also do

```
PixelConfigDBCommand --insertConfigAlias Physics dac 5 detconfig 3 tbm Physics
```

This will create a new top level configuration key and the alias 'Physics' will point to this key. Say that you then also do

```
PixelConfigDBCommand --insertConfigAlias PhysicsLowLumi dac 4 detconfig 3 tbm Physics
```

We have now created two aliases to two top level keys. Now, if you do

```
PixelConfigDBCommand --insertVersionAlias tbm 5 Physics
```

this should automatically create two new top level configurations for 'Physics' and 'PhysicsLowLumi'.

7.2.5 Use cases

Below we will give some examples for how we could be using the capabilities discussed in the previous two sections. To simplify this a little bit we will work with a configuration that only has 4 different types of objects (detconfig, tbm, dac, and masks).

The first thing we need to do is to load the configuration data. Assume that we do this on a completely empty database.

```
PixelConfigDBCmd.exe --insertData feccconfig feccconfig.dat
PixelConfigDBCmd.exe --insertData fedconfig fedconfig.dat
PixelConfigDBCmd.exe --insertData tkfeccconfig tkfeccconfig.dat
PixelConfigDBCmd.exe --insertData portcardmap portcardmap.dat
PixelConfigDBCmd.exe --insertDataSet fedcard fedcardlist.txt
PixelConfigDBCmd.exe --insertDataSet portcard portcardlist.txt
PixelConfigDBCmd.exe --insertDataSet tbm tbmlist.txt
PixelConfigDBCmd.exe --insertDataSet dac daclist.txt
```

So now we have version 0 for each of these objects. Now we can create aliases for the different objects.

```
PixelConfigDBCommand --insertVersionAlias detconfig 0 Physics
PixelConfigDBCommand --insertVersionAlias tbm 0 Default
PixelConfigDBCommand --insertVersionAlias dac 0 Default
PixelConfigDBCommand --insertVersionAlias mask 0 Default

PixelConfigDBCommand --insertConfigAlias Physics dac Default detconfig Physics
tbm Default mask Default
```

This command will create the first configuration key, number 0. Now say that a new set of DAC settings is created and loaded:

```
PixelConfigDBCommand --insertData dac daclist.txt
```

This will be version 1 of the DAC settings. Next if you make this alias the 'Default':

```
PixelConfigDBCommand --insertVersionAlias dac 1 Default
```

The code will automatically update all top level aliases that are using the 'Default' DAC settings. In this example it means that the top level 'Physics' alias will point to key 1.

7.3 Managing Configurations

We need higher level tools to manage configurations. The way we currently have the configurations implemented there are about 15 different objects that are needed to build a configuration. Two examples of such configurations are:

```
key 0
detconfig 2
nametranslation 0
fedconfig 0
fecconfig 0
fedcard 0
dac 2
mask 2
trim 2
calib 0
tbm 0
portcard 0
portcardmap 0
ttcciconfig 0
tkfecconfig 0
```

```

key 1
detconfig 2
nametranslation 0
fedconfig 0
fecconfig 0
fedcard 0
dac 2
mask 2
trim 2
calib 8
tbtm 0
portcard 0
portcardmap 0
ttcciconfig 0
tkfecconfig 0

```

For a moment we will not discuss the 'aliases', we will focus just on the data organization and the tools needed to manipulate the data. The different data types used in the configurations are discussed in Sect. 8.

8 Configuration Objects

This section describes the implementation of the configuration objects for the pixel online system. The interface for accessing this data was discussed in the previous section.

8.1 Introduction

Key goals for the design of the configuration data object include:

- The configuration has to be fast and reliable. The fewer components, read pieces of software, involved in the configuration step the more likely the system is to work reliably.
- The data volume should be small. I.e., the data has to be packed in an efficient way.
- Want to optimize the database access by accessing relatively few, but large, objects.
- Computers are good at manipulating data that is in memory, so the actual commands that are sent to the hardware can be built on the fly – assuming that all information to do this is accessible in memory.

- The data volume for the whole pixel system is $\mathcal{O}(100 \text{ MB})$ so holding this in memory in a single computer should not be an issue. In fact this will be spread over more than one computer as the FECs are in more than one crate.

The following classes are used to configure the online pixel applications:

PixelTrimAllPixels: This class stores the trim bits for the ROCs on one module. The trims are stored for each pixel.

PixelMaskAllPixels: This class stores the mask bits for the ROCs on one module. The masks are stored for each pixel.

PixelDACSettings: This class stores the DAC settings for all ROCs on one module.

PixelTBMSettings: This class stores the TBM settings for one TBM.

PixelNameTranslation: This class translates from the pixel naming scheme documents names of ROCs to the hardware addresses used by both the FEC and the FED to identify a ROC.

PixelDetectorConfig: This class lists the modules used in a configuration. The utility of this class is that it allows one to only use a small subset of the detector without having to create a new name translation.

PixelROCStatus: This class keeps track of the status of ROCs. The default assumption is that a ROC is working and this object allows us to list ROCs that are not working, or that we want to turn off.

PixelFECCConfig: This class lists the pixel FECs that are used.

PixelTKFECConfig: This class lists the tracker FECs that are used.

PixelFEDConfig: This class lists the pixel FEDs that are used in the configuration.

PixelFEDCard: This class stores the settings for one FED board.

PixelPortCard: This class stores the settings on a portcard, e.g. the delay25 settings and AOH settings.

PixelPortcardMap: This class maps the AOH channels on portcards to the FED channels.

PixelLTCCConfiguration: This class holds the configuration for the pixel LTC module.

PixelTTCciConfiguration: This class holds the configuration for the pixel TTCci module.

PixelCalibConfiguration: This class specifies how the calibrations are executed.

PixelDelay25Configuration: This class specifies how the delay25 calibration is executed.

With the exception of PixelCalibConfiguration and PixelDelay25Configuration the classes above are used to build the configuration of the hardware. The last two classes are used to configure the software applications to perform a given calibration.

The following sections describe what is implemented, starting with a common base class for all configuration objects and then a brief discussion of all classes currently implemented.

All configuration objects derive from a common base class called `PixelConfigBase`.

```
class PixelConfigBase {

public:

    //A few things that you should provide
    //description : purpose of this object
    //creator : who created this configuration object
    //date : time/date of creation (should probably not be
    //        a string, but we have no idea what CMS uses.
    PixelConfigBase(std::string description,
                    std::string creator,
                    std::string date);

    virtual ~PixelConfigBase(){}
    std::string description();
    std::string creator();
    std::string date();

    //Interface to write out data to ascii file
    virtual void writeASCII(std::string dir="") const = 0;

private:

    std::string description_;
    std::string creator_;
    std::string date_;

};
```

This class is intended to provide a common interface. Currently it simply holds information about when the objects were created.

8.2 Trim and mask bits

The trim and mask bits need to be set for each pixel. This makes these the largest configuration objects. In terms of configuration of the ROC, the mask and trim bits are loaded together. However, as the mask bits are used offline we will store the trim and mask bits in different objects as this will reduce the data that needs to be downloaded to access the mask bits offline. Also, this saves space as compared to storing the trim and mask bit for each channel as one byte. It is also likely that we would like to update the mask bits without changing the trim bits. This would happen for example when we discover a hot pixel and want to mask it off. Being able to change only the mask bits would then be useful.

For both the mask and trim bits we allow for different implementations. For example, we can have one implementation that allows us to use different settings for each pixel, or we can have a different configuration that uses the same settings for all pixels. The code handles this transparently by using inheritance.

For the mask bits the base class looks like:

```
class PixelMaskBase: public PixelConfigBase {  
  
public:  
  
    PixelMaskBase(std::string description,  
                  std::string creator,  
                  std::string date);  
  
    virtual ~PixelMaskBase();  
  
    void setOverride(PixelMaskOverrideBase*);  
  
    virtual const PixelROCMaskBits& getMaskBits(int ROCId) const =0;  
  
    virtual void writeBinary(std::string filename) const =0;  
  
    virtual void writeASCII(std::string filename) const =0;  
  
    friend std::ostream& operator<<(std::ostream& s, const PixelMaskBase& mask);  
  
private:  
  
    //Hold pointer to the mask override information.  
    PixelMaskOverrideBase* maskOverride_;
```

```
};
```

The concrete implementation that implements mask bits for each channel looks like:

```
class PixelMaskAllPixels: public PixelMaskBase {  
  
public:  
  
    PixelMaskAllPixels(std::string filename);  
  
    void writeBinary(std::string filename) const;  
  
    void writeASCII(std::string filename) const;  
  
    const PixelROCMaskBits& getMaskBits(int ROCId) const;  
  
private:  
  
    std::vector<PixelROCMaskBits> maskbits_;  
  
};
```

The file format that we use looks like:

```
ROC:      FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1  
col00:    11110000... (80 entries)... 0  
col01:    00000000... (80 entries)... 0  
col02:    00000000... (80 entries)... 0  
. . .  
col49:    00000000... (80 entries)... 0  
col50:    00000000... (80 entries)... 0  
col51:    00000000... (80 entries)... 0  
ROC:      FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2  
col00:    10000000... (80 entries)... 0  
col01:    20000000... (80 entries)... 0  
col02:    30000000... (80 entries)... 0  
. . .
```

Where the file contains the data for each of the ROCs in a module. Within each ROC the trim bits are listed for each column by the value of the trim bit as one hexadecimal character from 0 to F.

Similarly for the trim bits we have the base class:

```
class PixelTrimBase: public PixelConfigBase {  
  
public:  
  
    PixelTrimBase(std::string description,  
                  std::string creator,  
                  std::string date);  
  
    virtual ~PixelTrimBase();  
  
    void setOverride(PixelTrimOverrideBase* trimOverride);  
  
    //Build the commands needed to configure ROCs  
    //on control link  
  
    virtual void generateConfiguration(PixelFECConfigInterface* pixelFEC,  
                                       PixelNameTranslation* trans,  
                                       const PixelMaskBase& pixelMask) const =0;  
    virtual void writeBinary(std::string filename) const =0;  
  
    virtual void writeASCII(std::string filename) const =0;  
  
    virtual PixelROCTrimBits getTrimBits(int ROCId) const =0;  
  
    friend std::ostream& operator<<(std::ostream& s, const PixelTrimBase& mask);  
  
private:  
  
    PixelTrimOverrideBase* trimOverride_;  
};
```

And the concrete implementation looks like:

```
class PixelTrimAllPixels: public PixelTrimBase {  
  
public:
```

```

PixelTrimAllPixels(std::string filename);

//Build the commands needed to configure ROCs
//on control link

void generateConfiguration(PixelFECConfigInterface* pixelFEC,
    PixelNameTranslation* trans,
    const PixelMaskBase& pixelMask) const;

void writeBinary(std::string filename) const;

void writeASCII(std::string filename) const;

PixelROCTrimBits getTrimBits(int ROCId) const;

private:

    std::vector<std::string> rocname_;
    std::vector<PixelROCTrimBits> trimbits_;

};


```

We use basically the same format for the mask bits as we used for the trim bits:

```

ROC:    FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1
col00:  11111111... (80 entries)... 0
col01:  11111111... (80 entries)... 1
col02:  11111111... (80 entries)... 1
.
.
.
col49:  11111111... (80 entries)... 1
col50:  11111111... (80 entries)... 1
col51:  11111111... (80 entries)... 1
ROC:    FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2
col00:  11111111... (80 entries)... 1
col01:  11111111... (80 entries)... 1
col02:  11111111... (80 entries)... 1
.
.
.
```

Here the mask bits are either 0 or 1 for each pixel.

8.3 ROC DACs

The DAC settings for each readout chip are stored in the class

```
class PixelDACSettings: public PixelConfigBase {  
  
public:  
  
    PixelDACSettings(std::string filename);  
  
    PixelROCDACSettings getDACSettings(int ROCId) const;  
  
    //Generate the DAC settings  
    void generateConfiguration(PixelFECConfigInterface* pixelFEC,  
                               PixelNameTranslation* trans) const;  
  
    void writeBinary(std::string filename) const;  
  
    void writeASCII(std::string filename) const;  
  
    friend std::ostream& operator<<(std::ostream& s, const PixelDACSettings& mask);  
  
private:  
  
    std::vector<PixelROCDACSettings> dacsettings_;  
  
};
```

The format for the DAC settings in the ASCII format is given by

| | |
|-------------|---------------------------------|
| ROC: | FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1 |
| Vdd: | 6 |
| Vana: | 140 |
| Vsf: | 128 |
| Vcomp: | 15 |
| Vleak_comp: | 0 |
| VrgPr: | 0 |
| VallPr: | 35 |
| VrgSh: | 0 |
| VwllSh: | 35 |
| VHldDel: | 90 |
| Vtrim: | 29 |
| VthrComp: | 70 |

```

VIbias_Bus:      30
Vbias_sf:        6
VoffsetOp:       30
VIbiasOp:        115
VoffsetR0:       100
VIon:            115
VIbias_PH:       90
VIbias_DAC:      100
VIbias_ROC:      160
VIColOr:         99
Vnpix:           0
VsumCol:         0
VCal:            80
CalDel:          90
WBC:              120
ChipContReg:    0
ROC:             FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2
Vdd:              6
Vana:            140
Vsf:              128
.
.
.

```

Where the file contains the DAC settings for all the ROCs in a given module.

8.4 PixelDetectorConfig

Specifies which components of the detector are used in the configuration. The level of configurability is the module (or in FPIX language the panel). I.e., the components that are controlled by one TBM.

The file format that we use to specify this format looks like:

```

FPix_BmI_D1_BLD1_PNL1
FPix_BmI_D1_BLD1_PNL2
.
.
.
```

The format above is the 'old' format. After discussions with the database GUI developers we have decided to make this object specify the ROC and their status. All ROCs on a module have to be listed. Otherwise there is an internal inconsistency in the configuration. In addition to listing the ROC one can specify a status of the ROC. An example of the file is given below

```

Rocs:
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC0
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2 off
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC3
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC4 noAnalogSignal
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC5
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC6 off noHits
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC7
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC8
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC9

```

The status words are explained in more detail in Sec. 8.5.

The class `PixelConfigurationVerifier` checks the internal consistency of the detector configuration. For instance, it checks that if any ROC on a FED channel is marked with `noAnalogSignal`, then the entire FED channel is similarly marked. Also, it ensures the consistency of the FED card with the detector configuration. If an entire FED channel is marked as `noAnalogSignal`, then the corresponding FED channel is automatically disabled. Similarly, if a FED channel is disabled in the FED card, but enabled in the detector configuration, then the FED card is dynamically modified to conform to the detector configuration. In this way the detector configuration is the “master” flag for what parts of the detector are included in the configuration.

8.5 PixelROCStatus

The `PixelROCStatus` class is used to store the status of ROCs. The default assumption is that a ROC is working and is on. However, we are likely to have problems with some ROCs given the number of components we have. This structure should allow us to add new failure modes as we discover new problems. An example of the data would look like

```

FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2 noHits
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC7 off
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC8 off noHits

```

The different status flags that we can set are:

- `noHits` indicates that we can not generate hits on the ROC. For example this means that the ROC can not be calibrated e.g. for address levels. However, it is not preventing us from doing the UB equalization. In principle this flag should be handled on a calibration-by-calibration basis, but at the moment it is ignored.
- `off` indicates that the ROC is disabled via a control bit on the ROC. It is not used in any calibration and will not generate hits. However, even if a ROC is `off` it will be configured. Presently the use of this flag is not implemented (it is ignored by the code).

- `noInit` indicates that the ROCs in a module should not be included in the configuration. This is implemented.
- `noAnalogSignal` indicates that the ROC can be configured, but that something in the analog readout is broken. The ROC is included in the configuration but excluded from calibrations. This is implemented, and the corresponding FED channel is automatically disabled.

With the file-based configuration, we can set more than one of these flags at once. (Likely one would implement this as a bitmap.) However, the database configuration does not allow more than one flag to be set at a time.

8.6 PixelNameTranslation

This class generates the translation between the names used in the naming document and the hardware addresses. This includes both the FEC and the FED.

The data format used for the name translation is given by:

| # name | TBMchannel | FEC | mfec | mfecchannel | hubaddress | portadd | rocid | FED | channel | roc# |
|---------------------------------|------------|-----|------|-------------|------------|---------|-------|-----|---------|------|
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC0 | A | 1 | 8 | 1 | 31 | 0 | 0 | 1 | 12 | 0 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1 | A | 1 | 8 | 1 | 31 | 0 | 1 | 1 | 12 | 1 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC2 | A | 1 | 8 | 1 | 31 | 0 | 2 | 1 | 12 | 2 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC3 | A | 1 | 8 | 1 | 31 | 0 | 3 | 1 | 12 | 3 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC4 | A | 1 | 8 | 1 | 31 | 0 | 4 | 1 | 12 | 4 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC5 | A | 1 | 8 | 1 | 31 | 0 | 5 | 1 | 12 | 5 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC6 | A | 1 | 8 | 1 | 31 | 0 | 6 | 1 | 12 | 6 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC7 | A | 1 | 8 | 1 | 31 | 0 | 7 | 1 | 12 | 7 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC8 | A | 1 | 8 | 1 | 31 | 0 | 8 | 1 | 12 | 8 |
| FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC9 | A | 1 | 8 | 1 | 31 | 0 | 9 | 1 | 12 | 9 |

The name translation allows us to map the ROC name to the hardware addresses used in the configuration.

8.7 PixelFECConfig

This class specifies the location of the pixel FECs. This basically gives the VME base address to use. An arbitrary FEC number is used here. Also we refer to the crate number here. There will be one PixelFECSupervisor per crate. The PixelFECSupervisor is initialized with a crate number and will control the FEC cards that are in the crate.

The file format that we have to store this information looks like

```
#FEC number      crate      vme base address
1                  1          0x80000000
```

Each FEC is identified by a number. The FEC is identified by the crate and base address.

8.8 PixelTKFECConfig

This class specifies the location of the tracker FECs in the system. This specifies the VME slot and crate used for each TKFEC board. An arbitrary TKFEC ID string is used here.

There will be one PixelTKFECSupervisor per crate. The PixelTKFECSupervisor is initialized with a crate number and will control the TKFEC cards that are in the crate.

The file format that we have to store this information looks like

```
#TKFEC ID      crate      VME/PCI      slot/address
tkfec1          1           0x1c
```

Each TKFEC is identified by an ID string. (Currently this string is arbitrary but we should use an agreed upon convention.)

Optionally, to use a PCI TKFEC, the string “PCI” may be added between the crate number and slot number. “VME” may also be specified. If this parameter is not specified, it defaults to VME.

8.9 PixelFEDConfig

This specifies how the FEDs are configured. This includes the FED number and the VME base address. The FED number is the same as the FED id in the Slink data. Each PixelFEDSupervisor is initialized with a crate number corresponding to the crate it controls.

The file format that we have to store this information looks like

```
#FED number      crate      vme base address
1                  1           0x1c000000
```

Each FED is identified by a number, this is the same number as the FED id in the raw data. The FED is identified by the crate and base address. For now the crate is identified by an arbitrary number.

8.10 PixelCalibConfiguration

This class was formerly known as PixelCalib, but was renamed after it was moved to the CMSSW repository in order to make it more consistent with offline conventions. This class incorporates information about how a calibration is executed. In particular it handles calibrations where groups of pixels have charge injects. It specifies how we loop over pixels and pulse the detector in a calibration. The class is also used in the offline in order to analyze the calibration data, so that we know what event had what charge injected and what pixels were expected to be hit.

Below is an example of this file:

```
Mode: ThresholdCalDelay
Rows:
10 | 20
Cols:
10 | 20
VcalHigh
```

```

Scan: VcThr 0 255 8
Scan: CalDel 0 255 8
Set: Vcal 50
Repeat: 10
Rocs:
FPix_BmI_D1_BLD1_PNL1_PLQ1_ROC0
FPix_BmI_D1_BLD1_PNL1_PLQ1_ROC1
FPix_BmI_D1_BLD1_PNL1_PLQ2_ROC0
FPix_BmI_D1_BLD1_PNL1_PLQ2_ROC1
.
.
.

```

The **Scan** statement allows you to specify that you want to scan the settings of a DAC parameter in a range, above starting at 0 and incrementing in steps of 8 until it exceeds 255. In addition to this you can specify a non-uniform set of scan points using the following format

```

ScanValues: Vcal 10 20 30 35 40 42 44 46 48 50
               52 54 56 58 60 65 70 80 90 100 -1

```

Arbitrary parameters may be specified just before the “Rows:” line. For example,

```

Mode: AOHBias
Parameters:
TargetBMin      412
TargetBMax      612
printFEDRawData      no
printFEDOffsetAdjustments  no
printAOHBiasAdjustments  no
Rows:
.
.
.
```

These parameters are accessible in the calibration code. Each calibration may look for particular parameters to control its operation. Parameters not defined for a particular calibration are ignored.

In particular the parameter “ScanMode” can be defined. It can take the three values of “maskAllPixel”, “useAllPixel”, and “default”. In the default mode the trim and mask bits specified in the configuration is used during the scan. Charge is injected according to the pattern and the pixels that are enabled in the configuration are used during the calibration. In the maskAllPixels mode all pixels are disabled before the first event. Then the pixels are enabled corresponding to the pixel mask and the pattern that has charge injected. That is,

only pixels that have charge injection and that are not disabled in the configuration will be enabled. The pixels use the trim bits from the configuration. In the `useAllPixels` mode, all the pixels on the current pattern are enabled independently of what the mask bit is in the configuration. If `ScanMode` is not defined the mode `useAllPixels` will be used.

DACs to scan are selected with lines of the form:

```
Scan: [DAC name] [min scan value] [max scan value] [scan step size] [mix]
```

The last parameter is optional. If nothing is given here, then all ROCs will be set to the same DAC value at the same time. If the word `mix` is placed here (at the end of the line), then the ROCs on a particular channel will not have the same DAC value at the same time. Instead, the DAC values on different ROCs will be spread out to cover the entire range. This is useful when scanning `Vsf` or any other setting that affects the power drawn by the chip, as it prevents the ROCs from all drawing high power at the same time.

The list of ROCs may be specified completely, or it may be auto-generated. Auto-generation requires knowing which modules are in the configuration, and which ROCs are on those modules. Currently, this information is not available in offline (i.e. CMSSW) code, so offline code must have the ROC list specified completely. To do this, the format is:

```
Rocs:  
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC0  
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1  
. . .
```

In online running, where this information is available, it is preferable to use auto-generation. To auto-generate the list, the first line should be “`ToCalibrate:`” instead of “`Rocs:`”. (The user may specify a complete ROC list with “`ToCalibrate:`”. Compared to using “`Rocs:`”, this has the advantage of not adding ROCs which are not found in the configuration. So it is always recommended to use “`ToCalibrate:`” if possible.)

Probably the most common thing is to just add all ROCs and modules in the configuration. This simply requires the word `all`:

```
ToCalibrate:  
all
```

To specify, say, just 2 ROCs, use:

```
ToCalibrate:  
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC0  
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1
```

To specify, say, all the ROCs on just two modules, use:

```
FPix_BpI_D1_BLD1_PNL1  
FPix_BpI_D1_BLD1_PNL2
```

(The ROCs on each module are obtained from the name translation.)

You can mix and match ROCs and modules:

```
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO  
FPix_BpI_D1_BLD1_PNL2
```

You can also exclude particular ROCs and modules by adding them with a minus sign in front (separated by whitespace):

```
ToCalibrate:  
all  
- FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO  
- FPix_BpI_D1_BLD1_PNL2
```

This fills the ROC list with all ROCs in the configuration, except for FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO and for all the ROCs in FPix_BpI_D1_BLD1_PNL2.

For another example:

```
ToCalibrate:  
FPix_BpI_D1_BLD1_PNL1  
- FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO
```

This adds all ROCs on FPix_BpI_D1_BLD1_PNL1 except for FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO.

You may optionally add a “+” in front of things that you’re adding. (If no “+” or “-” is seen, “+” is assumed.)

```
ToCalibrate:  
+ all  
- FPix_BpI_D1_BLD1_PNL1  
+ FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO
```

This adds all ROCs on all modules, except for FPix_BpI_D1_BLD1_PNL1, on which only FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO is added.

Note that the order matters. The following:

```
ToCalibrate:  
+ all  
+ FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO  
- FPix_BpI_D1_BLD1_PNL1
```

would not include FPix_BpI_D1_BLD1_PNL1_PLQ2_ROCO in the ROC list, because it is later removed by the removal of FPix_BpI_D1_BLD1_PNL1.

(For completeness, even though it’s completely useless, you can use “- all” to clear the ROC list.)

The auto-generated ROC list only adds ROCs on modules given in the `PixelDetectorConfig`.

8.11 PixelFedCard

The `PixelFEDCard` class contains the settings for a FED board. The file is fairly long and we do not include an example here. You can find a sample file in `PixelFEDInterface/test/params_fed.dat`

8.12 PixelTBMSettings

This class holds the settings used by the TBM. The file contains the module name and gain settings as well as if the TBM should be configured in 'SingleMode' or 'DualMode'.

The format used for this information is given by

```
FPix_BpI_D1_BLD1_PNL1_PLQ2_ROC1
AnalogInputBias: 160
AnalogOutputBias: 110
AnalogOutputGain: 207
Mode: SingleMode
```

8.13 PixelPortcardMap

This class lists the portcards used in the configuration and which fed channels are on which AOH.

The format of this information is illustrated by the file:

```
# PortcardName      Module          AOH channel
FPix_BpI_D1_PRT2  FPix_BpI_D1_PNL4 A    1
FPix_BpI_D1_PRT2  FPix_BpI_D1_PNL5 A    2
```

For each portcard the modules are listed and the corresponding AOH channel is listed. The TBM channel, "A" or "B", is specified after the module name. (If no TBM channel is specified, it defaults to "A".)

The AOH channel numbering begins from 1. On each forward pixel port card, there is one AOH with 6 channels, so the numbering goes from 1 to 6. On barrel supply boards, there are 4 AOHs, each with 6 channels. The numbering goes from 1 to 24. The first AOH contains channels 1-6, the second contains channels 7-12, etc.

8.14 PixelPortCardConfig

This class holds the settings that are used on each portcard.

The format of the portcard configuration file is

```
Name: FPix_BpI_D1_PRT1
TKFECID: tkfec1
ringAddress: 0x8
```

```

ccuAddress: 0x7d
channelAddress: 0x10
i2cSpeed: 0x64
Delay25_GCR: 0x0
Delay25_SCL: 0x60
Delay25_TRG: 0x68
Delay25_SDA: 0x5c
Delay25_RCL: 0x60
Delay25_RDA: 0x60
AOH_Bias1: 0x19
AOH_Bias2: 0x1f
AOH_Bias3: 0x1f
AOH_Bias4: 0x1f
AOH_Bias5: 0x1f
AOH_Bias6: 0x1f
AOH_Gain1: 0x2
AOH_Gain2: 0x0
AOH_Gain3: 0x0
AOH_Gain4: 0x0
AOH_Gain5: 0x0
AOH_Gain6: 0x0

```

In the barrel, there are 4 AOH devices, denoted AOH1, AOH2, AOH3, and AOH4. Each of these 4 devices has 6 channels, each with its own bias and gain. For example, AOH1_Bias1, AOH1_Gain5, AOH2_Bias3, AOH4_Gain6.

In addition, values may be specified for the settings PLL_CTR1, PLL_CTR2, PLL_CTR3, PLL_CTR4or5, PLL_CTR4, PLL_CTR5, DOH_Ch0Bias_CLK, DOH_Dummy, DOH_Ch1Bias_Data, and DOH_Gain_SEU. The settings PLL_CTR4or5, PLL_CTR4, and PLL_CTR5 are a special case. PLL_CTR4or5 is the actual address that may be written. If bit 5 of PLL_CTR2 is set to zero, this address writes to PLL_CTR4; if that bit is set to one, this address writes to PLL_CTR5. When PLL_CTR4 or PLL_CTR5 is used in the configuration file, this bit will be automatically set correctly without needing to manually specify a setting for PLL_CTR2. It is recommended to use this feature, instead of using PLL_CTR4or5.

Prior to loading the settings in the file, an initialization sequence is sent to the port card:

```

PLL_CTR1:      0x8
PLL_CTR1:      0x0
PLL_CTR2:      0x20
PLL_CTR4or5:   0x15
DOH_Dummy:     0x0

```

Then the device settings in the file (all lines after i2cSpeed: 0x64) are loaded in the order given.

8.15 PixelDelay25Calibration

This class specifies the parameters used in the delay 25 scan for the delay settings of return data and send data. An example of the configuration files used for this calibration is given below

```
Mode:  
Delay25  
Portcards:  
FPix_BmI_D1_PRT2  
AllModules:  
0  
OrigSDa:  
64  
OrigRDa:  
64  
Range:  
64  
GridSize:  
8  
Tests:  
10  
StableRange:  
6  
StableShape:  
2
```

8.16 PixelGlobalDelay25

The PixelGlobalDelay25 is a class that contains one global delay setting that delays the signal in delay25 chip for the clock and data. In the FED it adds the same delay such that the digitization works independently of the delay setting in the global delay.

The format of this file is very simple, it is just one single number in hex

0x10

The delay specified here is in units of 0.499 ns. This corresponds to the steps of the delay in the delay 25 chip.

9 Usage of configuration data in XDAQ applications

This section contains a brief description of how the configuration data objects are used.

First we look at the PixelSupervisor. RCMS, via the PixelFunctionManager, will pass a string for a configuration alias to the PixelSupervisor during the configure transition. The PixelSupervisor gets the configuration alias and looks up the corresponding configuration key:

```
std::string alias=parametersReceived[0].value_;
unsigned int globalKey=PixelConfigDB::getAliases_map().find(alias)->second;
theGlobalKey_=new PixelConfigKey(globalKey);
```

Having obtained the configuration key this is what is used to extract configuration data. For example the PixelSupervisor extracts some objects:

```
PixelConfigDB::get(theCalibObject_, "pixel/calib/", *theGlobalKey_);
PixelConfigDB::get(theDetectorConfiguration_, "pixel/detconfig/", *theGlobalKey_);
```

The get method returns a pointer that the PixelSupervisor is responsible for deleting.

When the PixelSupervisor asks the other supervisors to configure it does this by passing the configuration key, not the alias, to them. This guarantees that the configuration used by the different supervisors is consistent. We only need to retain the configuration key as a record of how the detector was configured.

Inside the PixelFECSupervisor in the configuration method we have code like

```
PixelConfigDB::get(theNameTranslation_, "pixel/nametranslation/", *theGlobalKey_);
assert(theNameTranslation_!=0);

PixelConfigDB::get(theDetectorConfiguration_, "pixel/detconfig/", *theGlobalKey_);
assert(theDetectorConfiguration_!=0);

PixelConfigDB::get(theFECCConfiguration_, "pixel/feccconfig/", *theGlobalKey_);
assert(theFECCConfiguration_!=0);
assert(theFECCConfiguration_->getNFECBoards()==1); //FIXME

PixelConfigDB::get(theCalibObject_, "pixel/calib/", *theGlobalKey_);
calibStateCounter_=0;

// Loop over all modules in the Detector Configuration and instantiate FECInterfaces re
// Download TBM, DAC, Masks and Trim settings into hardware.
std::vector <PixelModuleName>::iterator module_name = theDetectorConfiguration_-
>getModuleList().begin();
for (;module_name!=theDetectorConfiguration_->getModuleList().end();++module_name)
{
    diagService_->reportError("Configuring module=" + module_name->modulename(),DIAGDEBUG)
    const PixelHdwAddress* module_hdwaddress=theNameTranslation_->getHdwAddress(*module_na
```

```

unsigned int fecnumber=module_hwaddress->fecnumber();
unsigned int feccrate=theFEConfiguration_->crateFromFECNumber(fecnumber);
unsigned int fecVMEBaseAddress=theFEConfiguration_->VMEBaseAddressFromFECNumber(fecnum

if (feccrate==crate_){
    PixelMaskBase *tempMask=0;
    PixelTrimBase *tempTrims=0;
    PixelDACSettings *tempDACs=0;
    PixelTBMSettings *tempTBMs=0;
    std::string modulePath=(module_name->modulename());

    PixelFECInterface* tempFECInterface=new PixelFECInterface(fecVMEBaseAddress, aBHandle
    assert(tempFECInterface!=0);
    tempFECInterface->setssid(4);

    PixelConfigDB::get(tempMask, "pixel/mask/"+modulePath, *theGlobalKey_);
    assert(tempMask!=0);
    theMasks_.insert(make_pair(*module_name, tempMask));

    PixelConfigDB::get(tempTrims, "pixel/trim/"+modulePath, *theGlobalKey_);
    assert(tempTrims!=0);
    theTrims_.insert(make_pair(*module_name, tempTrims));

    PixelConfigDB::get(tempDACs, "pixel/dac/"+modulePath, *theGlobalKey_);
    assert(tempDACs!=0);
    theDACs_.insert(make_pair(*module_name, tempDACs));

    PixelConfigDB::get(tempTBMs, "pixel/tbm/"+modulePath, *theGlobalKey_);
    assert(tempTBMs!=0);
    theTBMs_.insert(make_pair(*module_name, tempTBMs));

    tempDACs->generateConfiguration(tempFECInterface, theNameTranslation_);
    tempTBMs->generateConfiguration(tempFECInterface, theNameTranslation_);
    tempTrims->generateConfiguration(tempFECInterface, theNameTranslation_, *tempMask);

    FECInterface[fecVMEBaseAddress]=tempFECInterface;
}
}

```

Similar extractions of the configuration data is used by other supervisors. The supervisors cache the data received. In general, configuration data is cleared in the halt transition. However, we now hold on to most of the configuration data in the FECSupervisors, and only clear it before the next configuration if we see that the value of the global key has

changed. In this way, we can avoid reloading identical data from the database on the next configuration.

9.1 Global Delay 25 usage

The global delay 25 does not need to be included in a configuration. Then it is simply ignored if the PixelConfigInterface::get call returns a null pointer. However, if it exists in the configuration it has the following effects in the different applications:

In the PixelTKFECSupervisor it is checked that we are in a physics run and have the global delay 25 settings. If this is the case then SCL and TRG are delayed by the delay setting in the global delay 25. It is checked in the code if the calculated delay 25 setting is larger than 127. If this is the case then an error message is printed and the delay 25 without the global delay is applied. For the SDA, if the calculated delay25 setting is too large, it allows the SDA to wrap around.

In the PixelFEDSupervisor a similar logic is applied; if you have the global delay 25 and you are taking a physics run then the TTC RX chip adds a delay. This delay is calculated using the method `PixelGlobalDelay25::getTTCrxDelay`.

10 Configuration

The configuration of the detector starts with a global key, which encodes the set of configuration data that should be used. The PixelSupervisor then coordinates the actions of the other supervisors. Configuration data must be fetched from files or the database, and then the data must be loaded into the detector hardware. This happens in several steps that must happen in a particular order. The logic governing this process is described below.

10.1 Configuration Sequence

The configuration sequence is listed below.

- PixelSupervisor::Configure
 - Translate the configuration alias into the Global Key
 - Send preConfigure command and Global Key to the PixelFECSupervisors
- PixelSupervisor::transitionHaltedToConfiguring
 - Fetch other miscellaneous configuration data from the database
- PixelSupervisor::stateConfiguring
 - Configuration of PixelDCSFSMInterface
 - (wait for completion)
 - Parallel configuration of PixelTKFECSupervisors
 - (wait for completion)
 - Parallel configuration of PixelFECSupervisors
 - Parallel configuration of PixelFEDSupervisors
 - (wait for completion)
 - Configuration of TTCCiControl
 - Configuration of PixelLTCSupervisor
 - Configuration of PixelDCStoTrkFECDpInterface

The list covers the steps visible to the PixelSupervisor. Each individual supervisor also has several internal configuration steps that are not described here. Also note that the only configuration information passed from the PixelSupervisor to the other supervisors is the global key. Each supervisor is responsible for fetching its own configuration data from the database, based on this global key.

As shown in the list, for several of the configuration steps we must wait for them to complete before moving on to the next step. This is implemented using objects of type

`std::map<instance, state>` maintained by the PixelSupervisor. These objects track the state of the underlying supervisors⁶, and are called:

```
statePixelFECSupervisors_
statePixelFEDSupervisors_
statePixelTKFECSupervisors_
statePixelDCSFSMInterface_
```

These are initialized in `PixelSupervisor::Initialize` by actively asking each supervisor with a SOAP command “`FSMStateRequest`”.

Note that the first step in the configuration after loading the global key is to send it to the FECSupervisors. Although the FECSupervisors cannot begin programming the hardware until later in the configuration sequence, they can then begin to immediately fetch configuration data from the database. This step, called `preConfiguration`, is done outside of the finite state machine structure (the PixelFECSupervisors remain in the `Halted` state even as they do the `preConfiguration`). This step saves considerable time in the configuration process (roughly 25 seconds when configuring from the database).

In the `PixelSupervisor::stateConfiguring` method, the PixelSupervisor loops over all the PixelTKFECSupervisors, checks the FSM state of each as maintained in `statePixelTKFECSupervisors_`, and if it finds “`Halted`”, it tries to “`Configure`” it and updates the map to read FSM state “`Configuring`”. If any of these PixelTKFECSupervisors are not in the “`Configured`” state, the `::stateConfiguring` function exits. The PixelSupervisor then does nothing until it receives a “`FSMStateNotification`” SOAP message from PixelTKFECSupervisor with the message “`Configured`”. This `FSMStateNotification` signals the PixelSupervisor to make a transition from “`Configuring`” to “`Configured`”, thus triggering `::stateConfiguring` to run again. In this way the PixelTKFECSupervisors configure in parallel, and can take as long as they like to finish configuration.

Once the TKFECSupervisors are done, we can configure the pixel FECs and FEDs in parallel. The procedure is similar: if their FSM state is “`Halted`”, we give them the SOAP message to `Configure` and set their local map state to “`Configuring`”. And now we can configure the FEDs too in a similar, parallel, manner. (Note that we do not return out of the function directly, but rather if we detect one of the supervisors to not be in “`Configured`”, we set a bool called “`proceed`” to false, and this bypasses a large chunk of code that configures the LTC and TTC, and transitions the FSM of PixelSupervisor to “`Configured`”.) Only when all FECs and FEDs are configured, then we proceed to configure the TTC & LTC, and then the PixelSupervisor is pushed into the “`Configured`” state.

10.1.1 Possible modifications to the configuration Sequence

In principle there is nothing to prevent the TTC and LTC from being configured before the other Supervisors. Arguably, this would be more logical, since the trigger source should be configured before the other devices are programmed.

⁶Note that these objects are useful for more than just configuration, and are always kept updated with the state of each Supervisor

10.2 Configuration steps of the underlying supervisors

In this section we give an outline of the steps taken by each supervisor during configuration.

10.2.1 PixelDCFSMInterface and PixelDCSToTrkFECInterface

The PixelDCFSMInterface loads the detector configuration. If it finds that an entire ROG is in `noInit` or `noAnalogSignal`, then it ignores that ROG when summarizing the power state of that section of the detector. Note that initially this was true only for “`noInit`”, since a ROG marked `noAnalogSignal` will still be configured and thus needs LV to be on. This logic was changed to accomodate turning on only a small fraction of the HV during the first beam operations.

Note that presently the relationship between detector modules (as listed in the `detconfig`) and DCS ROGs is hard-coded in this class. See Sec. 15.9 for more information.

10.2.2 PixelTKFECSupervisor

- Create a `FecAccess` object and issue a VME bus reset
- Issue a `resetPlxFec` to reset the CCU and portcard devices
- Load the portcard configuration data from the database and program the portcard devices (AOH, DOH, Delay25, etc)

Note that in this last step the data is programmed to the hardware in the same order that it is provided by the class that loads the data from the configuration files or database. That class passes file-based data directly, in the same order as it appears in the files. For the database, it applies a crude sorting algorithm to put PLL settings first, Delay25 settings second, and AOH settings last. Any other settings come between the Delay25 and AOH settings.

10.3 Quick reconfiguration for the fine delay scan

A special “Reconfiguration” option is available for changing the settings that are directly relevant to the global delay of the detector with respect to the trigger and the clock.

10.3.1 Relevant settings

The following settings are modified during the reconfiguration process:

- PixelTKFECSupervisor: the SCL, TRG, and SDA registers of the Delay25 chip
- PixelFECSupervisor: the WBC DAC setting. Note that in the current implementation, **all** DAC settings are reloaded and reprogrammed during reconfiguration.
- PixelFEDSupervisor: the TTCrx delay

10.3.2 Implementation

This feature is implemented outside of the normal finite state machine. Instead, the PixelSupervisor displays a list of configuration aliases when we are in either the Configured state (i.e. after Configuration is done, but before receiving the Start command) or the Paused state. This list of aliases is only shown if we have been configured for Physics, and if a special environment variable is set. Only aliases that begin with the string “Physics” are shown.

The user can then select one of the displayed aliases and click “Reconfigure”. The reconfiguration process takes about 20 seconds, and when it is finished the global key shown in PixelSupervisor will update to the new value.

10.3.3 Limitations

The current implementation has the following limitations:

- The alias used for Reconfiguration must have a `globaldelay25` object defined as part of the configuration. Note that this is stricter than the requirement for regular configuration.
- If the Reconfiguration fails for some reason, then the PixelSupervisor will likely still update the displayed global key. The user should be careful to check the LogReader or log file for an error message.

11 Event builder configurations

We do not want to build complete events on one node; rather we want to make sure that the event fractions from each FED is always shipped to the same filter unit.

To accomplish this the network fabric of the event builder will need to be reconfigured for each pixel calibration run. Note that this is likely to be done between each LHC fill. (This does not mean that we will determine all constants at this rate, but that we will want to monitor quantities such as address levels frequently.)

We have developed algorithms that analyze the pixel raw data for the purpose of doing gain calibrations, pixel alive, and Scurves. Initially this work is done using data recorded in plain files. To progress towards the final goal we should first try to run this with a filter unit that is fed data from a source that reads files. This will allow us to test and develop the application in the form it will run in the HLT. The second step is to actually run this with data produced from a FED. This can be either real data or emulated playback generated in a DAQ test stand.

We should also test and deploy code so that we can take data using local DAQ (VME) and send it out to an event builder.

12 Directory Structure

Class PixelSupervisorConfiguration Base helps to organize the output directory, for example: the output directories with runNumber between 2000 and 2999 will be stored in directory Run_2000. PixelSupervisorConfiguration, PixelFEDSupervisorConfiguration and PixelTK-FECSupervisorConfiguration are derived from this class.

12.1 Tree structure

Tree structure is added in AddressLevel, ROCUBEequalization, VsfAndVHldDel, LinearityVsVsf, PHRange, VcThrCalDel, VcThr and CalDel calibrations (among others). It helps to view some important information directly through Histoviewer.

Two TTree pointers, tree and tree_sum, are used to store the information of "pass", and the summary information, respectively. Table 5 lists the pass condition and summary information for the above calibrations. The implementation of the tree structure is to declare the struct in the header file of each kind of calibration. In the struct there must be a char array to store each ROC's name. This ensures that all the information can be viewed by Histoviewer. The following code example is from ROCUBEequalization calibration:

```
typedef struct branch{
    float pass;
    char rocName[38];
};

typedef struct branch_sum{
    float new_VIBias_DAC;
    float delta_VIBias_DAC;
    char rocName[38];
};
```

Table 5: Pass condition and summary information stored in the trees. If the condition for failure is satisfied, then the "Pass" state is stored as 0 in the summary tree. Otherwise the "Pass" state is stored as 1.

| Calibration | Condition for Failure | Summary Information |
|----------------|--|--|
| Address Levels | Recommended Level 0 [†] is less than UltraBlack high threshold or Number of ROC peaks found is not equal to 6 | number of peaks, maximum peak RMS minimum peak separation, RMS of the black levels |
| ROC UB | The measured UB does not cross the target level | new VIBias value and change of VIBias |
| Vs VHldDel | no hits found or no good Vsf found during scan | new Vsf value and the change of Vsf new VHldDel and its change |
| LinearityVsVsf | no hits found or no good Vsf found during scan | new Vsf value and its change |
| PHRange | no hits found or no settings give PHInRange or only one Vcal produces hits | new VIBias.PH, VOffsetOp and their changes new VIon and VOffsetRo and their changes |
| VcThrCalDel | there is not valid settings | new VcThr, CalDel and their changes |
| VcThr | there is not valid VcThr | new VcThr and its change |
| CalDel | there is not valid CalDel slope | new CalDel and its change |
| Iana | new Vana \geq 249 or maximum Iana < 25 | new Vana, change in Vana, new Iana, maximum Iana (fit Iana at Vana = 250), χ^2 of the fit |

13 Calibration Algorithms

The two main purposes of the online software are the configuration for data taking and the online calibrations that are performed between runs for the calibration of the detector. There are a large number of different calibration tasks that need to be performed. This section describes the algorithms implemented to carry out these calibrations.

A brief summary of the most fundamental calibrations is given here, roughly in order of use. In the next section, Sec. 14 we document our experience calibrating the detector using these calibrations.

The Delay25 calibration is used to ensure correct FEC communication; it is very fundamental but once the settings are found they do not need to be readjusted often. We usually start by running the FEDBaselineCalibration, because it runs quickly and provides an informative plot of the data buffer for each FED channel. It also needs to be run relatively frequently to adjust the black level. If the BaselineCalibration fails to converge, one can run the AOHBias calibration to more coarsely adjust the black levels, then subsequently try the BaselineCalibration again. The ClockPhase calibration is used to adjust the timing of the signal digitization; a bad clock phase can also cause problems for most other calibrations. The ultrablack levels are adjusted for the TBM using the TBMUB calibration, and then the ROC UB levels are adjusted with the ROCUB calibration. (Note that it is never harmful to repeat the BaselineCalibration as one progresses through these steps.) The AddressLevels calibration is then run in order to be able to decode pixel hits properly. If this calibration has difficulty seeing hits, it may be necessary to adjust the threshold and calibration injection delay using the VcThrCalDel scan (or a related calibration). Once the AddressLevels are found one can test the detector using the PixelAlive or SCurve scans.

More details are given below, along with descriptions of additional calibrations not included in the above list.

13.1 AOH and FED channel mapping test

The AOH and FED channel mapping test is not really a calibration, but rather a test to see whether the connections between AOH channels and FED channels in the configuration are correct. The idea is to change the AOH bias for a single channel, and look to see whether the black level on the corresponding FED channel changes. (FED automatic baseline correction must be turned off.) This process is repeated for each channel. Those for which the black level changes are correctly connected; those for which there is no change are not connected correctly or have a mistake in the configuration files.

13.1.1 Mapping test steps

The steps of this calibration are listed in order below.

1. Set the FED channels to the 2V peak-to-peak range in order to provide maximum dynamic range for the AOH bias scan.

2. Turn off the FED automatic baseline correction.
3. Set all FED optical receiver input offsets to 8, and channel offsets to 255. These settings tend to put the signal in a range where the black level will not be stuck at the top or bottom of the range as AOH bias is varied. (Maybe this step should be removed, and the default values should be used instead.)
4. Issue `ClrCal` to all ROCs, and disable hits with the control register on all ROCs, to ensure that no hits are output.
5. Issue `LRES` and `CLRES` commands to all FEDs to clear the transparent data. This ensures that no stale data is sitting in the buffer.
6. On the first channel, loop over AOH bias values. On each trigger, inspect the transparent buffer and look for the start of the signal. Record the black levels in the time slots preceding the start of the signal. Keep track of the mean black level and the standard deviation for each AOH bias value. A plot of this scan is written to a ROOT file.
7. Inspect the scan of black level vs. AOH bias. If the black level varies as AOH bias changes, this channel is properly connected. If the black level is constant, this channel is flagged as improperly connected. (Quantitatively, the χ^2 of the scan data is evaluated for the hypothesis of a flat line. If $\chi^2/\text{d.o.f.} > 100$, the scan data is considered to vary; otherwise it is considered constant.) One exception is if the scan is constant at the top or bottom of the range (1023 or 0). In this case the signal is out of range, and so the connectivity test is inconclusive.
8. Set AOH bias back to the default value (from the configuration). This ensures that channels are not falsely flagged as being stuck at the top of the range. If, say, two FED channels were swapped in the configuration, the first one scanned would show up as incorrectly connected. With the AOH left at the highest scan point, when the FED channel that is actually connected to that AOH is checked later, the black level will appear to be stuck at the top of the range. Setting that AOH back to its default value should put the black level on that channel back into the middle of the range.
9. Repeat the AOH bias scan for each channel in the configuration.
10. Print a summary to the screen. This summary gives the number of channels which are properly connected, improperly connected, or for which the test was inconclusive. Improperly connected and inconclusive channels are then listed.

13.1.2 Parameters

A number of parameters in `calib.dat` may be used to control the test.

Table 6: Optional parameters for AOH and FED channel mapping test.

| Parameter | Default | Description |
|-----------------|---------|--|
| ScanMin | 0 | Low end of AOH bias scan range |
| ScanMax | 50 | High end of AOH bias scan range |
| ScanStepSize | 5 | Step size for AOH bias scan |
| printFEDRawData | no | Whether to print decoded transparent buffer |
| printScan | no | Whether to print the AOH bias scan to the screen |

Only two “standard” parameters are used. The “Repeat:” parameter determines the number of triggers at each AOH bias scan point. The channel list is used to determine which channels are calibrated. Note that the rows, columns, and DAC scan settings in `calib.dat` are completely ignored.

Some optional parameters may also be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 6.

13.2 FED phase and delay scan

The FED has a delay (0 to 15) that goes in steps of 25/16 ns. In addition to the delay there is a phase that controls when the data is latched. Certain combinations of the delay and phase are invalid and result in garbage ADC values.

In this calibration the 32 values of the phase and delay are scanned. For each setting of the phase and delay a fixed number of events (typically around 10) is read out in transparent mode.

A detailed description of this algorithm, is presented in Appendix A.

13.2.1 Output

The FED phase and delay calibration produces new `fed_params.dat` files that are updated with the new FED settings for the phase and delay. The files should otherwise be identical to the input files specified in the configuration. In addition to the `fed_params.dat` files there are several canvases generated in the output ROOT file. For each channel you get four output canvases:

```
PhaseAndDelayRaw_37_1
PhaseAndDelayPurged_37_1
PhaseAndDelayOrdered_37_1
PhaseAndDelayFinal_37_1
```

Where the numbers in the file name represent the fed id and the fed channel. Appendix A gives examples of these plots.

13.2.2 Example configuration

An example of a configuration file to run the phase and delay scan is given below.

```
Mode: ClockPhaseCalibration
Rows:
Cols:
Vcal:
100 100 5
Repeat:
10
ToCalibrate:
all
```

As described in Appendix A one can change the algorithm used to find the best phase and delay using

```
Parameters:
oldMode Yes
```

13.3 Delay25 settings for send data and return data

This calibration scans the delay for the send data and return data. For each set of values commands are sent to the TBM and the return status in the (pixel) FEC is checked. This calibration uses the PixelFECInterface::testDelay25 method that sends 4 different types of commands. This algorithm combines the results from the 4 different types of commands and only if all 4 succeeded will you get 100% efficiency. Note that one of the commands to be tested is read from the file `infeccmd.dat` in the `PixelRun` directory. If necessary, one can customize this file.

Note that this algorithm does not check that the TBM or ROC actually received the command, it just checks the return status.

13.3.1 Output

The main output from this calibration is new delay settings that are stored in the portcard files. The only changes made to the port card settings are for the send data and return data delays.

In addition to generating the new portcard files with delay settings the scans are plotted in a ROOT file. An example plot is shown in Fig. 13.

A plot from the Cornell test stand is shown in Fig. 14.

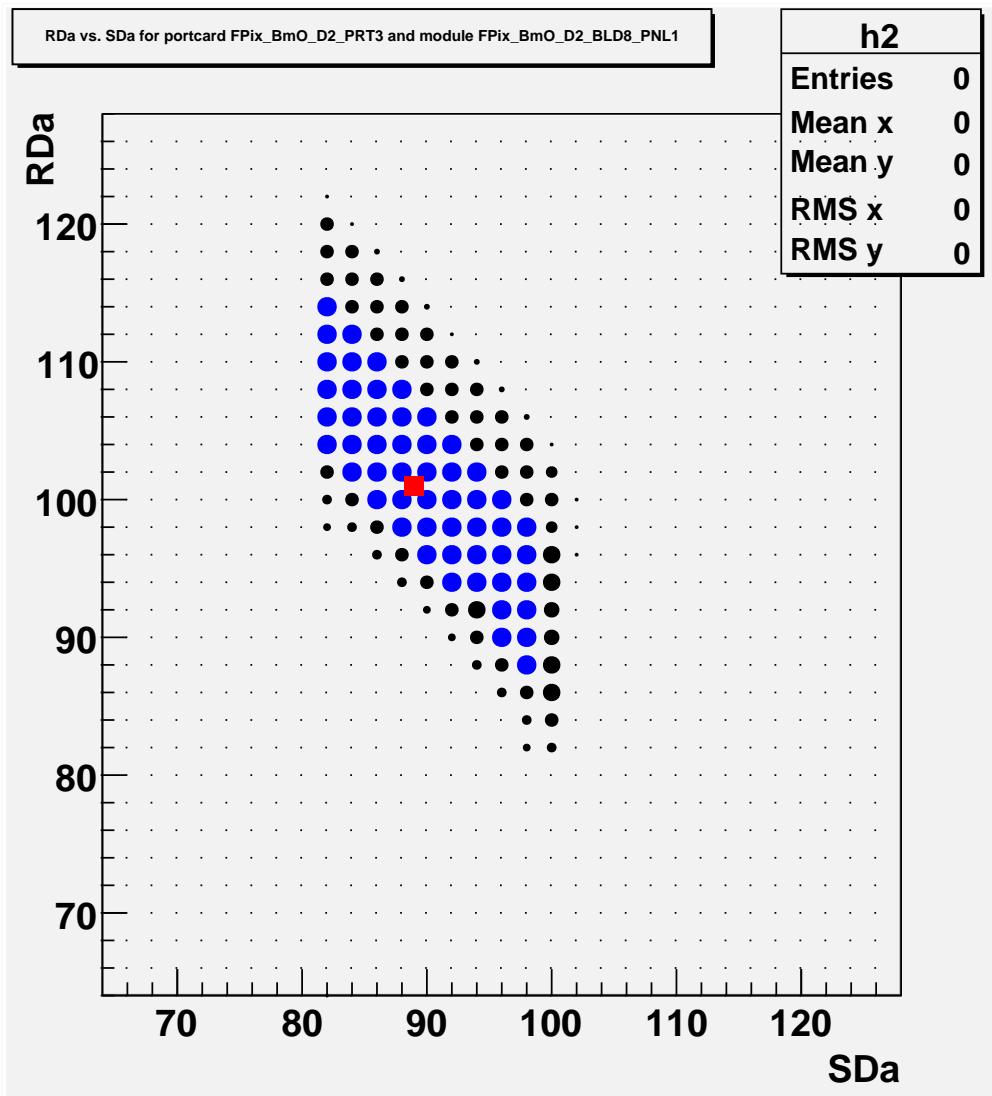


Figure 13: This plot shows efficiency as a function of RD_a and SD_a. The blue dots indicates areas with 100% transmission efficiency. The black dots indicated partial efficiency, larger dots have higher efficiency. The red square indicates the point chosen by the algorithm.

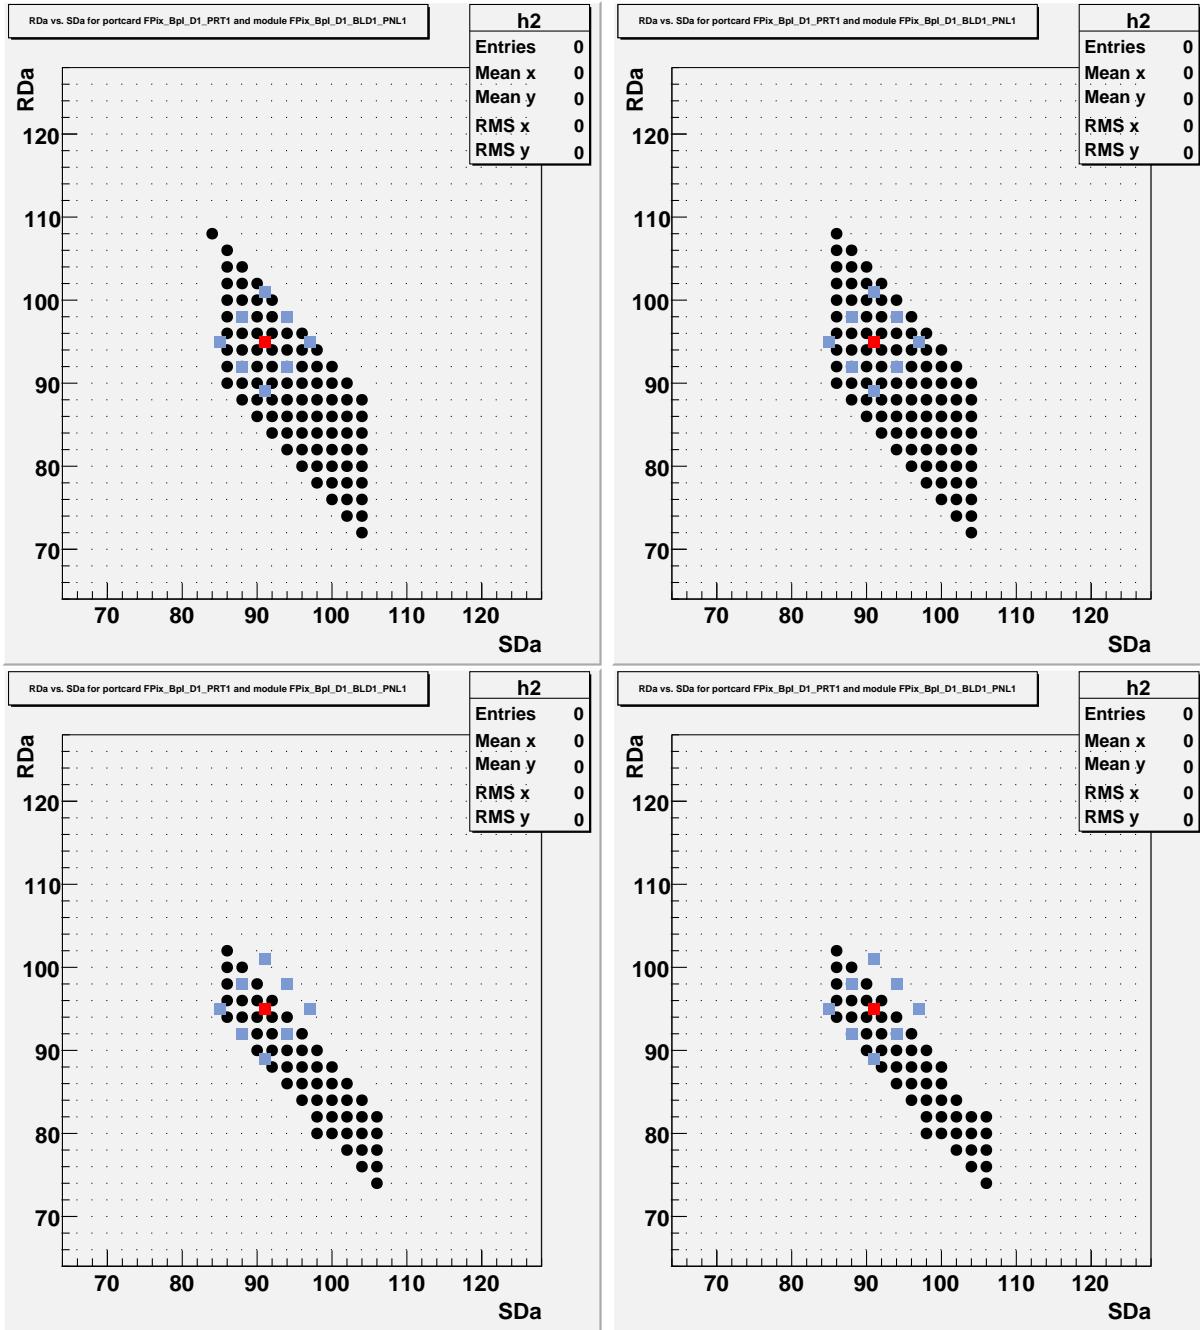


Figure 14: This figure shows the results of scans for delay25 settings on the Cornell test stand. The upper left plot shows (large black dots) the region for which the return data was valid when sending a ROC command (CalPix). The upper right plot shows the valid region when sending a TBM command (tbm speed). The lower left shows the region of success when sending a roc init and the lower right shows the region of success with a roc trim load command. As is seen, the working region is smaller for the long commands. The red and blue points indicates the algorithm used to select the operating point.

13.4 Delay25 trigger setting

Scans the trigger delay to make sure that the triggers are received correctly. The delay setting is scanned and the corresponding FED channel is read to make sure the trigger has arrived. This has to be done before the phase and delay scan of the FED.

13.5 FED baseline calibration

This calibration adjusts the input offset and channel offsets of the optical receivers in the FED such that the black level is adjusted to be near a given target value, normally 450, which is near the midpoint of the dynamic range of the ADC. During this calibration the baseline correction in the FED is turned off.

Besides determining the input offset and the channel offsets the algorithm determines address levels for the black and ultra-black levels.

13.6 AOH bias settings

13.6.1 Introduction and discussion

The AOH bias is a setting on the port card which controls the amount of light sent to the FED. There is one AOH bias setting per FED channel. As AOH bias increases, more light is sent, and the ADC values on the FED increase. At low values of AOH bias, both black and ultrablack do not change with AOH bias, and there is no separation between black and ultrablack levels. At some threshold, the black level begins to increase approximately linearly. At a higher threshold, the ultrablack level also starts to increase linearly with approximately the same slope. This behavior is illustrated in Fig. 15.

Note that the maximum black-ultrablack separation depends on how the TBM DACs are set. At low DAC settings, the TBM outputs a signal with relatively low separation; as these settings increase, the separation also increases. In the AOH bias scan, the black level is independent of the TBM settings. However, the linear rise of the ultrablack level begins at a later point for higher TBM settings, and hence the black-ultrablack difference saturates at a higher AOH bias value when the TBM DAC settings are higher.

The goal of the AOH bias calibration is to determine an AOH bias setting for each channel that is just high enough to saturate the black-ultrablack difference. The calibration measures this difference, using black and ultrablack levels from the TBM header and trailer, as a function of AOH bias. It is important, though, that during the scan the TBM DACs are set at least as high as they will be set in later calibrations and physics runs. Otherwise, the AOH bias value determined from the saturation point will be too low. TBM settings above those used in this scan will not increase the B-UB separation because the AOH cannot provide more separation.

Temperature variations alter the response of the AOH, essentially shifting the curves in Figure 15 to the left or right. In order to provide a margin of error for temperature changes, the AOH bias should be set higher than the saturation value. A temperature increase of 5

FED 1, channel 12: TBM B & UB vs AOH bias

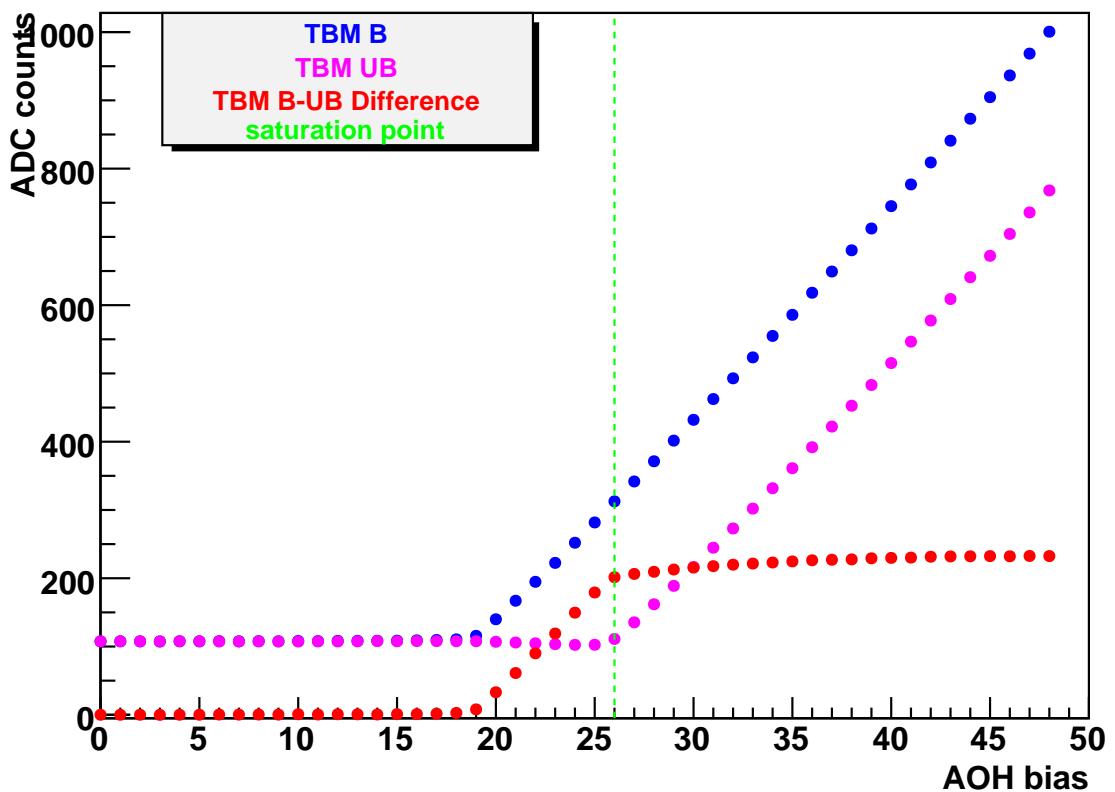


Figure 15: Black and ultrablack levels as a function of AOH bias.

degrees Celcius will shift the curves by about 4 AOH bias counts. Therefore, by default the chosen AOH bias setting will be 4 counts higher than the saturation value. This offset is a configurable parameter.

It is also important that the AOH bias not be too high; otherwise the FED offsets could not bring the signal into the dynamic range of the FED.

The last part of the AOH bias calibration is to do a coarse baseline adjustment. The FED channel offsets are set to the center of the range (127), and then the FED optical receiver offsets and AOH bias settings are adjusted to bring all FED baselines into a wide target range. AOH bias is not decreased below the saturation value unless it is absolutely necessary. The end result is a configuration of AOH bias and FED offset values that puts all FED baselines near the center of the dynamic range, with AOH bias values that allow for a large B-UB separation. After the AOH bias calibration, the FED baseline calibration should be run to perform fine adjustments of the baseline (using the freedom to move each channel offset).

13.6.2 AOH bias calibration steps

This calibration involves many distinct steps. They are listed in order below. Each step is performed on each channel being calibrated.

1. Set the FED channels to the 2V peak-to-peak range in order to provide maximum dynamic range for the AOH bias scan.
2. Turn off the FED automatic baseline correction.
3. Issue `ClrCal` to all ROCs, and disable hits with the control register on all ROCs, to ensure that no hits are output.
4. Set all TBM DACs to high values. These values may be specified as parameters in `calib.dat`.
5. Set all FED optical receiver input offsets to the highest useful value (the setting that minimizes the FED ADC values, without impairing the B-UB separation). This value is configurable, and it defaults to 8 (on a scale from 0 to 15). Also, all channel offsets are set their maximum value, 255.
6. Issue `LRES` and `CLRES` commands to all FEDs to clear the transparent data. This ensures that no stale data is sitting in the buffer.
7. Loop over AOH bias values. Attempt to decode the transparent buffer on each trigger, assuming the correct number of ROCs and no hits. That is, find the TBM header, and verify that the TBM trailer is in the right place. When decoding is successful, record the start slot of the TBM header and trailer. On each channel, this slot should be the same for all triggers. The reason for this step is to find the time slots for later use in reading out TBM B and UB levels, even at low AOH bias settings where full decoding

would fail. If no reliable time slots are found⁷ on a channel, this channel is considered failed, and it is ignored in the rest of the routine.

8. Now scan over AOH bias again, this time to record the TBM B and UB levels at each setting, using the time slots recorded in the previous step⁸. During the scan, if the signal goes out of range high or low, the FED channel offset is adjusted to bring it back into range, if possible. Since only the B-UB difference is of interest, coherent shifts in B and UB do not matter. The FED optical receiver offset is not adjusted during the scan – it remains at the high setting described in step 5.
9. Output plots of the TBM black and ultrablack, and the difference, as a function of AOH bias, to a ROOT file. Figure 15 is an example of the plots produced. Find the AOH bias value at which the B-UB difference saturates (defined as a reduction in the slope to less than 20% of its maximum value). Set each AOH to its saturation value plus the offset defined by the parameter `SaturationPointOffset` (which defaults to 4).
10. Set FED channels back to the 1V peak-to-peak range, for those channels which were originally set at 1V in the FED configuration file. This is done so that the coarse baseline adjustment will be done with the range that will be used for future data-taking.
11. Set FED optical receiver input offsets to 0 (lowest value, corresponding to highest FED ADC values), and all channel offsets to 127 (middle of the range). The channel offset will be left at 127 for the rest of the routine.
12. Measure black levels on all channels. On each FED optical receiver, if any channel has a black level above the target range, increase that receiver's offset by one. If not, do nothing. Repeat this until all channels have black levels in or below the target range, or have a receiver offset equal to the maximum value described in step 5 (defaults to 8). The idea here is to ensure that no AOH bias value will have to be decreased to place the black level in the target range (unless this is absolutely necessary because the receiver offset cannot be increased further).
13. Again measure the black levels on all channels. If a channel's black level is within the target range, that channel is done. If it is above or below the target range, decrease or increase AOH bias. Repeat this until all channels are within range. (Or, if two adjacent AOH bias values produce black levels that straddle the target range, choose

⁷A time slot is considered found if the most common start slot occurs on at least 95% of the triggers. Alternatively, if the two most frequent time slots differ by 1 (due to a jumping clock) and together account for at least 95% of triggers, then the time slots are considered found, and TBM B and UB are sampled only in those time slots which are known to be B or UB for either of the two start slots.

⁸If good time slots were not found on a given channel, only the black level will be recorded and plotted. It is taken from the first 10 slots of the transparent buffer, rather than from the TBM header and trailer. This plot is for diagnostic purposes only; it is not used in generating new configuration settings.

the one that is closer to the target.) AOH bias should not have to decrease unless the FED receiver input offset was at maximum. If it does decrease, a warning is generated in the summary at the end of the calibration.

14. Turn FED automatic baseline correction back on.
15. Write new config files for all port cards and FEDs with at least one successfully calibrated channel.
16. Print a summary to the screen. Thus summary gives the number of successful and failed channels, statistics on the new settings for successful channels, descriptions of the errors on unsuccessful channels, and warnings for channels on which the final AOH bias setting is below the saturation point.

13.6.3 Parameters

A number of parameters in `calib.dat` may be used to control the calibration.

Only two “standard” parameters are used. The “`Repeat:`” parameter determines the number of triggers in each step (at each AOH bias scan point and when measuring black levels in the coarse baseline adjustment). The channel list is used to determine which channels are calibrated. Note that the rows, columns, and DAC scan settings in `calib.dat` are completely ignored.

Many other optional parameters may be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 7.

13.7 AOH gain calibration

13.7.1 Introduction and discussion

The AOH gain is a setting for each optical link (from detector to FED) that has just 4 possible settings (0, 1, 2, 3). This setting does not change the black level. Instead, it scales the size of deviations from the black level, expanding or shrinking the signal. Larger settings correspond to larger deviations. In particular, the separation between black and ultrablack levels will be larger at a larger gain. Settings on the TBMs and ROCs will be the primary means of adjusting the ultrablack to the desired level, but the AOH gain must be set large enough in order to make possible a low enough ultrablack. However, AOH gain should not be set too high since larger settings will increase the power drawn, and larger settings are intended to be used to compensate for radiation damage over time.

The aim of this calibration is to set the AOH gain at the lowest level that will allow the TBM UB level to be low enough. The three TBM DACs (described in Sec. 13.8) are set to high values. Then, for each FED channel, TBM UB is recorded as a function of AOH gain. To choose an AOH gain setting, the calibration selects the smallest value that produces an UB below a user-defined threshold.

Table 7: Optional parameters for AOH bias calibration.

| Parameter | Default | Description |
|---------------------------|---------|--|
| ScanMin | 0 | Low end of AOH bias scan range |
| ScanMax | 50 | High end of AOH bias scan range |
| ScanStepSize | 1 | Step size for AOH bias scan |
| TargetBMin | 412 | Allowed range for the coarse baseline adjustment at the end of this calibration |
| TargetBMax | 612 | |
| SaturationPointOffset | 4 | After finding the saturation point, this number is added to get the AOH bias setting (but before adjusting the black level). |
| MaxFEDReceiverInputOffset | 8 | Largest allowed value of the FED receiver input offset, which can range from 0 to 15 |
| SetAnalogInputBias | 200 | TBM settings to use |
| SetAnalogOutputBias | 120 | for all channels – |
| SetAnalogOutputGain | 200 | all 3 should be set high |
| printFEDRawData | no | Whether to print decoded transparent buffer |
| printFEDOffsetAdjustments | no | Whether to print when the FED offsets change |
| printAOHBiasAdjustments | no | Whether to print a message when the AOHBias is changed during the baseline adjustment |

13.7.2 AOH gain calibration steps

The calibration consists of the following steps:

1. Issue `ClrCal` to all ROCs, and disable hits with the control register on all ROCs, to ensure that no hits are output.
2. Issue `LRES` and `CLRES` commands to all FEDs to clear the transparent data. This ensures that no stale data is sitting in the buffer.
3. Set all TBMs to high values. The user may specify these values; otherwise defaults will be used.
4. Scan over values of AOH gain. Attempt to decode the transparent buffer on each trigger, assuming the correct number of ROCs and no hits. That is, find the TBM header, and verify that the TBM trailer is in the right place. If decoding is successful, record the UB values in the TBM header and trailer (5 values per trigger, 3 from the header and 2 from the trailer). If decoding is unsuccessful, ignore that trigger, and don't record anything. For each channel, a plot of this scan is output to a ROOT file.
5. Analyze the recorded scan data on each FED channel, looking for the smallest AOH gain where the UB is below the threshold specified by the user. This is the recommended AOH gain value. (If all values were above the threshold, set the AOH gain to 3 and generate an error.)

Table 8: Optional parameters for AOH gain calibration.

| Parameter | Default | Description |
|---------------------|---------|--|
| SetAnalogInputBias | 160 | Value of AnalogInputBias to set on all channels. |
| SetAnalogOutputBias | 110 | Value of AnalogOutputBias to set on all channels. |
| SetAnalogOutputGain | 240 | Value of AnalogOutputGain to set on all channels. |
| MaxTBMUB | 100 | TBM UB threshold. AOH gain will be set so that TBM UB is below this level. |
| printFEDRawData | no | Whether to print decoded transparent buffer |
| printScan | no | Whether to print TBM UB levels for each AOH gain |

6. On dual TBMs, check whether we can reduce the difference in UB levels by raising one of the two AOH gains by 1. If so, do it. This minimizes the spread between ultrablack levels on dual TBMs. (Since the two channels share the same TBM DAC settings, there is no other way to independently adjust the TBM UB levels.)
7. Write out new port card configuration files with the new AOH gain settings.
8. Print out a summary.

13.7.3 Parameters

A number of parameters in `calib.dat` may be used to control the calibration.

Only two “standard” parameters are used. The “`Repeat:`” parameter determines the number of triggers in each step of the scan. The channel list is used to determine which channels are calibrated. Note that the rows, columns, and DAC scan settings in `calib.dat` are completely ignored.

Many other optional parameters may be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 8.

13.8 TBM UB calibration

13.8.1 Introduction and discussion

With the black level set at 512 by the baseline calibration and automatic baseline correction, the next step is to set the ultrablack levels appropriately. We first adjust DACs on the TBM to set the TBM header and trailer ultrablack to an appropriate value. (Experience from forward testing suggests that a value of 120 to 150 is good.)

There are three DAC settings on the TBM, all of which affect the ultrablack level. Higher values of these DACs correspond to lower ultrablack (and greater B-UB separation). Figure 16 shows how these DACs affect the output of the TBM. The `AnalogOutputGain` setting

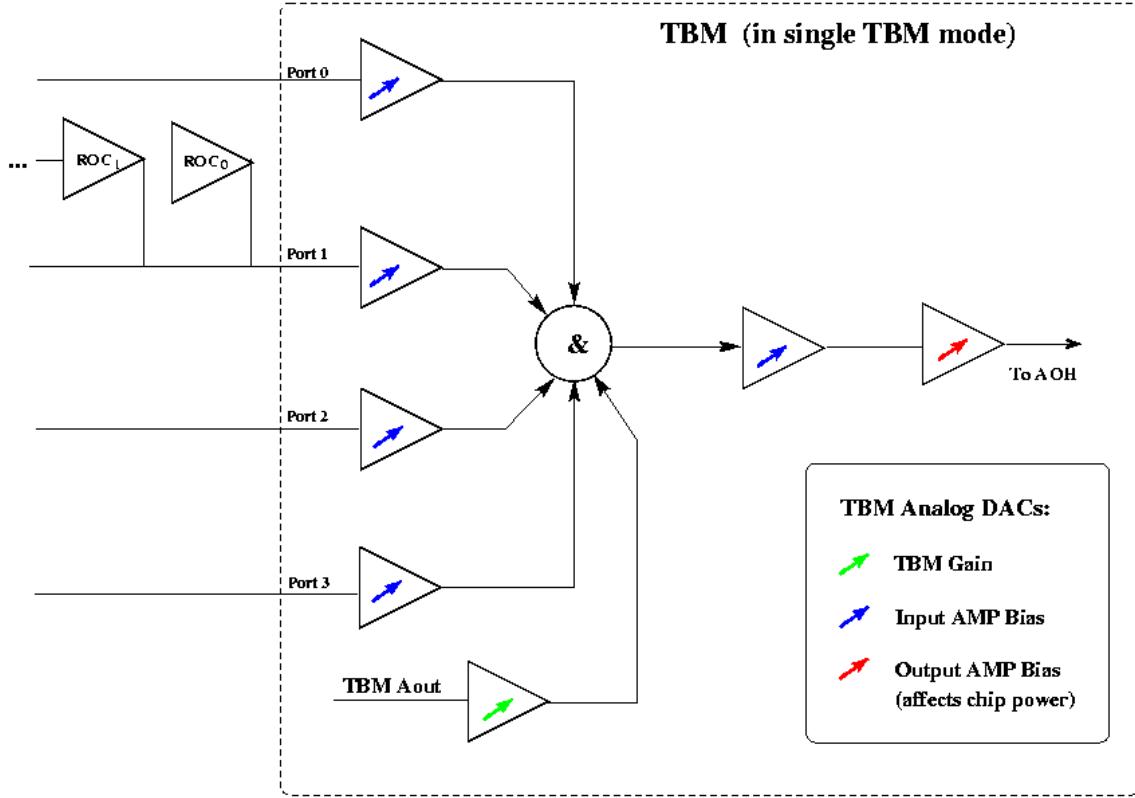


Figure 16: Diagram illustrating how the TBM DACs affect the output of the TBM. AnalogInputBias is referred to as Input AMP Bias (blue arrows), AnalogOutputBias is referred to as Output AMP Bias (red arrows), and AnalogOutputGain is referred to as TBM Gain (green arrows).

alters the levels of the TBM header and trailer, but not the signals from the ROCs. The `AnalogInputBias` setting affects both ROC and TBM signals, but it changes ROC signals more than TBM signals. The `AnalogOutputBias` setting affects ROC and TBM signals equally. The TBM UB level may be adjusted by any or all of these three DAC settings.

Although higher DAC values generally produce lower ultrablack levels, at very high DAC values the ultrablack level may actually increase. In particular, `AnalogOutputBias` should not exceed ~ 128 . If the calibration finds multiple settings that give the target ultrablack level, it will choose the lower settings.

The TBM UB calibration routine can run in either of two modes, as selected by the user:

1. Fix two DAC settings and scan the third. The user selects which DAC to scan and specifies the scan range. Optionally, the other two DACs may be given particular settings for all ROCs; otherwise the values stored in the configuration database will be

used.

2. Scan all three settings simultaneously. The user selects a scan range for each DAC, as well as the number of scan points. During the scan, all three settings are moved simultaneously and proportionally through their ranges. Specifically, for scan point $i = \{0, 1, \dots, N - 1\}$, the value of DAC d is

$$v_d(i) = min_d + \frac{i}{N - 1} \times (max_d - min_d) \quad (1)$$

where N is the number of scan points and min_d and max_d are the minimum and maximum scan values for DAC d .

In either mode, the calibration searches for DAC settings that put the TBM UB at the target level. (This level defaults to 135, and may be changed by the user.) The second mode, where all three settings are varied simultaneously, is recommended.

Dual TBMs are a special case. The two channels on a dual TBM share the same DAC settings, so they cannot be adjusted independently. Therefore the TBM settings cannot be adjusted to simultaneously put both channels' UB at the target level. For dual TBMs, we set the DACs so that one channel is at the target UB level, and the other is below.

13.8.2 TBM UB calibration steps

The calibration consists of the following steps:

1. Issue **ClrCal** to all ROCs, and disable hits with the control register on all ROCs, to ensure that no hits are output.
2. Issue **LRES** and **CLRES** commands to all FEDs to clear the transparent data. This ensures that no stale data is sitting in the buffer.
3. Scan over DAC values as specified by the user, either varying just one DAC or varying all three simultaneously. Attempt to decode the transparent buffer on each trigger, assuming the correct number of ROCs and no hits. That is, find the TBM header, and verify that the TBM trailer is in the right place. If decoding is successful, record the UB values in the TBM header and trailer (5 values per trigger, 3 from the header and 2 from the trailer). If decoding is unsuccessful, ignore that trigger, and don't record anything. For each FED channel, a plot of this scan is output to a ROOT file.
4. Analyze the recorded scan data on each FED channel, looking for the point where the measured UB crosses the target level. When the crossing is found, interpolate between the nearest scan points for the best DAC settings. If no crossing is detected, the channel is considered failed. If multiple crossings are detected, a warning is generated for that channel. New DAC settings will still be generated if the scan contains some data – either the value that puts the UB closest to the target, or the leftmost crossing point (lowest DAC values) if there are more than one.

5. On single TBMs, the new DAC values are the values at this crossing point. On dual TBMs, the new DAC values are the higher of the two sets of values found on the two FED channels for that TBM; this ensures that both UB levels will be at or below the target. If only one of the two channels was on the list of channels to calibrate, then the new DAC value is based on the one channel, and a warning is generated. A warning is also generated if both channels were successful, but the DAC values were too far apart.
6. Write out new TBM configuration files for the modules with new DAC settings (including any settings which were not varied during the scan).
7. Print out a summary.

13.8.3 Parameters

A number of parameters in `calib.dat` may be used to control the calibration.

Only two “standard” parameters are used. The “Repeat:” parameter determines the number of triggers in each step of the scan. The channel list is used to determine which channels are calibrated. Note that the rows, columns, and DAC scan settings in `calib.dat` are completely ignored.

One other parameter is required. The parameter `DACToScan` must be set to `AnalogInputBias`, `AnalogOutputBias`, `AnalogOutputGain`, or `All`. Selecting `All` will scan all three DACs simultaneously; otherwise just one is scanned.

Many other optional parameters may be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 9.

13.9 ROC UB equalization calibration

13.9.1 Introduction and discussion

This calibration sets the ultrablack level for each ROC equal to the corresponding TBM’s ultrablack level.

There are two DAC settings on the ROC which affect the ultrablack level. Higher values of these DACs correspond to lower ultrablack (and greater B-UB separation). `VIbias_rock` affects UB, address levels, and pulse height. `VIbias_DAC` affects UB and address levels, but not pulse height. In this routine, `VIbias_DAC` is scanned while `VIbias_rock` is fixed.

13.9.2 ROC UB equalization calibration steps

This calibration is very similar to the TBM UB calibration, except that each individual ROC is analyzed. It consists of the following steps:

1. Issue `LRES` and `CLRES` commands to all FEDs to clear the transparent data. This ensures that no stale data is sitting in the buffer.

Table 9: Parameters for TBM UB calibration. All except DACToScan are optional.

| Parameter | Default | Description |
|------------------------|---------|--|
| DACToScan | N/A | Which of the three TBM DACs to scan, or “All” |
| | | DACToScan ≠ All |
| ScanMin | 0 | Low end of DAC scan range |
| ScanMax | 255 | High end of DAC scan range |
| ScanStepSize | 5 | Step size for DAC scan |
| SetAnalogInputBias | [empty] | If not set, leave AnalogInputBias unchanged. Otherwise, set it to the specified value and change in the output configuration file. |
| SetAnalogOutputBias | [empty] | If not set, leave AnalogOutputBias unchanged. Otherwise, set it to the specified value and change in the output configuration file. |
| SetAnalogOutputGain | [empty] | If not set, leave AnalogOutputGain unchanged. Otherwise, set it to the specified value and change in the output configuration file. |
| | | DACToScan = All |
| MinAnalogInputBias | 0 | Minimum AnalogInputBias value during scan |
| MaxAnalogInputBias | 255 | Maximum AnalogInputBias value during scan |
| MinAnalogOutputBias | 0 | Minimum AnalogOutputBias value during scan |
| MaxAnalogOutputBias | 128 | Maximum AnalogOutputBias value during scan |
| MinAnalogOutputGain | 0 | Minimum AnalogOutputGain value during scan |
| MaxAnalogOutputGain | 255 | Maximum AnalogOutputGain value during scan |
| | | Any DACToScan Setting |
| ScanNumSteps | 52 | Number of scan steps – will override ScanStepSize setting, if present |
| TargetUB | 135 | Desired TBM UB level |
| DualTBMMaxScanStepDiff | 6 | If the difference in scan steps between the final DAC settings on the two channels of a dual TBM exceeds this value, a warning is generated. |
| printFEDRawData | no | Whether to print decoded transparent buffer |
| printScan | no | Whether to print TBM UB levels for each DAC setting |

2. Scan over values of `VIbias_DAC`. (If the `calib.dat` file specified a fixed value of `VIbias_rock`, this is also set.) Each time after `VIbias_DAC` is changed, issue `ClrCal` to all ROCs, and disable hits with the control register on all ROCs, to ensure that no hits are output (so any hits specified in the configuration file are ignored). Attempt to decode the transparent buffer on each trigger, assuming the correct number of ROCs and no hits. That is, find the TBM header, and verify that the TBM trailer is in the right place. If decoding is successful, record the UB values in the TBM header and trailer, and for each ROC. If decoding is unsuccessful, ignore that trigger, and don't record anything. For each ROC, a plot of this scan is written to a ROOT file.
3. Analyze the recorded scan data on each ROC, looking for the point where the measured ROC UB equals the UB level on the corresponding TBM. When the ROC UB crosses the TBM UB, interpolate between the nearest scan points for the best `VIbias_DAC` setting. If no crossing is detected, the ROC is considered failed. If multiple crossings are detected, a warning is generated for that ROC. A new `VIbias_DAC` setting will still be generated if the scan contains some data – either the value that puts the ROC UB closest to the TBM UB, or the leftmost crossing point (lowest `VIbias_DAC` value) if there are more than one.
4. Write out new ROC configuration files with the new `VIbias_DAC` settings. If `VIbias_rock` was modified by the user, the output files will contain this new value, too.
5. Print out a summary.

13.9.3 Parameters

A number of parameters in `calib.dat` may be used to control the calibration.

The various “standard” parameters are used in this calibration. The “`Repeat:`” parameter determines the number of triggers in each step of the scan. The ROC list is used to determine which ROCs are calibrated.

A line like

`Scan: VIbias_DAC 80 200 5`

must be included to specify the scan range (first two numbers) and step size (last number).

Optionally, `VIbias_rock` may be set with a line like

`Set: VIbias_rock 150`

Any hits specified in `Rows:` and `Columns:` will be ignored. (This is done by explicitly clearing them.) The `Rows:` and `Columns:` parameters should be empty; otherwise the calibration will work but will print a warning message.

Some optional parameters may also be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 10.

Table 10: Optional parameters for ROC UB equalization calibration.

| Parameter | Default | Description |
|-----------------|---------|--|
| printFEDRawData | no | Whether to print decoded transparent buffer |
| printScan | no | Whether to print TBM UB levels for each DAC setting |

13.10 Address level determination

The address level determination determines the values used by the FED to encode the transparent data. This include address levels for the pixel addresses, TBM header and trailer levels, and the black and ultrablack levels.

Pixels are scanned to make sure that we probe combinations of address levels that could potentially cause problems, such as transitions from high to low levels and vice versa.

Due to limitations in the size of the FIFO1 transparent data size of 512 clocks, we cannot even take one hit per ROC unless the timing is adjusted such that the start of the pulse train comes relatively early, i.e. in the first 20 or so clock cycles. Recall, each ROC, with one hit takes 9 clock cycles to read out. Then with 24 ROCs in a module we need 216 clock cycles just to read out the hits. Plus about 16 clock cycles to read out the TBM headers.

13.10.1 Data volume and time estimate

The worst case scenario for an address level calibration is one BPIX crate. We have 16 FED cards in one crate, and we can assume that all of them are fully populated. This means that we have $16 \times 36 = 576$ links. As we read out this data in transparent mode we will read 1024 words per event (even though the transparent data is only 512 bytes long). This means a total of 235926 bytes. If we want to go through all pixel we have to read out 4160 times this data volume for a total of 9,4 GB. At a speed of 10 MB/s this will take 940 s.

Assuming that we change the transparent data FIFO to be 1024 bytes. Then we can easily fit 4 hits into each ROC. This then cuts the readout time by a factor of 4 as we will need exactly 1/4 the number of triggers. This then becomes 235 seconds.

We can also use different patterns that use fewer pixel hits. But one has to take care to have all relevant transitions.

13.11 Pixel alive, Scurve, and Gain calibration

The pixel alive calibration loops over pixels and injects charge for a fixed VCal setting. The data is analyzed to produce an efficiency map that displays the efficiency for each pixel on a plaquette.

The Scurve and gain calibrations are variations of the pixel alive test where we in addition loop over the VCal setting for each pixel. For the Scurve the data is analyzed to produce an efficiency as a function of the VCal setting while for the gain calibration we look at the charge (ADC value) as function of the VCal.

To analyze the data, go to PixelAnalysisTools/test and do one of the following:

```
./bin/linux/x86/PixelAnalysis.exe PixelAlive <runnum>
./bin/linux/x86/PixelAnalysis.exe SCurve <runnum>
./bin/linux/x86/PixelAnalysis.exe Gain <runnum>
```

Writing e.g. PixelAlive will pick up the default configuration xml file, configurations/PixelAliveAnalysis.xml. You can also specify the exact configuration file you would like to use, e.g. configuration/PixelAliveAnalysis_FPix.xml.

You can scan over multiple WBCs when taking this data and analyze the data from all the WBCs summed together or only one of them using the ChooseWBC option in the configuration xml file. To measure absolute thresholds, use the in-time WBC and the following one. To measure the in-time threshold, use the in-time WBC only.

13.11.1 Practical issues

If you get a lot of “Time Out Error” messages it could be because you don’t have the right channels enabled on the FED. Or more specifically that you have enabled channels that are not connected.

13.12 Iana vs. Vana Calibration

A calibration called “Iana” is implemented to scan Vana and measure the analog current, Iana. As the power distribution for the low voltage is done per portcard for the forward pixels, this calibration works in parallel on a ROC at the time on each of the portcards. The current changes we are measuring is of the order 50 mA as you go from Vana=0 to Vana=255. The A4603 power module has a resolution of about 7 mA. We perform this measurement when all other ROCs are configured to their default values. Hence, we measure the current changes on top of a current of several amperes. (We performed trials where we lowered Vana on all other ROCs to reduce the current, but it does not really improve the measurements.) Hence, we have to take the data and fit it in order to interpolate between the points. (It takes a long time to do these measurements as one has to wait up to 6 seconds before the current measurement is stable.)

To analyze the data we fit it to a functional form. After some experimentation we came up with the following. We divide the range of Vana values into 3 ranges, compare Fig. 17. In the high range Iana is taken to be a constant. In the low range Iana is taken to be a constant plus an exponential. In the middle range Iana depends linearly on Vana, and the function is required to be continuous. In addition the function is required to have a continuous derivative as it goes between the low and middle region. This function has 5 free parameters that are determined in the fit.

First the raw data is fit to this form. After doing this fit the offset at Vana=0 is subtracted and the data is refit and plotted. These are the fits shown in Fig. 17.

A black line is used to indicate the current for the default Vana setting (Vana=140). In red is shown where you should set Vana to get an analog current of 25 mA.

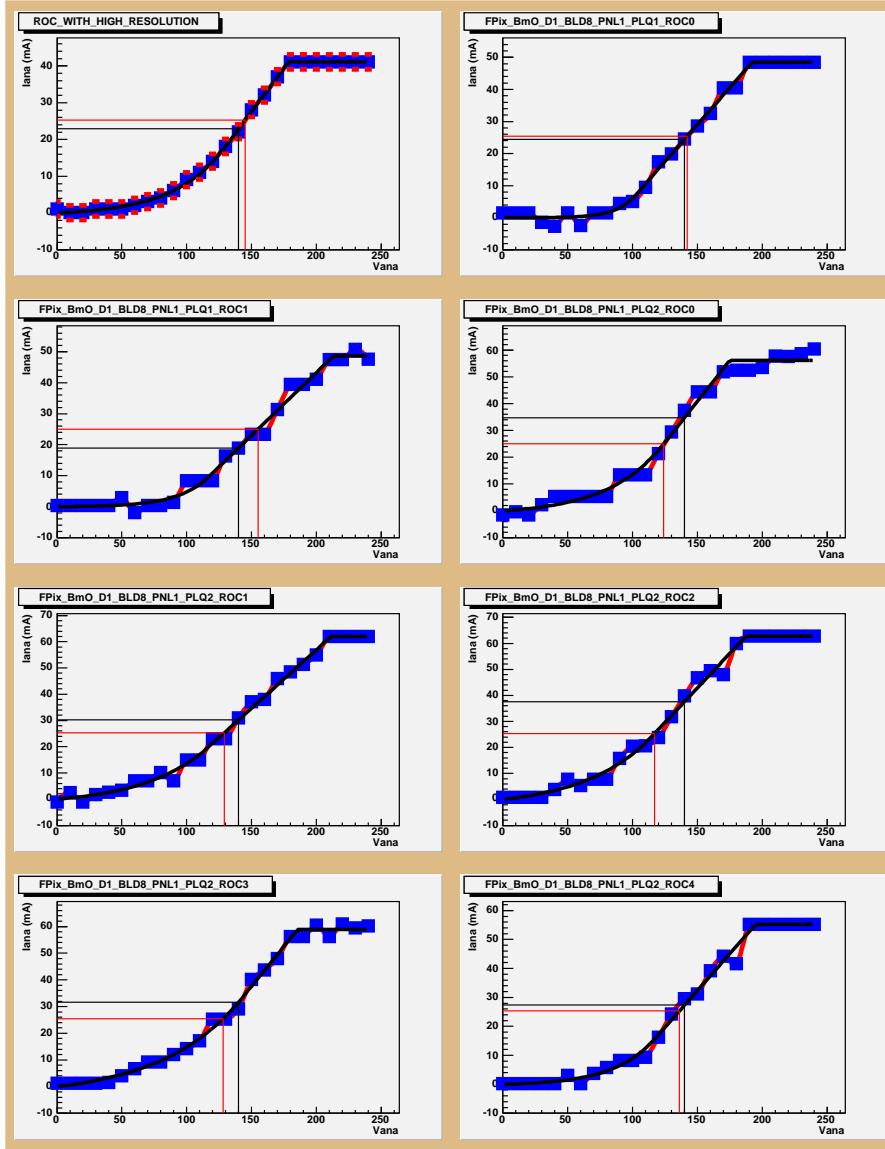


Figure 17: This figure shows the measured analog current, I_{ana} , vs the V_{ana} setting for a few ROCs. The top left plot shows the I_{ana} measured for one ROC with the A1715 power supply that has a better resolution. The black lines indicate the default setting ($V_{ana}=140$) used in the configuration when this data was taken. The red lines indicate the setting that gives a analog current of 25 mA.

13.13 Settings of CalDelay and VcThreshold

These settings are ROC wide and a few pixels are selected and a scan over the CalDelay and VcThreshold setting is performed. For each setting, triggers are taken and data is looked at using FIFO3 to see how many hits we had on each ROC. In order for this algorithm to work the address levels for black and ultra-black has to be set. (There is also an implementation that uses the FIFO1, it does not need address levels.)

We point out that CalDel is *only* relevant for calibration data taken with charge injection. The CalDel setting controls the time in which the charge is injected into the pixels. For real data we will have to adjust the timing, e.g. using the delay 25 settings from the trigger. However, for data taken with charge injection this delay is important.

This calibration performs a 2-D scan of CalDel vs VcThr. For large VcThr, which corresponds to low thresholds, we get lots of noise and the ROC digital readout basically shuts down and we don't see any hits. For lower values of VcThr we get in the range of CalDel values that corresponds to the WBC used. For larger thresholds this curve 'bends' to lower values of CalDel. As this corresponds to a higher threshold the signal reaches the threshold later and hence we need to use a smaller CalDel, i.e. inject the signal earlier.

13.13.1 Output

The primary output of this calibration are new ROC DAC settings. The only settings that should be changed in this calibration are for CalDel and VcThr.

In addition to the DAC settings the calibration produces the output file `VcThrCalDelaySummary.txt` which lists all the ROCs that were used in the calibrations. For each of the ROCs there is a corresponding file, named like `ThrCalDelScan_FPix_Bm0_D1_BLD8_PNL2_-PLQ1_ROC0.dat`. Also, plots and a summary tree are put into a ROOT file. An example of a plot is shown in Fig. 18.

13.13.2 Example

An example configuration file is shown below.

```
Mode: ThresholdCalDelay
```

```
Rows:
```

```
10 | 20
```

```
Cols:
```

```
10 | 20
```

```
VcalHigh
```

```
Scan: VcThr 0 255 8
```

```
Scan: CalDel 0 255 8
```

```
Set: Vcal 50
```

```
Repeat: 10
```

```
ToCalibrate:
```

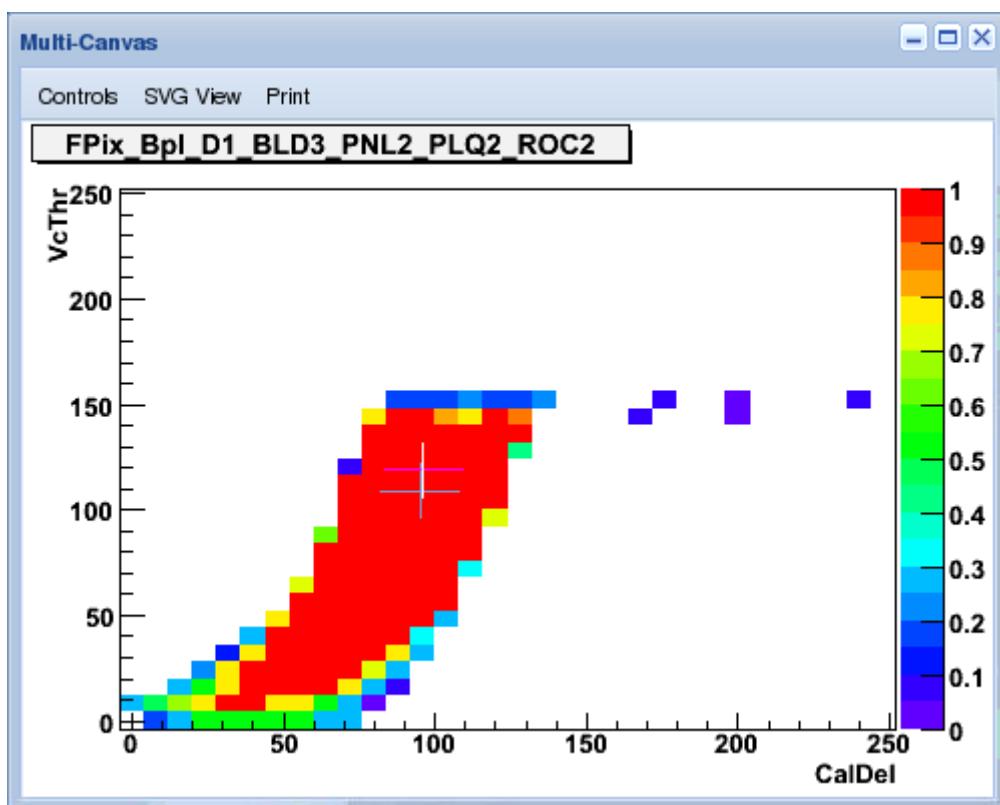


Figure 18: The efficiency for detecting a hit is shown as a function of VcThr vs. CalDel. Too large values of VcThr, corresponding to a low threshold, generates much noise that saturates the digital circuit and no hits are seen. The optimal point is indicated in black and the blue point indicates the old point from the configuration.

all

13.14 CalDel calibration

Again the note from Sect. 13.13 applies in that CalDel is not relevant for physics data.

In this calibration we scan CalDel for a series of different WBC settings. This creates a pattern as indicated in Fig. 19. As a change of one unit of WBC corresponds to 24.95ns (40.079 MHz) we can use the change in CalDel for different WBC settings to calibrate absolutely what a change of CalDel corresponds to in absolute time.

We run this calibration with an injection of a large signal (255 on the high Vcal range). We then pick a Vcal setting that is near the *early* start of the efficient range for the WBC setting. We pick an early time so that we retain efficiency for signals that are not as strong and takes time to get over threshold. How close to the edge the CalDel setting is taken can be specified by the parameter “Fraction”. Setting this to 0 means that you pick the point right at the early edge, setting the parameter to 1 means that you select the late edge.

13.14.1 Output

The CalDel calibration produces new ROC DAC settings that update the CalDel settings, all other settings should be left unchanged. The calibration also generates output to display the region of good efficiency in the WBC vs CalDel. The output can be processed to generate plots using the root script `caldel.c`. An example plot is shown in Fig. 19.

13.14.2 Example

An example of a configuration file for the CalDel calibration is shown below:

```
Mode: CalDelCalibration
Rows:
10 | 20
Cols:
10 | 20
VcalHigh
Scan: WBC 117 123 1
Scan: CalDel 0 255 1
Set: Vcal 250
Repeat: 2
ToCalibrate:
+ all
```

If you want to just scan over WBC to find the right WBC value you can use a configuration that sets CalDel to a fixed value:

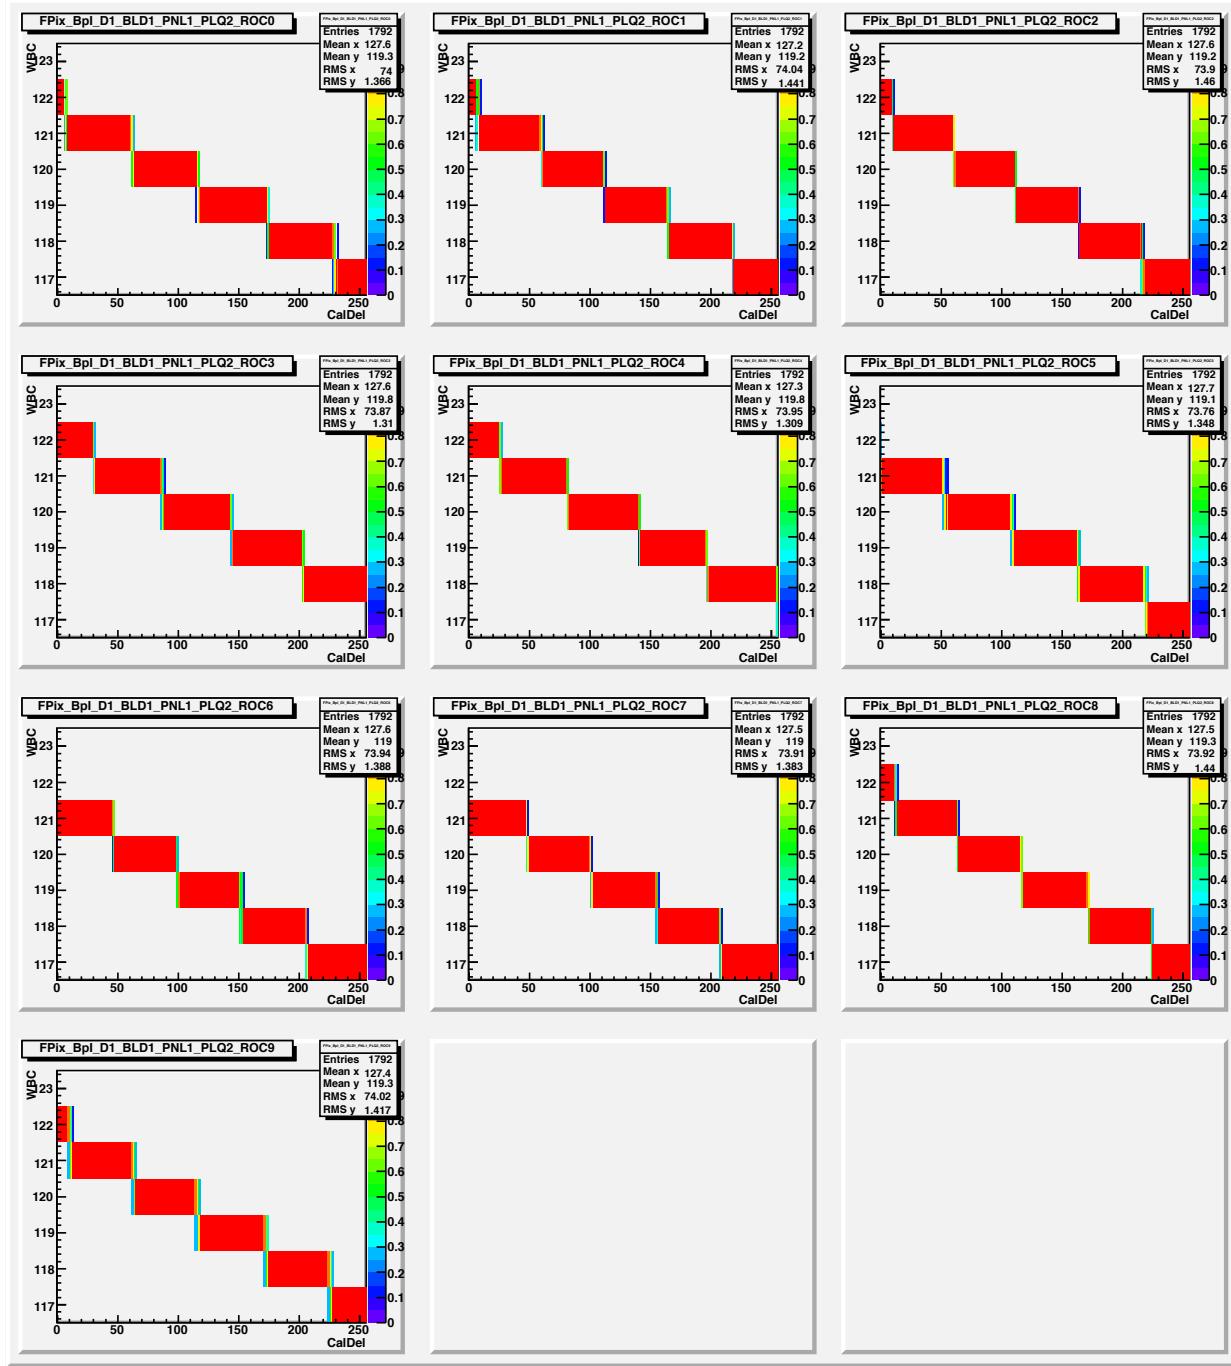


Figure 19: The efficiency as a function of WBC vs CalDel.

```

Mode: CalDelCalibration
Rows:
10 | 20
Cols:
10 | 20
VcalHigh
Scan: WBC 117 123 1
Scan: CalDel 100 100 1
Set: Vcal 250
Repeat: 2
ToCalibrate:
+ all

```

13.15 Idigi vs. Vsf

In this 'calibration' we scan Vsf and measure the digital current. This is not really a calibration as we don't determine any settings from this. Rather what we do is to find a maximum Vsf setting that we will use. If Vsf is too large the digital current increases. In the configuration for this calibration you can set the maximum increase in Idigi that you allow using the parameter IdigiMax (units are in mA). IdigiMax is currently set to 7 mA.

13.15.1 Output

This calibration produces the object PixelMaxVsf which is used in the calibration that determines Vsf to optimize the linearity. In addition this calibration produces output in a file called idigi.dat. Using the root script idigi.c you can generate plots as shown in Fig. 20.

13.15.2 Example

The configuration used for the digital scan looks like

```

Mode: Idigi
Rows:
Cols:
Vcal:
200 200 5
Repeat:
1
ToCalibrate:
+ all

```

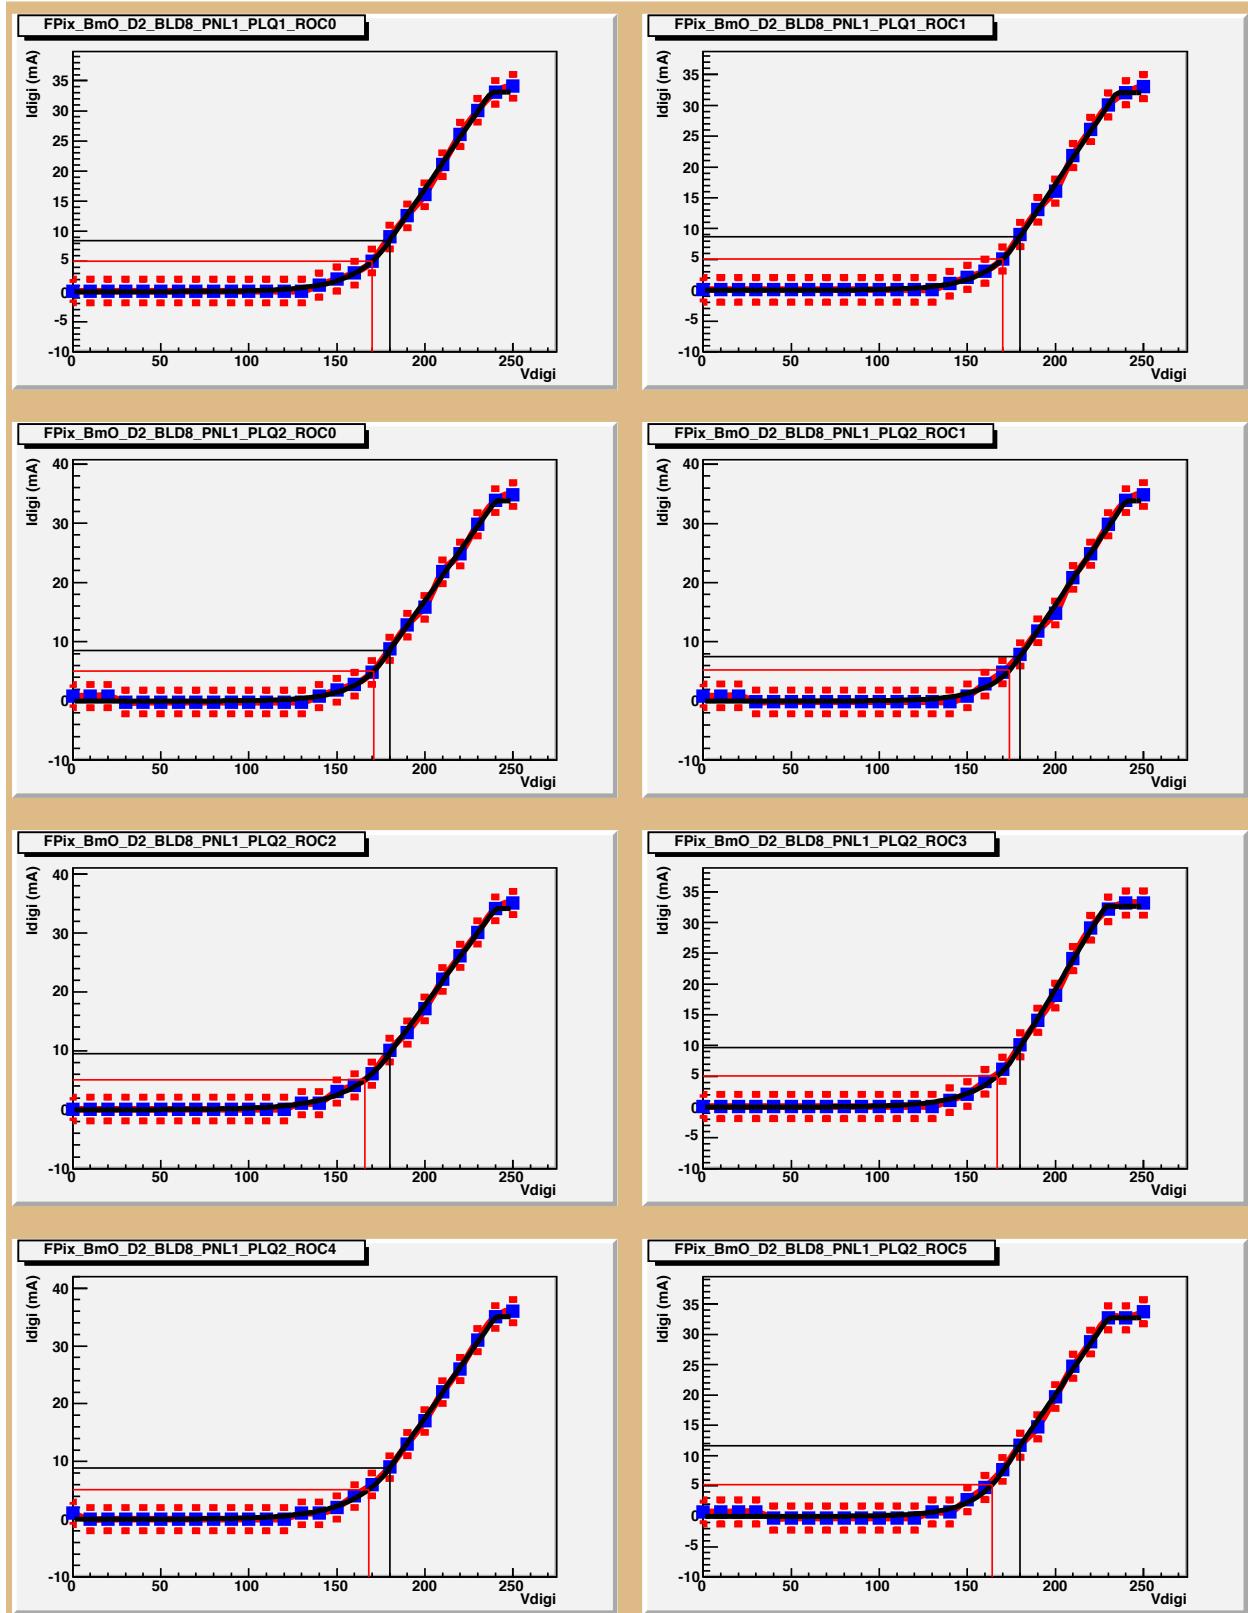


Figure 20: The digital current⁸⁶ is shown as a function of Vsf.

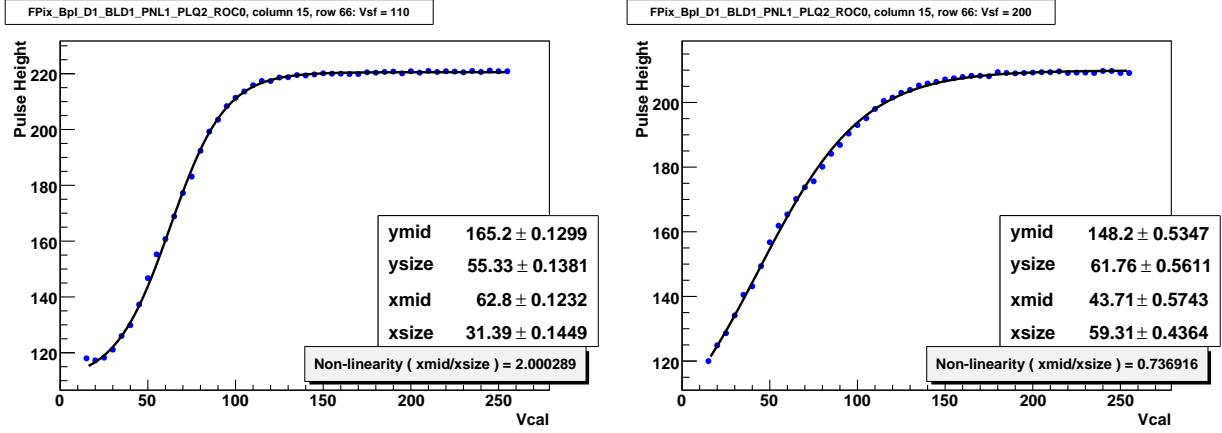


Figure 21: Scans of pulse height vs. Vcal at different values of Vsf. The scan on the left has poor linearity, and the scan on the right has good linearity. (Note that these scans come from a detector which has no voltage bias, so the received charge and the slope of the linear section are lower than in a biased detector.)

13.16 Linearity vs. Vsf

The ROC DAC setting V_{sf} affects the linearity of the pixel response vs. received charge; larger values improve linearity. V_{sf} also affects the digital current, with higher values increasing the current. We have implemented two algorithms to set V_{sf} at a value that gives good linearity without drawing excessive power. The first, described in this section, is the algorithm developed at PSI.

13.16.1 Introduction and discussion

This algorithm measures linearity from scans of pulse height vs. V_{cal} at different values of V_{sf} . Examples of these scans are shown in Fig. 21.

To quantify the degree of nonlinearity, the scan data are fit with a function

$$PH = f(V_{cal}) = y_{\text{mid}} + y_{\text{size}} \times \tanh\left(\frac{V_{cal} - x_{\text{mid}}}{x_{\text{size}}}\right) \quad (2)$$

where PH is the recorded pulse height, V_{cal} is the input, $(x_{\text{mid}}, y_{\text{mid}})$ is the point at the center of the quasi-linear rise region of the hyperbolic tangent, x_{size} is the horizontal scale of the quasi-linear region, and y_{size} is the vertical scale of that region.

From this fit, the degree of nonlinearity can be quantified in different ways. The simpler nonlinearity parameter, used by PSI and used by default in this calibration, is $x_{\text{mid}}/x_{\text{size}}$. When this parameter is small, it means that $V_{cal} = 0$ lies within the quasi-linear rise region, and hence the response is linear for small amounts of injected charge. PSI has used $x_{\text{mid}}/x_{\text{size}} = 1.4$ as the cutoff – values below 1.4 indicate good linearity.

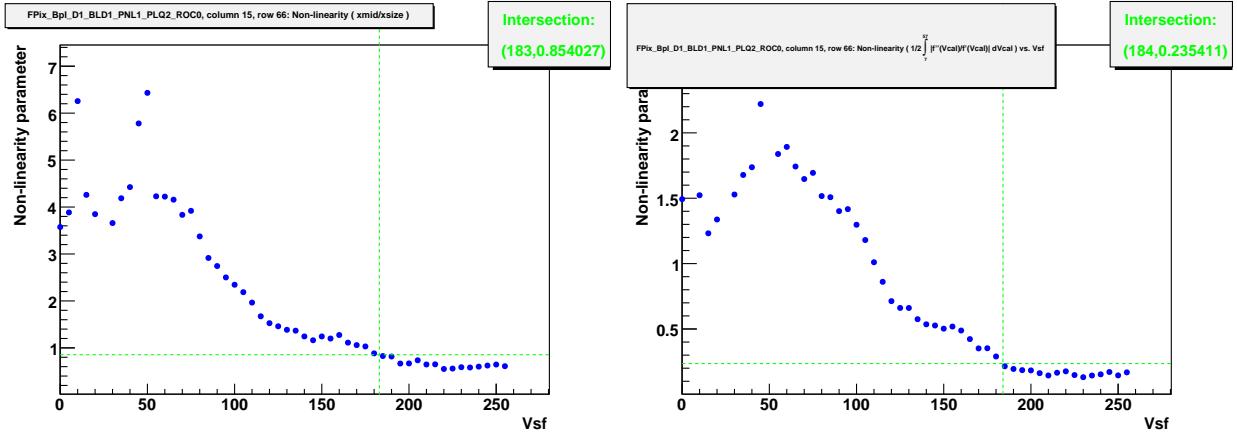


Figure 22: Plots of nonlinearity vs. Vsf . On the left, nonlinearity is measured by $x_{\text{mid}}/x_{\text{size}}$, and on the right, it is measured by the integral in Eq. (3).

An alternate nonlinearity parameter is

$$\frac{1}{2} \int_{Vcal_{\min}}^{Vcal_{\max}} dVcal \left| \frac{f''(Vcal)}{f'(Vcal)} \right| \quad (3)$$

where $f(Vcal)$ is the hyperbolic tangent function in Eq. (2). This integral is a measure of the vertical change due to curvature divided by the vertical change due to slope over the range of the integral. The limits of the integral should be chosen to include the range for which we want good linearity. By default the limits go from $Vcal = 50$ to $Vcal = 400$ in low-scale $Vcal$ units, or from $50/7$ to $400/7$ in high-scale units. This integral can be evaluated analytically.

Figure 22 shows scans of both measures of nonlinearity vs. Vsf . As seen in the plots, they give similar shapes. This calibration allows the user to use either measure; the default is $x_{\text{mid}}/x_{\text{size}}$.

13.16.2 Linearity vs. Vsf calibration steps

This calibration consists of the following steps:

1. Scan $Vcal$ and Vsf according to the calibration's configuration. On each trigger, read out data from each FED's FIFO3. Record the pulse height, keeping separate scans for different pixels.
2. For each scan of pulse height vs. $Vcal$ (for a given ROC, pixel, and Vsf value), fit with the function in Eq. (2). A plot of each scan and fit is written to a ROOT file (see Fig. 21). If the fit is successful, compute the nonlinearity parameter (either $x_{\text{mid}}/x_{\text{size}}$ or the integral in Eq. (3)) and add it to a scan of nonlinearity vs. Vsf for that ROC and pixel. This scan is also written to a ROOT file (see Fig. 22).

3. For each scan of nonlinearity vs. `Vsf`, determine an optimal `Vsf`. This is done by finding the highest `Vsf` where the scan intersects a nonlinearity threshold. This threshold may be either absolute or relative, according to the parameter “`absoluteNonlinearityThreshold`”. If absolute, the threshold is set by the parameter “`nonlinearityThreshold`”. If relative, the threshold is the value of the nonlinearity at the highest `Vsf` in the scan, multiplied by the parameter “`nonlinearityThreshold`”.
4. For each ROC, examine the optimal `Vsfs` on the various pixels. If there are at least 4 pixels, discard any `Vsf` outliers. After any discarding, average the `Vsf` values to determine the `Vsf` setting for that ROC.
5. Compare the `Vsf` setting on each ROC to the `maxVsf` setting in the configuration. If the setting from this calibration exceeds the maximum, reduce `Vsf` to the maximum allowed.
6. Write out new ROC configuration files for the successfully-calibrated ROCs with the new `Vsf` settings.
7. Print out a summary.

13.16.3 Parameters

A number of parameters in `calib.dat` may be used to control the calibration.

The various “standard” parameters are used in this calibration. The “`Repeat:`” parameter determines the number of triggers in each step of the scan. The ROC list is used to determine which ROCs are calibrated.

The scan ranges are specified by lines like

```
Scan: Vcal 0 255 5
Scan: Vsf   0 255 5 mix
```

which specify the scan range and step size. The “`mix`” flag must be enabled for `Vsf` (or alternatively SingleROC mode must be used); otherwise the calibration will abort.

The hits specified in `Rows:` and `Columns:` will be used. One scan is taken for each pixel with hits. At most two hits per pixel pattern may be specified (or alternatively SingleROC mode must be used); otherwise the calibration will abort. The reason for this restriction is that the FED’s spy FIFO3 will overflow if each ROC connected to it has more than two hits.

Some optional parameters may also be set. All have default values which will be used if the parameter is not set. These parameters, their defaults, and their functionality are given in Table 11.

Table 11: Optional parameters for linearity vs. V_{sf} calibration.

| Parameter | Default | Description |
|-----------------------------------|---------|---|
| absoluteNonlinearityThreshold | yes | Whether the nonlinearity threshold is an absolute number, or a multiple of the nonlinearity parameter at maximum V_{sf} |
| nonlinearityThreshold | 1.4 | Value of the nonlinearity threshold |
| useIntegrated2ndOver1stDerivative | no | Whether to use the integral in Eq. (3) for the nonlinearity parameter, instead of $x_{\text{mid}}/x_{\text{size}}$ |
| integralMinVcal | 50/7 | Lower V_{cal} bound of the integral (if used) |
| integralMaxVcal | 400/7 | Upper V_{cal} bound of the integral (if used) |

13.17 V_{sf} and $VHldDel$

13.17.1 Introduction and discussion

We have also implemented the Renaissance algorithm for determining the ROC DAC settings V_{sf} and $VHldDel$. Figure 23 shows plots of pulse height vs. $VHldDel$ at low, medium, and high values of V_{sf} , with $V_{cal} = 250$ on the low scale. A good V_{sf} value is one for which this curve rises and then falls so that the pulse heights at the two endpoints (lowest and highest $VHldDel$) are equal. Figure 23 also includes a plot of these endpoints as a function of V_{sf} ; the rightmost intersection point is the V_{sf} value chosen. Low values of V_{sf} , below ~ 90 , produce garbage output.

After choosing a V_{sf} value, $VHldDel$ should be set to the value that maximizes the pulse height.

When charge is injected with $V_{cal} = 250$ on the low scale, these V_{sf} and $VHldDel$ settings are found to give good linearity in pulse height vs. injected charge, while not drawing too much power.

This procedure may be done in one calibration run, or it may be split into two – the first to determine V_{sf} and the second to determine $VHldDel$. Splitting into two runs is better because (1) it requires fewer scan points and (2) when the best V_{sf} is a value interpolated between two scan points, the scan over $VHldDel$ will take place at that value instead of at the nearest V_{sf} scan point. The same calibration code can calibrate both DACs at once, or either one of them; calibration parameters determine which mode is used. This is described in more detail below.

If the PSI algorithm described in Sec. 13.16 is used to determine V_{sf} , $VHldDel$ should be set with the algorithm used here – maximizing pulse height at the chosen V_{sf} .

13.17.2 V_{sf} and $VHldDel$ calibration steps

This calibration consists of the following steps:

1. Scan V_{sf} and $VHldDel$ according to the calibration’s configuration. On each trigger, read out data from each FED’s FIFO3. Record the pulse height. If multiple pixels are

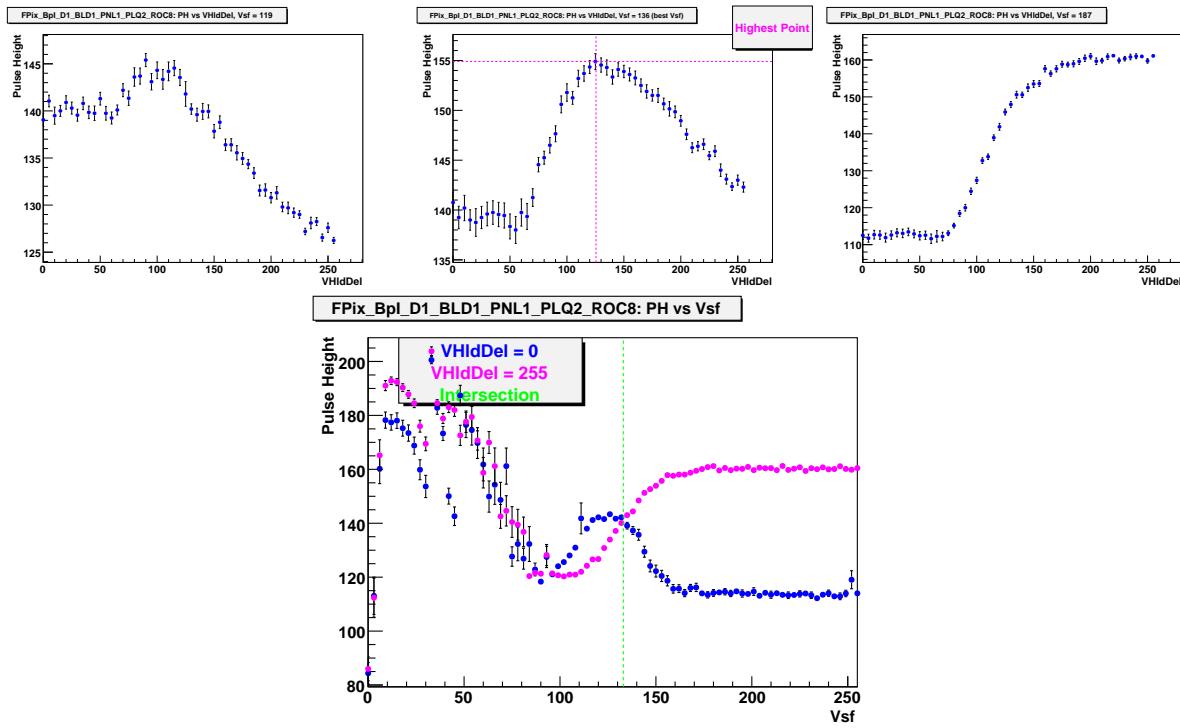


Figure 23: *Top row*: Pulse height vs. VHldDel at low, medium, and high values of Vsf, with Vcal = 250 on the low scale. As Vsf increases, the right endpoint increases. The best Vsf value is the one for which the pulse heights measured at VHldDel = 0 and VHldDel = 255 are equal. *Bottom plot*: Pulse height at the endpoints of these plots, as a function of Vsf. Low values of Vsf produce garbage output.

enabled, average all of their pulse heights on each ROC.

2. If **Vsf** was scanned, determine the best **Vsf**. To do this, make plots of pulse height vs. **Vsf** for the lowest **VH1dDel** scan value and for the highest **VH1dDel** scan value. Take the rightmost intersection (the one at the largest **Vsf**) as the best **Vsf** value. (When finding this intersection, interpolate between scan points.) If the intersection point exceeds the **maxVsf** setting in the configuration for that ROC, reduce **Vsf** to the maximum allowed.
3. If there are 3 or more **VH1dDel** scan points, determine the best **VH1dDel**. If **Vsf** was scanned, examine the scan of pulse height vs. **VH1dDel** for the **Vsf** scan point closest to the value chosen. If **Vsf** was not scanned, a scan of pulse height vs. **VH1dDel** was taken at the **Vsf** value stored in the configuration; examine this scan. Find the scan point with the highest pulse height. Calculate the parabola that passes through this point and the two adjacent points, and take the maximum of this parabola as the best **VH1dDel** value. (If the highest point is an endpoint, use the endpoint as the best **VH1dDel**.)
4. Write out new ROC configuration files for the successfully-calibrated ROCs with the new **Vsf** and/or **VH1dDel** settings.
5. Print out a summary.

13.17.3 Parameters

The standard parameters in `calib.dat` may be used to control the calibration.

The “**Repeat:**” parameter determines the number of triggers in each step of the scan. The ROC list is used to determine which ROCs are calibrated.

The hits specified in **Rows:** and **Columns:** will be used. At most two hits per pixel pattern may be specified (or alternatively SingleROC mode must be used); otherwise the calibration will abort. The reason for this restriction is that the FED’s spy FIFO3 will overflow if each ROC connected to it has more than two hits.

The scan ranges determine whether the calibration will determine only **Vsf**, only **VH1dDel**, or both. To determine just **Vsf**, use something like:

```
VcalLow
Scan: Vsf      0 255 3    mix
Scan: VH1dDel 0 255 255
Set:   Vcal 250
```

This sets **Vcal** = 250 on the low scale, just as Renaissance does. The **Vsf** scan is specified in the usual way. The “**mix**” flag must be enabled for **Vsf** (or alternatively SingleROC mode must be used); otherwise the calibration will abort. **VH1dDel** is also scanned, but with just two scan points – the smallest and largest values.

To determine just **VH1dDel**, use something like:

```
VcalLow  
Scan: VH1dDel 0 255 5  
Set: Vcal 250
```

Vsf is not scanned, and a full scan is specified for **VH1dDel**.

To determine both **Vsf** and **VH1dDel**, use something like:

```
VcalLow  
Scan: Vsf      0 255 3      mix  
Scan: VH1dDel 0 255 5  
Set: Vcal 250
```

Note that this scan is more time-consuming than running first for **Vsf** and then for **VH1dDel**. It also does not allow for the use of interpolated **Vsf** values when determining **VH1dDel**. Therefore, it is recommended to run two separate scans.

This calibration has no additional parameters aside from the standard ones.

13.18 Pulse height range calibration

13.18.1 Introduction and discussion

A number of ROC DAC settings affect the scaling of the pulse height signal that is sent to the FED. (This refers not to the actual charge collection, but to the translation of that charge into the signal sent out to the FED.) We want the range of this signal – the difference in recorded pulse height between small and large amounts of charge – to be large. However, the pulse height signal should not go low enough to be confused with the ultrablack level, nor high enough to exceed the FED’s dynamic range.

The ROC DAC setting **VIbias_PH** is intended to be used for adjusting the pulse height, so this DAC should be scanned. Other DACs that affect the pulse height are **VOffsetRO**, **VIon**, and **VOffsetOp**. Previous studies found that adjusting **VIbias_PH** and **VOffsetOp**, while fixing **VOffsetRO** and **VIon** worked well [6].

13.18.2 Pulse height calibration procedure

The algorithm for this calibration is rather minimal. In the configuration, the user specifies any number of arbitrary ROC DACs to be scanned, and also “low” and “high” amounts of injected charge (i.e. low and high **Vcal** settings). The pulse height is recorded at each scan point for the low and high charges. (If multiple pixels are enabled, their pulse heights are averaged on each ROC.) A scan point is discarded if either the low or high reading falls outside a user-specified range. From the remaining scan points, the point chosen is the one with the largest pulse height difference between high and low charges. These settings are written out.

If either 1 or 2 DACs are scanned, plots are generated of the high PH, low PH, and PH range (difference between high and low) as a function of the DAC setting(s).

Table 12: Optional parameters for pulse height range calibration.

| Parameter | Default | Description |
|-----------|------------|------------------------------|
| minPH | 75 = 300/4 | Minimum allowed pulse height |
| maxPH | 254 | Maximum allowed pulse height |

13.18.3 Parameters

The standard parameters in `calib.dat` may be used to control the calibration.

The “Repeat:” parameter determines the number of triggers in each step of the scan. The ROC list is used to determine which ROCs are calibrated.

The hits specified in `Rows:` and `Columns:` will be used. At most two hits per pixel pattern may be specified (or alternatively SingleROC mode must be used); otherwise the calibration will abort. The reason for this restriction is that the FED’s spy FIFO3 will overflow if each ROC connected to it has more than two hits.

The scan ranges are used to determine which DACs will be adjusted. For example, to scan over `VIBias_PH` and `VOffsetOp`, use something like:

```
VcalHigh
Scan: VIBias_PH 0 255 5
Scan: VOffsetOp 0 255 5
Scan: Vcal 25 255 230
```

Note that a `Vcal` range should also be given. The pulse height range at a given scan point will be calculated from the highest and lowest `Vcal` for which hits were found. If the final scan contains more than two `Vcal` values with hits found, the points in the middle are ignored in determining the pulse height range. Therefore, usual practice is to specify just two `Vcal` scan points. Note that if the lower `Vcal` is below threshold, it will generate no hits; therefore the user should select a minimum that is above threshold. If just one `Vcal` value produces hits on a ROC, the calibration fails on that ROC.

Two optional parameters may also be set – `minPH` and `maxPH`. These specify the minimum and maximum allowable pulse height values; DAC settings that give pulse height outside of this range are rejected. These parameters are specified on the scale of pulse height readings, which runs from 0 to 255. This is in contrast to FED ADC counts, which run from 0 to 1023. (Multiply the pulse height by 4 to get FED ADC counts. For example, a PH of 75 corresponds to 300 FED ADC counts.) These parameters default to 75 and 254; they may be changed in `calib.dat`. The parameters are summarized in Table 12.

13.19 Vana and time walk

This calibration has not been completely implemented. The idea is to run a scan of Vana vs CalDel. For a large charge we will assume that the response of different sensors is similar. As for a large charge (the maximum that can be injected) the rise time is small it is probably

a reasonable assumption that we can assume that we have a small variation between ROCs. However, for a smaller charge, like 250 on the low Vcal range, this is not obviously true. This calibration finds the time (absolute) that the Vcal=250 signal is over threshold and also finds the slope dt/dI_{ana} , or dt/dV_{ana} . Using this information the ROCs can be unified to one time at which a pulse of this strength passes over threshold.

13.20 Trim bit determination

The algorithm described below is iterative. It uses a linearized model to determine the parameters. Successive iterations of this algorithm should provide better values of the trim parameters.

The threshold (for a pixel) is a function of V_{cthr} , V_{trim} , and $Trim$.

$$\text{Threshold} = \text{Threshold}(\text{pix}_i, V_{cthr}, V_{trim}, \text{Trim}_i) \quad (4)$$

We can expand this as:

$$\begin{aligned} \text{Threshold}(\text{pix}_i, V_{cthr}, V_{trim}, \text{Trim}_i) &= \text{Threshold}(\text{pix}_i, V_{cthr_0}, V_{trim_0}, \text{Trim}_{0i}) \\ &+ \text{delta}V_{cthr} \frac{d\text{Threshold}}{dV_{cthr}} \\ &+ \text{delta}V_{trim} \frac{d\text{Threshold}}{dV_{trim}} \\ &+ \text{delta}Trim_i \frac{d\text{Threshold}}{dTrim} \end{aligned}$$

Simplifying the notation we have

$$\text{Threshold}(\text{pix}_i, V_{cthr}, V_{trim}, \text{Trim}_i) = A_i + B_i * \text{delta}V_{cthr} + C_i * \text{delta}V_{trim} + D_i * \text{delta}Trim_i \quad (5)$$

Now we want to set all of these equal to some value, ThrTrim .

$$\text{ThrTrim} = A_i + B_i * \text{delta}V_{cthr} + C_i * \text{delta}V_{trim} + D_i * \text{delta}Trim_i \quad (6)$$

So we have N equations and $N + 2$ unknowns. Generally there are many solutions. But recall that the trim bits only takes on values from 0 to 15. So we need more constraints. We use the following:

- We need to decide how many pixels we include in the trimming. I.e., setting V_{trim} large enough so that the trim bits can adjust the threshold. Based on the 'initial' parameters find the rms of the threshold distribution. Find $n - \sigma$ and using an average C_i calculate $\text{delta}V_{trim}$ needed to include pixels out to $n - \sigma$. So this gives the first parameter. (Use C_i for the smallest V_{trim} .)

- Require that the average trim is 7.5, that is

$$\sum_i \text{deltaTrim}_i = \sum_i (7.5 - \text{Trim}_i)$$

This yields a simple constraint:

$$\frac{\text{ThrTrim}}{D_i} = \frac{A_i}{D_i} + \frac{B_i}{D_i} \text{deltaVcthr} + \frac{C_i}{D_i} \text{deltaVtrim} + \text{deltaTrim}_i$$

and now summing over pixels give:

$$\sum_i \frac{\text{ThrTrim}}{D_i} = \sum_i \frac{A_i}{D_i} + \sum_i \frac{B_i}{D_i} \text{deltaVcthr} + \sum_i \frac{C_i}{D_i} \text{deltaVtrim} + \sum_i (7.5 - \text{Trim}_i)$$

as we already have determined deltaVtrim we now also get deltaVcthr.

Last all we have to do is to loop over all pixels and calculate deltaTrim_i . This algorithm can be iterated to determine a better approximation of the trim bits.

The derivatives used above have to be determined numerically by running several S-curve runs.

13.21 Trim bit analysis

This section describes the tools that we have available today for trim bit determination.

13.21.1 Running the threshold analysis

The inputs to the threshold determination is Scurve runs with different parameters. After taking one of these Scurve runs it has to be analyzed to produce a file that contains the thresholds for each pixel.

After taking an SCurve (trim) run (TrimDefaultShort, TrimVcThrShort, TrimVtrimShort, TrimOnShort, TrimOffShort, TrimDefault, TrimVcThr, TrimVtrim, TrimOn, TrimOff) and analyze the run using

```
./bin/linux/x86/PixelAnalysis.exe_SCurve_jrunnumi
```

The configuration/SCurveAnalysis.xml file needs to contain:

```
<OutputTrimFile      Write="Yes"/>
```

in order to generate the output needed in the trim analysis.

Currently this produces a very long output name. Rename this file:

```
mv $POS_OUTPUT_DIRS/Run_100000/Run_100141/TrimOutputFile_Fed_32-33-34\  
35-36-37-38-39.dat $POS_OUTPUT_DIRS/Run_100000/Run_100141/TrimDefault.dat
```

For a 'TrimDefault' run. For the other types of runs listed above they should be named:

```
TrimVcThr.dat  
TrimVtrim.dat  
TrimOn.dat  
TrimOff.dat
```

13.21.2 Quality of trimming

There is a command line tool that looks at the quality of the trimming. It can be run using:

```
[pixelpro@vmepecS2B18-17 test]$ ./bin/linux/x86/PixelTrimAnalysis.exe 100141 | more
Usage: PixelTrimAnalysis.exe runTrimDefault
trimDefault:/nfshome0/pixelpro/TriDAS_build7/pixel/PixelRun/Runs/Run_100000/\
Run_100141/TrimDefault.dat
All ROCs
FPix_Bm0_D1_BLD1_PNL1_PLQ1_ROC0 81 64.1024 2.61066
FPix_Bm0_D1_BLD1_PNL1_PLQ1_ROC1 81 69.239 1.39128
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC0 81 74.0942 1.6798
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC1 81 65.1567 1.36941
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC2 81 65.5652 1.4101
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC3 81 77.1227 2.36945
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC4 78 111.641 18.149
FPix_Bm0_D1_BLD1_PNL1_PLQ2_ROC5 81 70.3416 1.4199
FPix_Bm0_D1_BLD1_PNL1_PLQ3_ROC0 81 65.2256 1.4215
.
.
.
FPix_BpI_D2_BLD12_PNL2_PLQ3_ROC6 81 68.3048 1.82446
FPix_BpI_D2_BLD12_PNL2_PLQ3_ROC7 77 75.3988 1.54956
FPix_BpI_D2_BLD12_PNL2_PLQ3_ROC8 81 70.205 1.43413
FPix_BpI_D2_BLD12_PNL2_PLQ3_ROC9 81 59.5848 4.26486
Problem ROCs
FPix_Bm0_D2_BLD1_PNL2_PLQ1_ROC3 81 49.3475 3.26693
FPix_Bp0_D1_BLD9_PNL1_PLQ3_ROC2 11 65.2924 43.158
```

Where the criteria for a problem ROC was set in the example above as a ROC with a threshold below 50 or with less than 50 hit pixels (out of 81).

13.21.3 Determining new VcThr settings

If you just want to adjust VcThr to a fixed threshold you can run the following. First, take one 'TrimDefault(Short)' and one 'TrimVcThr(Short)'. In this example these are runs 65624 and 65626 respectively. The 'TrimVcThr' run adjusts VcThr (down) by a fixed amount for each ROC, e.g. 5 DAC settings.

To run the analysis code

```
./bin/linux/x86/PixelTrimVcThr.exe 10630 65624 65626 > log_TrimVcThr_65624
```

10630 is the configuration key for run 65624. The output is a new set of DAC settings in the current directory with new VcThr values. The target Vcal is set in `src/common/PixelTrimVcThr.cc` on the line

```
trimROC.setThrTrim(60.0);
```

This is the target threshold in Vcal units.

The produced logfile is huge. One useful thing to do is:

```
grep deltaVcthr log_TrimVcThr_65624
```

to see what the changes in VcThr are for the different ROCs.

13.21.4 Make simple plots of thresholds

To make simple plots of the thresholds do:

```
[pixelpro@vmepcS2B18-17 test]$ root -l  
root [0] .L read.c  
root [1] .x plot.c
```

plot.c specifies the input file.

13.21.5 Full trim bit determination

To do full trimming, i.e., adjusting the trim bits, Vtrim, and VcThr to trim the ROCs to a preset threshold the following steps are to be followed.

First you need to take five runs:

```
TrimDefaultShort  
TrimVcThrShort  
TrimVtrimShort  
TrimOnShort  
TrimOffShort
```

Normally, we use the ‘short’ configuration files that only use about 50 to 100 pixels on each ROC. This is enough pixels to find the values needed for VcThr and Vtrim on a given ROC. In a later step we adjust just the trim bits for the remaining pixels. The TrimDefaultShort run is a run with the ‘default’ or current settings. The procedure is iterative and the parameters produced in this iteration will be the default parameters for the next iteration. The TrimVcThrShort and TrimVtrimShort runs adjusts the VcThr and Vtrim parameter by a default amount (typically -5 and +10 for VcThr and Vtrim respectively). The last two runs, TrimOnShort and TrimOffShort, turns the trim bits on and off respectively for each of the pixels.

After taking these runs, they are analyzed as described in Sect. 13.21.1. Having done this the new DAC and trim bits are determined using

```
./bin/linux/x86/PixelTrim.exe <key> <runTrimDefault> <runTrimVcThr> \  
<runTrimVtrim> <runTrimOn> <runTrimOff>
```

Running this command produces new DAC settings and trimbits. In addition, it produces a file called ’rocder.dat’. The use of this file is described in the next section.

13.21.6 Determining all trimbits

Having carried out the procedure described in the previous section a file called 'rocder.dat' should have been produced. This file contains three numbers for each ROC determined from the analysis described above. First is the derivative of the threshold with respect to the VcThr setting. The second is the derivative of the threshold with respect to Vtrim. The last is the change in the threshold with respect to changing trim bits, averaged over the pixels on the ROC. The change in threshold with respect to the trimbits is used in the adjustment of the trim bits.

A run is taken using the 'TrimDefault' configuration. This run is analyzed as described in Sect. 13.21.1.

```
./bin/linux/x86/PixelTrimBits.exe <key> <runTrimDefault>
```

Where <key> is the global key used and <runTrimDefault> is the run number for the trim default run. This produces new trim bit settings. The quality of the trimming can be checked as described in Sect. 13.21.2.

13.22 Emulated physics

The FED allows us to program emulated analog data as inputs to the FPGAs in order for us to test the veracity of the Finite State Machines within them that decode the analog signals. In order to do this, the FED must identify the baseline Black level and all other address levels from the emulated data. This does not involve the FEC at all, as no real data from the ROCs are received. Therefore, this requires independent but similar calibration procedures.

14 Calibration Procedure

In this section, we describe a way to do a full calibration of the detector “from scratch”. Detailed descriptions of each algorithm are in the previous section. The procedure must be followed *in order*, because of the various dependencies between settings.

This section is based on our experience calibrating the FPix in 2008-2009 before beam and then again at a colder temperature in 2012 (7.4 C changed to 0 C). Large portions of the procedure were also used for the BPix.

We will try to include tips for common problems.

14.1 Useful Tools, Background Knowledge

This section assumes some background knowledge:

1. **POS:** You will obviously need to know how to run the Pixel Online Software. The results of the calibration are usually summarized in a log file (which one depends on the calibration, unfortunately). Most calibrations produce a root file in the Run directory which you can view in root or with the very useful HistoViewer tool.

2. **Configuration Base:** It is important that you are familiar with the structure of the configuration base and the ways to make new configurations. You should know how to change the aliases using the PixelConfigDBCmd.exe tool.
3. **Scripts to change DACs:** changeSome.pl can be used to shift the value of a particular DAC on a subset of ROCs by a specified value. The script needs a list of all DAC files and a list of ROCs to change. Run it by doing “perl changeSome.pl”. There is a similar script, called changeAll.pl, to change all ROCs. Both scripts are in TriDAS/pixel/PixelAnalysisTools/test/additionalTools.
4. **Analyze SCurve/PixelAlive/Gain:** Know how to analyse these calibrations as described in Sec. 13.11.
5. **Loop scripts:** In TriDAS/pixel/PixelAnalysisTools/test/additionalTools, there are root macros that look at the output of the analysis mentioned in the previous point. These are run by doing e.g. “root -l SCurveLoop_FPix.C”.

14.2 Starting Point

The calibrations we run only affect a subset of the detector’s many configurable settings. Many settings are fixed to initial recommendations found in specifications or early studies and will never have to change. Other settings have stayed fixed so far but may have to change as the detector suffers from radiation damage (as of January 2012, we haven’t changed any of these). To obtain good values for the settings that are not calibrated in this procedure, one should start with an old configuration, perhaps one in the P5 configuration database. There are also recommendations in the ROC specifications.

14.3 AOH and FED Channel Mapping

If the fibers have been moved, the AOH and FED Channel Mapping calibration can be run to confirm that the mapping is correct.

Here is the `calib.dat` file we have used at P5:

```
Mode: AOHAndFEDChannelMappingTest
Parameters:
ScanMin      0
ScanMax      50
ScanStepSize  5
printFEDRawData    no
printScan      no
Rows:
Cols:
VcalHigh
Repeat: 10
```

```
ToCalibrate:  
all
```

14.4 Achieving Successful Programming

Run the Delay25 calibration.

Here is the `delay25.dat` (effectively the `calib.dat`) file we have used at P5:

```
Mode:  
Delay25  
Portcards:  
All  
AllModules:  
0  
OrigSDa:  
64  
OrigRDa:  
64  
Range:  
64  
GridSize:  
1  
Tests:  
10  
Commands:  
0
```

Tip: If your efficient band is very thin, it might be because the bias settings of the DOH (DOH_Ch0Bias_CLK and DOH_Ch1Bias_Data), which reads back the result of the programming test, are too low. The plot can also look bad if the fibers are dirty.

14.5 Basic Signal Properties

Before you can optimize the performance of the ROC, you have to adjust the very basic properties of the signal so that the FED can interpret the data. There is a little bit of a bootstrapping problem here because even these calibrations need to start with a somewhat recognizable signal. Sometimes, you will have to do a preliminary calibration of setting “A” so that you can calibrate “B” before you do the final calibration of “A”. You may even have to manually tune a setting at the beginning, depending on your starting point.

14.5.1 FEDBaseline

We run this calibration all the time. It has to be run any time there is a significant change in the temperature of the AOH. At P5 in 2010-2011, we determined that the automatic

baseline correction could cope with changes of 1-2 degrees C. We monitored the size of the automatic baseline correction with the PixelMonitor tool and reran the FEDBaseline calibration whenever the size of the correction surpassed 100 units.

The calibration is also useful as a debugging tool. The calibration is fast and can often provide the feedback you need to assess the status of the detector.

We target a black level of 450 ADC units.

Here is the `calib.dat` file we used at P5:

```
Mode: FEDBaselineWithPixels
```

```
SingleROC
```

```
Rows:
```

```
0 | 1 | 2 | 3 | 4 | 5 |
6 | 7 | 8 | 9 | 10 |
11 | 12 | 13 | 14 | 15 |
16 | 17 | 18 | 19 | 20 |
21 | 22 | 23 | 24 | 25 |
26 | 27 | 28 | 29 | 30 |
31 | 32 | 33 | 34 | 35 |
36 | 37 | 38 | 39 | 40 |
41 | 42 | 43 | 44 | 45 |
46 | 47 | 48 | 49 | 50 |
51 | 52 | 53 | 54 | 55 |
56 | 57 | 58 | 59 | 60 |
61 | 62 | 63 | 64 | 65 |
66 | 67 | 68 | 69 | 70 |
71 | 72 | 73 | 74 | 75 |
76 | 77 | 78 | 79
```

```
Cols:
```

```
0 | 1 | 2 | 3 | 4 |
5 | 6 | 7 | 8 |
9 | 10 | 11 | 12 |
13 | 14 | 15 | 16 |
17 | 18 | 19 | 20 |
21 | 22 | 23 | 24 |
25 | 26 | 27 | 28 |
29 | 30 | 31 | 32 |
33 | 34 | 35 | 36 |
37 | 38 | 39 | 40 |
41 | 42 | 43 | 44 |
45 | 46 | 47 | 48 |
49 | 50 | 51
```

```
VcalHigh:
```

```
100 100 5
```

```
Repeat:  
10  
ToCalibrate:  
all
```

Tip: Multiple channels share some settings, and sometimes it's impossible for the algorithm to converge on a setting it considers satisfactory for all of the channels at once. When this happens, you will see certain channels' baselines jumping up and down from iteration to iteration while everything else looks like it has converged. If it doesn't converge within the maximum number of iterations allowed by the calibration but ends on something close to your target (say less than 75 ADC units away, maybe more), you can just keep the settings it produces because the automatic baseline correction will make up the difference. If you see the baselines jumping up and down and end on the less favorable of the two, you can run the calibration a second time. If none of this works, you'll have to do an AOHBias calibration and then retry the FEDBaseline calibration.

14.5.2 AOH Settings

We perform the AOHBias calibration and then the AOHBias gain calibration. Both of these calibrations require that the TBM UB can be located in the signal. This means that you must have decent FEDBaseline and ClockPhase settings (this is an example of the bootstrapping problem mentioned above). Both the bias and the gain have the ability to change the temperature of the AOHBias. Therefore, you should run a FEDBaseline calibration after each of them (after waiting a few minutes so that the temperature is stable). To summarize, to calibrate the AOHBias, you should do: AOHBias, FEDBaseline, AOHBiasGain, FEDBaseline.

Here is the AOHBias `calib.dat` file we used at P5:

```
Mode: AOHBias  
Parameters:  
printAOHBiasAdjustments yes  
ScanMin      0  
ScanMax      50  
ScanStepSize  1  
TargetBMin    375  
TargetBMax    525  
MaxFEDReceiverInputOffset 10  
printFEDRawData      no  
Rows:  
Cols:  
VcalHigh  
Repeat: 10  
ToCalibrate:  
all
```

Here is the AOHGain `calib.dat` file used at P5:

```
Mode: AOHGain
Parameters:
SetAnalogInputBias 220
SetAnalogOutputBias 110
SetAnalogOutputGain 220
MaxTBMUB          100
printFEDRawData no
printScan         yes
Rows:
Cols:
VcalHigh
Repeat: 10
ToCalibrate:
all
```

Note that the vast majority of the AOHGain settings at P5 are 0, with a few at 1 and hardly any above that. This may change with time.

14.5.3 Ultrablacks

It's now time to set the UB level of the TBM and ROC. First run the TBMUB12and3 calibration (this is more robust than the single-variable version), and then the ROCUB calibration. These calibrations rarely fail. As with the AOH calibrations, it's important to have a decent Clock Phase calibration before running these.

As you will see in the `calib.dat` files below, we target a UB of 150 ADC for both the TBM and the ROC. In the FPix, there are a few channels that are noisy because of a dirty fiber. Since the noise is introduced at the fiber, we can separate the rather wide address levels by applying a gain at the front end. We do this by repeating the TBMUB12and3 and ROCUB calibrations for these channels with a UB target of 50 ADC units.

Here is the TBMUB12and3 `calib.dat` file that we use at P5:

```
Mode: TBMUB
Parameters:
DACToScan      All
ScanNumSteps   52
MinAnalogInputBias 0
MaxAnalogInputBias 255
MinAnalogOutputBias 0
MaxAnalogOutputBias 128
MinAnalogOutputGain 0
MaxAnalogOutputGain 255
TargetUB       150
```

```

DualTBMaxScanStepDiff 6
DualTBMaxSettingDiff 150
printFEDRawData no
printScan no
Rows:
Cols:
VcalHigh
Repeat: 10
ToCalibrate:
all

```

Here is the ROCUB calib.dat file that we use at P5:

```

Mode: ROCUBEEqualization
Parameters: ScanMode default
TargetUB 150
Rows:
Cols:
VcalHigh
Scan: VIBias_DAC 15 255 5
Set: Vcal 100
Repeat: 17
ToCalibrate:
all

```

14.5.4 Clock Phase

To calibrate the detector from its bootstraps, a Clock Phase calibration has probably already been run, but now you are finally ready to take the final calibration.

Here is the calib.dat file we use at P5:

```

Mode: ClockPhaseCalibration
SingleROC
Parameters:
oldMode Yes
Rows:
0 | 1 | 2 | 3 | 4 | 5 |
6 | 7 | 8 | 9 | 10 |
11 | 12 | 13 | 14 | 15 |
16 | 17 | 18 | 19 | 20 |
21 | 22 | 23 | 24 | 25 |
26 | 27 | 28 | 29 | 30 |
31 | 32 | 33 | 34 | 35 |
36 | 37 | 38 | 39 | 40 |

```

```

41 | 42 | 43 | 44 | 45 |
46 | 47 | 48 | 49 | 50 |
51 | 52 | 53 | 54 | 55 |
56 | 57 | 58 | 59 | 60 |
61 | 62 | 63 | 64 | 65 |
66 | 67 | 68 | 69 | 70 |
71 | 72 | 73 | 74 | 75 |
76 | 77 | 78 | 79

```

Cols:

```

0 | 1 | 2 | 3 | 4 |
5 | 6 | 7 | 8 |
9 | 10 | 11 | 12 |
13 | 14 | 15 | 16 |
17 | 18 | 19 | 20 |
21 | 22 | 23 | 24 |
25 | 26 | 27 | 28 |
29 | 30 | 31 | 32 |
33 | 34 | 35 | 36 |
37 | 38 | 39 | 40 |
41 | 42 | 43 | 44 |
45 | 46 | 47 | 48 |
49 | 50 | 51

```

VcalHigh:

100 100 5

Repeat:

10

ToCalibrate:

all

14.6 Getting Hits

The rest of the calibrations require the use of the charge injection feature of the ROC. For the injected charge to be read out as a hit, it has to cross threshold (which involves amount of charge and the threshold) and be validated by the trigger (which involves the timing of the injection). You can get hits by doing the VcThrCalDelFIFO3 calibration, which chooses compatible values for VcThr and CalDel. This calibration will not choose the optimal settings; rather it produces initial settings that will allow you to collect the data you need to find the final ones.

Here is the `calib.dat` file we use at P5:

Mode: ThresholdCalDelay

Rows:

10 | 20

```
Cols:  
10 | 20  
VcalHigh  
Scan: VcThr 0 255 8  
Scan: CalDel 0 255 8  
Set: Vcal 40  
Repeat: 10  
ToCalibrate:  
all
```

14.7 Address Levels

Now that you have hits, you can take an address level calibration. This will allow the FEDs to decode the hit related data. We run on a subset of pixels to save time. The subset is chosen to exercise the various address level transitions.

Always run a FEDBaseline calibration afterwards to reset some levels related to the UB-B discrimination.

Here is the `calib.dat` file we use at P5:

```
Mode: FEDAddressLevelWithPixels  
Rows:  
76 |  
12 |  
51 |  
72 |  
41 |  
12  
Cols:  
10 |  
21 |  
31 |  
10 |  
20 |  
31  
VcalHigh:  
100 100 5  
Repeat:  
30  
ToCalibrate:  
all
```

14.8 Initial Trimming

The Vana calibration described in the next subsection assumes that your ROCs are decently trimmed (RMS). The trimming that you should do in this step is the same as the trimming you should do following the Vana calibration, so please refer to that subsection.

14.9 Vana Calibration

The Vana calibration procedure described here was developed for the FPix as a way to calibrate Vana without a measurement of the analog current drawn by an individual ROC. The only measurement of analog current available now that the detector is fully assembled is the total current drawn from a single power supply, which services more than one-hundred ROCs.

We calibrate Vana to give a certain timewalk, which we quantify for each ROC with a variable we call “the difference in thresholds”, or DT. The DT of a ROC is defined as the average in-time threshold minus the average absolute threshold. It quantifies the speed of the amplifiers and therefore the current drawn by the ROC.

The target value of DT is chosen such that the average analog current per ROC in each power group is near specification ($\sim 25\text{mA}$). The correct target for DT depends on radiation damage (and temperature?), so we cannot give you a number to target. Instead, you should tune the target based on the analog current reading you have. For reference, we targeted 12 VCal for the FPix in 2009 (at 7.4C and before beam). In 2012, we targeted 8 VCal for the FPix and 13-14 VCal for the BPix, depending on the power group (and settled on value ~ 1 VCal higher due to rounding).

The Vana calibration is an iterative procedure. Several of the calibrations described in the previous section and a standalone analysis tool are used in each iteration.

Before beginning the iterative part of the procedure, you should do a VcThrCalibration targeting an absolute threshold of 70 VCal. We choose 70 VCal because it is usually high enough for you to avoid the problems associated with the thresholds being too low (failures in VcThrCalibration and CalDelCalibration). If you have these problems at 70 VCal, you can do the procedure at 80 VCal instead.

Next, you do iterations of the following list

1. **CalDelCalibration to 20%.** With the Vana and VcThr DACs changing in each iteration, it is important that you recalibrate CalDel. This ensures that the DT is a comparable quantity from iteration to iteration.
2. **SCurve99By3 (two WBCs).** Take an SCurve run on ~ 100 pixels/ROC (81 is our usual number) scanning over two WBCs (the in-time one and the one after it, which is smaller by 1 DAC unit). Run the SCurve analysis on the in-time WBC data and rename the TrimOutput file as `TrimDefault.dat`. Also run the SCurve analysis on the data from both WBCs and rename the TrimOutput file as `TrimDefault_abs.dat`.

3. **Run PixelInTimeVana to produce new DAC settings.** First make sure that `PixelAnalysisTools/test/src/common/PixelInTimeVana.cc` has `assumeGlobal=false` and the correct DT targets (with `assumeGlobal=false`, the only thing that matters is the difference between `globalThr` and the `targetInTimeThr` variable(s)). The relationship between a change in Vana and a change in DT is assumed to be the same for each ROC and is hardcoded in this file. We found it best to do a few iterations with $\Delta DT / \Delta Vana = 1$ and then change to it to 2 to avoid overshooting. The code must be compiled by doing ‘make’ in `PixelAnalysisTools/test`. Run it by doing `./bin/linux/x86/PixelInTimeVana.exe <key> <run number> > mylog.dat`. Keeping the log file is worthwhile because it allows you to watch the size of the corrections in each iteration to see if you are converging. The program will produce new DAC files for the ROCs you ran on. Copy the Default dacs to a new version, copy these files in, and make this new version the new Default.

We have done as many as ~ 20 iterations to get final values. The final distribution of `deltaVana` was ~ 0.5 with no DTs more than a few VCal away from the target. This can be very time-consuming.

Some scripts we have used to plot the $\Delta Vana$ distribution, currents, and other things are in `TriDAS/pixel/PixelAnalysisTools/test/additionalTools/VanaTools`.

Note: The first incarnation of the iterative part of the Vana calibration was slightly different. We found that it worked well for the vast majority of ROCs, but prevented some ROCs from converging (at least in 2012). The first incarnation was used in 2009 and the beginning of the 2012 recalibration. We include a description of it here just in case it is useful. We have never done a Vana calibration using only the second incarnation described above, simply because it was invented during the most recent calibration of the detector.

In this first incarnation, we assumed that the absolute threshold of each ROC was equal to the target of the `VcThrCalibration` rather than using a measurement of the mean absolute threshold provided by an `SCurve` calibration. The ROCs that had problems converging were ROCs for which this assumption was not very good (this can happen since the `VcThrCalibration` uses only a few pixels and is therefore more sensitive to the trimming).

Here is the first incarnation of the iterative part:

1. `VcThrCalibration` to 70 Vcal to reset the absolute thresholds (since not measured!)
2. `CalDelCalibration` to 20%
3. `SCurve99By3` to measure in time threshold only
4. Run `PixelInTimeVana` with `assumeGlobal=true` to produce new DAC settings

14.10 Final Trimming

The BPix group determined optimal trimbit settings for their detector before beam and since then have used measurements of `Vtrim` vs temperature to maintain a well-trimmed detector.

There are no such measurements for the FPix ROCs, so we redo the full FPix trimming when recalibrating the detector. We have developed a relatively fast trimming procedure.

The main purpose of the trimming is to make all of the thresholds within one ROC the same. This is important for the Vana calibration because we assume that the average threshold is a meaningful quantity. It is also important in the ROC-based threshold minimization. Our trimming algorithm makes all of the thresholds across the detector the same, and as a consequence, achieves the desired goal of making the thresholds within each ROC the same.

Begin by taking a VcThrCalibration with a target of 80 Vcal. We choose 80 Vcal because it is high enough to avoid problems associated with thresholds being too low. 70 Vcal would also work, we think. Take an SCurve99By3 run to measure the in-time threshold and set the target threshold in PixelTrim.cc and PixelTrimBits.cc to the mean in-time threshold (this way you don't waste time shifting your distribution and just spend time narrowing it).

Take iterations of the procedure described in Sec. 13.21.5 until the RMS of the thresholds seen in a TrimDefaultShort run is less than ~ 1.5 VCal. Each iteration takes about one and a half hours on the FPix. Then complete the full trimbit determination as described in Sec. 13.21.6. The long TrimDefault run takes about 8 hours on the FPix.

Note that the SCurves in the section are only taken with the in-time WBC, so we are trimming the in-time threshold. BPix trimmed the absolute threshold. Either way is fine.

14.11 Threshold Minimization

The goal of the threshold minimization is to lower the thresholds on each ROC to just above the point it starts failing. Failure is identified using SCurve and PixelAlive runs. To change the threshold in this procedure we only use VcThr, which shifts all of the thresholds on the ROC coherently. Raising VcThr lowers the thresholds.

If you are calibrating a test stand and are not concerned with optimizing the thresholds, you can just run a VcThrCalibration to 60 or 70 Vcal and skip the rest of this section. This will save you a lot of time.

Begin the procedure by doing a VcThrCalibration to 50 VCal. The vast majority of ROCs should work at this threshold (less than 10 out of all 15k should fail). This gives you the starting point for the VcThr setting of each ROC. The few ROCs that fail at this setting will have to have their threshold manually raised.

Next, take SCurve runs with VcThr shifted for all ROCs by -20, 0, +2, +4, +6, +8, +10, +12, +14, and +16. This can be done very easily using the `SetRelative` option of the `calib.dat` file, as shown in the example file below. When the thresholds are lowered, the digital current will increase significantly, causing the temperature to increase. When we ran at VcThr+16, the temperature increased so much that the automatic baseline correction went outside of the acceptable range (150). We therefore prepared a special DAC file for this run where ROCs that had already failed were kept at their nominal threshold (aka VcThr+0) while the rest were changed to VcThr+16.

`calib.dat` file for SCurve with shifted VcThr:

Mode: SCurve

```

Rows:
 0 | 9 |
18 | 27 |
36 | 45 |
54 | 63 |
72

Cols:
 3 16 29 | 42 4 17 |
30 43 5

VcalLow
Scan: Vcal 10 120 1
Scan: WBC 155 156 1
SetRelative: VcThr -20

Repeat:
20

ToCalibrate:
all

```

You should now have SCurve data for VcThr-20, 0, +2, +4, +6, +8, +10, +12, +14, and +16. Proceed by analyzing the data using the usual PixelAnalysisTools to get the absolute thresholds (so analyze both WBCs). Analyze the FPix and the BPix separately to produce one root file for each. The root files are used in the following steps.

Now, go to the TriDAS/pixel/PixelAnalysisTools/test/additionalTools/ directory. You will make use of several ROOT macros and perl scripts in this directory.

First, look at the VcThr-20 run. This run was taken at a high threshold where there should really be no problems due to thresholds being too low, so you can use it to determine how many pixels on each ROC out of the 81 you injected with charge are working. To do this, run SCurveLoop_BPix.C and SCurveLoop_FPix.C on the VcThr-20 ROOT files with some special options set: PRINT=0, print=1, plot=0, hold=0, useBase=0. Also, make sure to change **fFile** to the correct root file and **out_name** to something that indicates which run it is (e.g. “minus20”). Run by doing `root -l -b -q SCurveLoop_FPix.C`.

You should now have two “..._print.dat” files (one for FPix and one for BPix) that contain the number of pixels to expect from each ROC. Change **printedFile** in SCurveLoop_BPix(FPix).C to open the BPix(FPix) file and set the following options: PRINT=0, print=0, plot=0, hold=0, useBase=1. Now when the macros are run, the number of pixels expected from a ROC will be taken from the VcThr-20 run.

Now run SCurveLoop_BPix(FPix).C on the remaining SCurve runs (VcThr+0,2,4,6,8,10,12,14,16). Remember to change **out_name** and **fFile** for each run. Combine the FPix and BPix “..._fail.dat” files for each run by doing e.g. `cat plus6_BPIX_fail.dat plus6_FPIX_fail.dat > plus6_fail.dat`.

Create a copy of the Default DACs. Use the `changeSome.pl` script to shift the VcThr of every ROC included in the configuration (so, not noInit or noAnalogSignal ROCs) by +14 (two units less than the farthest you went, as a safety margin). `iterate.pl` will tell you how much the thresholds of each ROC must be raised away from these settings in order

for them to be above their failure point. First change the input files in iterate.pl to the “...fail.dat” files you just created. Then run it by doing “perl iterate.pl”. This will create lists of ROCs grouped by the required VcThr shift. The lists are put into text files with names that indicate the required VcThr shift. For that example, **sub6.dat** contains the list of ROCs that must have 6 subtracted from their VcThr value. Use changeSome.pl with all of these “subX.dat” files to apply all of the necessary shifts.

You now have a preliminary guess of the minimized thresholds. You should now rerun an SCurve with the new VcThr values (and no shift). Some ROCs will fail even though according to the runs you took earlier they should pass (We guess this is because the ROCs are not fully independent from one another... or they are unstable near the failing point). Analyze the SCurve using SCurveLoop_BPix(FPix).C and correct any ROCs that fail (i.e. that are listed in the “...fail.dat” files). To correct them, we recommend shifting their VcThr value by -4. Repeat this until everything passes.

Next, you can try running on a different subset of pixels (it’s best to shift at least the double columns used) to make sure all the ROCs still pass. You’ll have to run at VcThr-20 again to make a new “useBase” file. You may find a few ROCs fail because just one or two pixels within the ROC barely fail the min pixel threshold cut. This is just because they are in the tail of the distribution and is not an indication of the threshold being too low. You can correct their VcThr by -4, but we wouldn’t recommend doing this for more and more subsets if this is the only kind of failure you see. Obviously you should correct any ROC that fails in a more serious way.

Now do a second round of tests using the PixelAlive calibration. Take PixelAlive runs using the **calib.dat** file below and look for any signs of thresholds being too low. Correct the VcThr of any ROC that shows these signs and repeat. You can compare the efficiency map to that from a run taken at VcThr-20 using the **SetRelative** option if you are not sure if an inefficiency is due to the threshold being too low or e.g. a broken double column that will appear with any VcThr.

After a tedious day or two of repeating these SCurve and PixelAlive runs with manual corrections being made after each one, you should finally reach a point where every ROC reliably passes. If you find that many ROCs are not acting reliably (passing in one run and then failing in the next even though the settings are the same), consider shifting every ROC’s VcThr by an additional -2 to save yourself some time. We did this to the BPix in 2012.

Mode: PixelAlive

Parameters:

| | |
|----------|---------|
| ROCReset | yes |
| ScanMode | default |
| Rows: | |
| 0 | |
| 16 | |
| 32 | |
| 48 | |
| 64 | |

```

1 |
17 |
33 |
49 |
65 |
...
79
Cols:
0 |
13 |
26 |
39 |
1 |
14 |
27 |
40 |
...
51
VcalLow
Set: Vcal 150
Repeat:
10
ToCalibrate:
+
all

```

Note: The “...” are just to save paper – you should use every pixel. You should use the Default masks with this one (to have whatever is in Physics in and have noisy pixels masked)! VCal 150 was carefully chosen to be above the threshold of all pixels without being so large that it adds unnecessary cross-talk in the charge injection.

14.12 Pulse Height Optimizations

14.12.1 Vsf

Here is the `calib.dat` file we use at P5:

```

Mode:  VsfAndVHldDel
Rows:
3 12
Cols:
14
VcalLow

```

```
Scan: Vsf      0 255 3    mix
Scan: VHldDel 0 255 255
Set: Vcal 150
Repeat: 10
ToCalibrate:
all
```

Tip: If a ROC fails, it may be because its threshold is too low. Decrease VcThr by 5-10 units and try again.

14.12.2 VHldDel

Here is the `calib.dat` file we use at P5:

```
Mode:  VsfAndVHldDel
Rows:
12
Cols:
14 25
VcalLow
Scan: VHldDel 0 255 5
Set: Vcal 250
Repeat: 40
ToCalibrate:
all
```

Note that the value the calibration chooses depends somewhat on the VCal used. The value of 250 (low scale) has been used for a long time and seems to work well enough.

14.12.3 PHRange

Here is the `calib.dat` file we use at P5:

```
Mode:  PHRange
Parameters:
minPH 70
maxPH 235
Rows:
12
Cols:
14 25
VcalHigh
Scan: VIBias_PH 0 200 10
Scan: VOffsetOp 0 150 5
```

```
Scan: Vcal 17 167 150
```

```
Repeat: 4
```

```
ToCalibrate:
```

```
all
```

Tip: If a ROC fails this calibration, it may be because the range of the scan is not wide enough. We try to keep the scan narrow and with only a few points because this calibration uses a lot of memory for some reason.

14.12.4 Gain Calibration

The last step is to take a gain calibration. We usually take a short one on a subset of pixels to verify that everything is okay before doing the final one on all pixels that takes more than six hours.

The main purpose of taking the short version is to verify that the PH range calibration did its job. To do this, open up the resulting root file and look at the histogram called ADCOfAllPixelsSummary. This is a histogram of all of the ADC values obtained from all hits on all ROCs in the calibration. Make the y-axis log-scale and confirm that there are no entries at zero or the maximum. If there are, it means that one or more ROCs is outside the range of the FED's ADC. You can run GainLoop.C in TriDAS/pixel/PixelAnalysisTools/test/additionalTools to find out which ROCs are responsible. You will have to redo the PHRange calibration of those ROCs.

The full gain calibration is analyzed by the offline group and put into the data reconstruction.

Here is the `calib.dat` file we use at P5 for the short gain calibration:

```
Mode: GainCalibration
```

```
Rows:
```

```
0 | 9 |
```

```
18 | 27 |
```

```
36 | 45 |
```

```
54 | 63 |
```

```
72
```

```
Cols:
```

```
0 13 26 | 39 1 14 |
```

```
27 40 2
```

```
VcalHigh
```

```
ScanValues: Vcal 2 4 6 8 10 12 14 15 16 17 18 21 24 28 35 42 49 56 63 70 77 84 91 98 105  
1
```

```
Scan: WBC 155 156 1
```

```
Repeat:
```

```
5
```

```
ToCalibrate:
```

```
all
```

Here is the `calib.dat` file we use at P5 for the full gain calibration:

```
Mode: GainCalibration
Rows:
 0 |  1 |  2 |  3 |  4 |  5 |  6 |  7 |  8 |  9 |
10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 |
20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 |
30 | 31 | 32 | 33 | 34 | 35 | 36 | 37 | 38 | 39 |
40 | 41 | 42 | 43 | 44 | 45 | 46 | 47 | 48 | 49 |
50 | 51 | 52 | 53 | 54 | 55 | 56 | 57 | 58 | 59 |
60 | 61 | 62 | 63 | 64 | 65 | 66 | 67 | 68 | 69 |
70 | 71 | 72 | 73 | 74 | 75 | 76 | 77 | 78 | 79 |
Cols:
 0 13 26 | 39 1 14 |
27 40 2 | 15 28 41 |
 3 16 29 | 42 4 17 |
30 43 5 | 18 31 44 |
 6 19 32 | 45 7 20 |
33 46 8 | 21 34 47 |
 9 22 35 | 48 10 23 |
36 49 11 | 24 37 50 |
12 25 38 | 51
VcalHigh
ScanValues: Vcal 6 8 10 12 14 15 16 17 18 21 24 28 35 42 49 56 63 70 77 84 91 98 105 112
1
Repeat:
5
ToCalibrate:
all
```

15 XDAQ-DCS Interaction

Close interaction between the DAQ and DCS (Detector Control System) software systems are needed for the powering of the FPix, and also so that special settings can be applied to the ROCs in case the high voltage is off.

From the DAQ side, the startup procedure is embedded within the “Initializing” and “Configuring” FSM transitions prescribed by Run Control. In working through these states, the DAQ system must use information from the DCS system about the voltages being applied across various detector and front-end elements.

The XDAQ Supervisor applications mainly involved with this are:

- PixelFECSupervisor: Controls a crate of Pixel FEC boards

- PixelTKFECSupervisor: Controls a crate of Tracker FEC boards
- PixelDCSFSMInterface: Reports the voltages applied by the DCS system (simplified to LV_OFF/ LV_ON_REDUCED/ LV_ON for low voltage and HV_OFF / HV_ON for high voltage).

15.1 PSX Server

The gateway between DCS and our XDAQ applications is the PSX server. This server is a XDAQ application that runs on the machine `cmspsx` at P5. This machine runs independent PSX server applications for all of the different subdetectors. Both the machine and the software running on it are administered by the central DAQ experts.

We interact with two separate PSX server application – the “Pixel PSX server” running on port 9923, and the “Tracker PSX server” running on port 9922. The Pixel PSX server is used for making connections to DCS to learn the status of the power supplies (the topic of this section). The Tracker PSX server is used to send DCU data to DCS. Note that the Tracker PSX server is also used by the strip tracker, and for performance reasons it has proven critical to use the independent Pixel PSX server for making the connections to the power supply status.

15.2 PixelDCSFSMInterface

The PixelDCSFSMInterface provides an interface between the PSX server and our other pixel XDAQ applications. Changes in the states of each power supply module are sent in real time from the PSX server to the PixelDCSFSMInterface application. The PixelDCSFSMInterface summarizes this information into one state per half cylinder/shell (thus there are 8 summarized states, 4 in the FPix and 4 in the BPix). The logic for making the summary is as follows:

- If all A4603 power supplies in the summary group are ON (HV and LV is ON), then the summarized state reflects this
 - A4603 power supplies in state HVMIXED (one HV channel is on and the other is either off or ramping), are treated as if they were ON⁹.
- When forming the summary, if an entire power supply has been marked as noInit or noAnalogSignal in the detconfig, then that power supply is ignored when forming the summary¹⁰

⁹This configuration was deployed while planning for the first beam data, because of a desire to have the high-voltage treated as ON, even if the FPix inner radius HV was OFF.

¹⁰Initially the code would only ignore the power supply if the status was marked as noInit (because a noAnalogSignal power group still required LV power). However, for the partial configurations prepared for the first beam running, we wanted to be able to ignore a power supply even if it was in state noAnalogSignal.

- If at least one power supply in the summary group has the HV in the off state, then the summarized state treats the whole group as though it is in the state with HV off and LV on.
- If at least one power supply in the summary group has the LV off, then the summarized state treats the whole group as off.

When the summarized state of one of the half cylinders/shells changes, the PixelDCSFSMInterface sends a SOAP message to the relevant PixelTKFECSupervisor/PixelFECSupervisor. The nodes to be readout from DCS are defined in `PixelDCSInterface/xml/interface.xml`.

15.3 Use of DCS information by the supervisors

SOAP messages are sent from the PixelDCSFSMInterface to the PixelTKFECSupervisors and PixelFECSupervisors in real time whenever the summarized state of the power supplies changes. The Supervisors receive these messages asynchronously (in other words, they handle them immediately in a separate thread). When a supervisor receives an update about the power status, the only action taken is to update member data that stores the summarized state. No other action is taken at that moment.

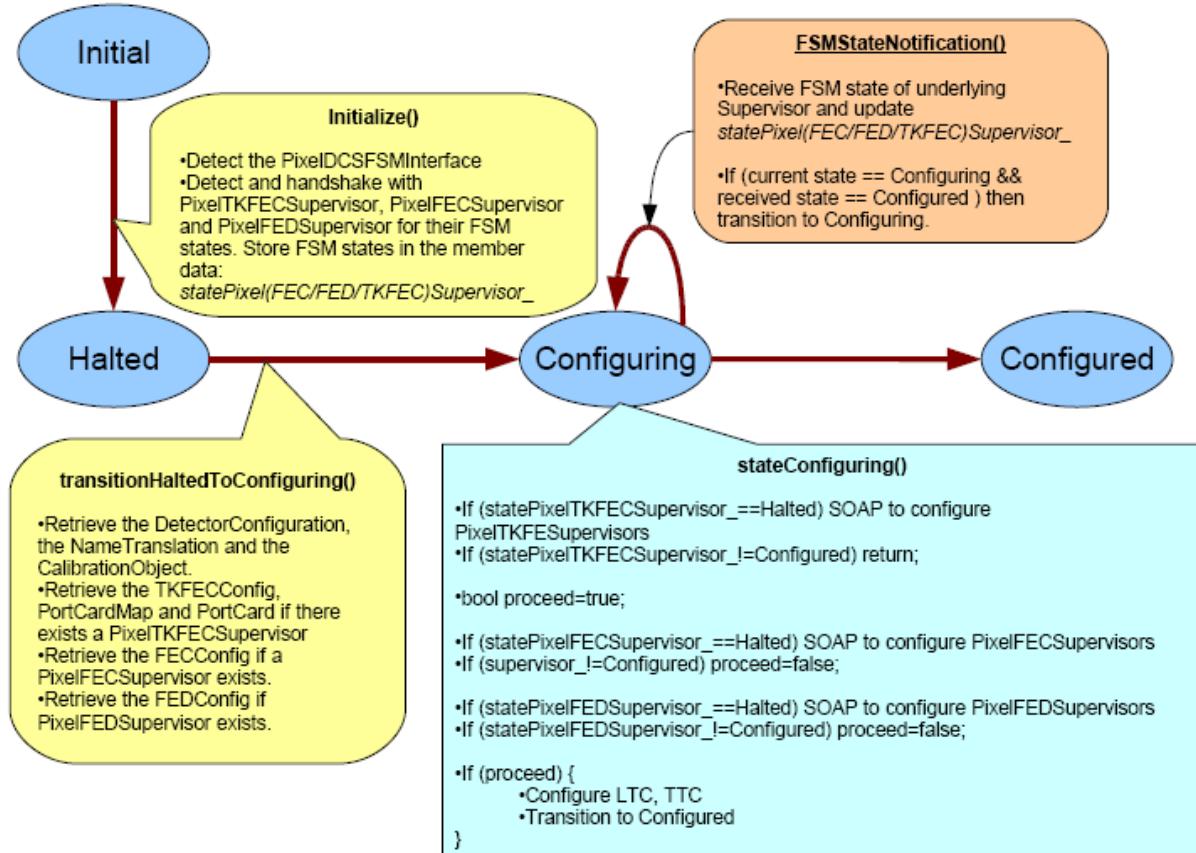
The PixelFECSupervisors receive and track the states of the low and high voltage of the A4603 power supplies. The low voltage has three possible states: LV_OFF, LV_ON_REDUCED, and LV_ON. The high voltage has two possible states: HV_OFF and HV_ON.

The PixelTKFECSupervisors receive and track the states of the low voltage provided by the A4602 power supplies. There are two possible states: LV_OFF and LV_ON.

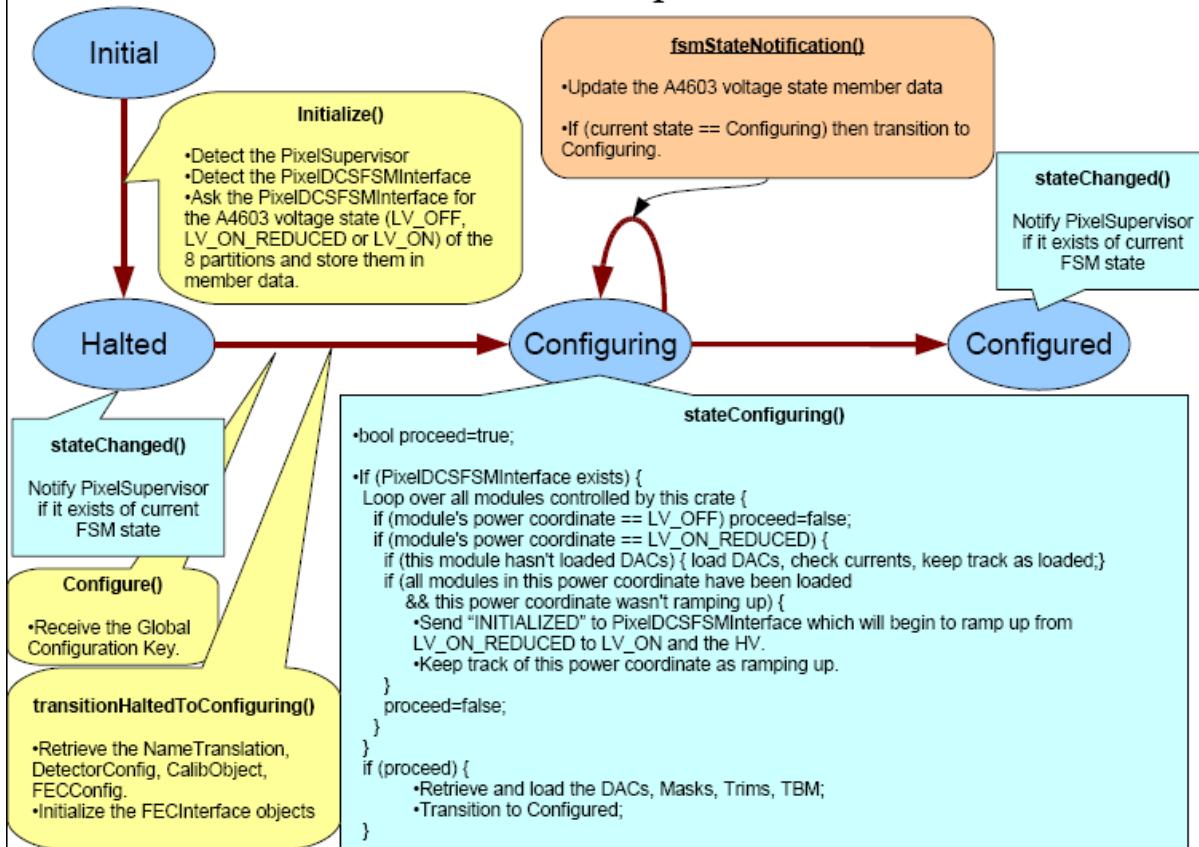
15.4 Configuration

The progress through the DAQ's FSM states and transitions for initialization and configuration is envisioned as follows:

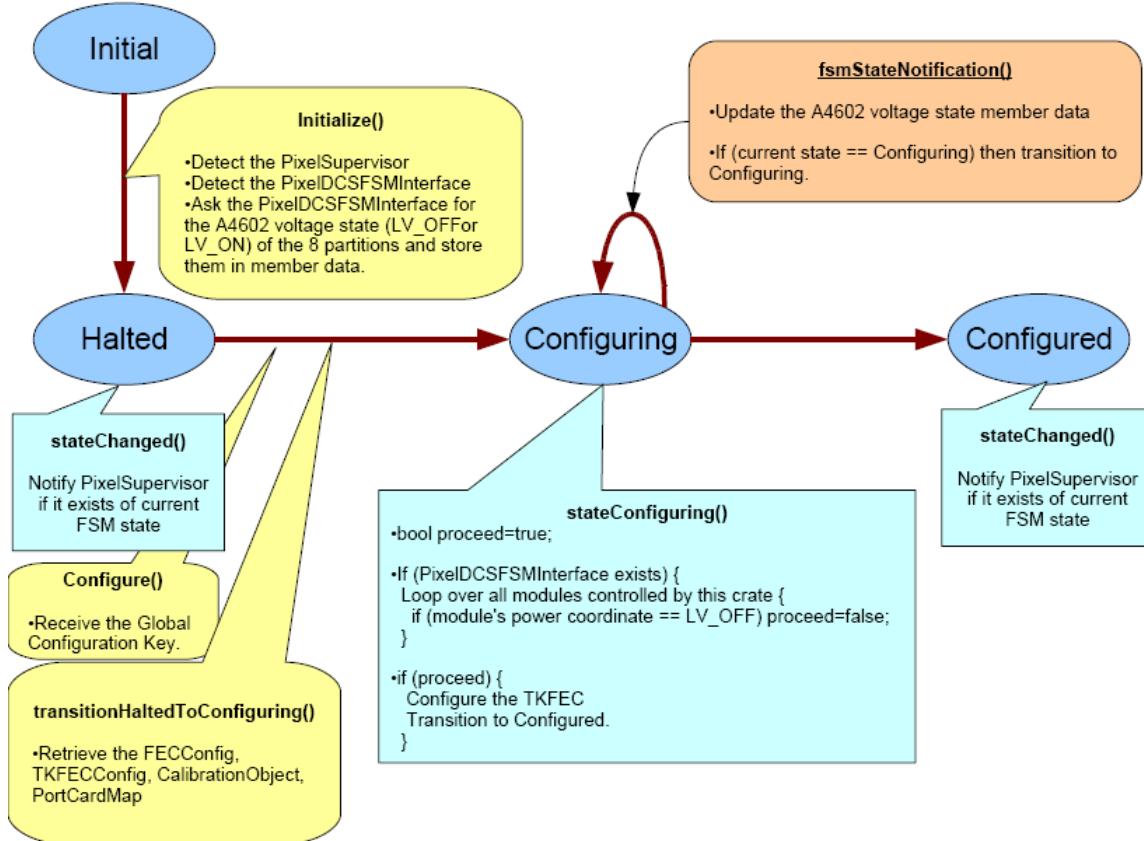
PixelSupervisor



PixelFECSupervisor



PixelTKFECSupervisor



1. PixelFECSupervisor(s) and PixelTKFECSupervisor are initially in their "Initial" state and await the "Initialize" command.
2. On receiving the "Initialize" command, PixelFECSupervisor(s) and PixelTKFECSupervisor independently query PixelDCSFSMInterface.
3. The PixelDCSFSMInterface responds with information regarding the voltage states.

Note that in practice, the correct power state is never reported in the "Initialize" step, because at P5 the "Initialize" transition is always triggered immediately after starting the software, at which point the PixelDCSFSMInterface has not yet had time to make connections with the PSX server and be notified of the state of the power supplies. Once the PSX server connections are made, the correct power state is sent via the standard asynchronous mechanism.

4. The PixelFECSupervisor(s) and PixelTKFECSupervisor update their tri- and bi-state member data (for LV and HV as appropriate) according to the response from PixelDCSFSMInterface. This concludes the "Initializing" transition and both the PixelFECSupervisor(s) and PixelTKFECSupervisors land in their "Halted" states.

5. On receiving the "Configure" command, PixelFECSupervisor(s) and PixelTKFECSupervisor check the status of their member data. If the appropriate voltages are set then they go ahead with step 6. If the LV_OFF state is found, configuration does not proceed.
6. If the PixelFECSupervisor finds that the LV is in state LV_ON_REDUCED, the ROCs' DAC settings are loaded and programmed to the ROCs. (It has been proposed to query the low-voltage currents before and after this loading of DAC settings, in order to ensure that the currents change when the DAC settings are loaded. However this check is not implemented at present.) If the LV is already in state LV_ON (as is the case for the BPix, or on subsequent configurations of the FPix), this step is skipped.
7. PixelFECSupervisor sends a SOAP message to PixelDCSFSMInterface to notify it that the FPix DAC settings have been loaded. The PixelDCSFSMInterface forwards this message to DCS via the PSX server.
8. PixelFECSupervisors wait for an LV_ON SOAP message from PixelDCSFSMInterface.
9. PixelFECSupervisors load their ROCs' DAC, mask and trim settings. When programming the DACs, we check the state of the data member containing the HV status. The PixelTKFECSupervisors load the slow I2C settings for CCU boards, Port Cards etc.
10. The PixelFECSupervisor(s) and PixelTKFECSupervisor(s) land in their "Configured" state.

During the configure step in PixelFECSupervisor, we follow the "HV-dependent" DAC programming procedure if and only if we are configuring for a calibration. If we are configuring for Physics running, then the ROCs are always disabled (treated as if the HV was off).

15.5 HV-dependent DAC programming

At Configure (for calibrations only), and at the Start and Resume transitions, we check the data members that store the HV status, and we check the data members that track whether the ROCs were most recently enabled or disabled. If we find that the HV is on and the ROCs were previously disabled, we enable the ROCs. If we find that the HV is off and the ROCs were previously enabled, then we disable the ROCs. If we find that the HV is on but the ROCs are already enabled, then we do nothing; similarly if the HV is off but the ROCs are already disabled.

To disable the ROCs, we overwrite the VcThr DAC setting with 0 (high threshold), then disable the ROCs using the ROC control register. These procedures are implemented in the PixelDACSettings class. This class caches the DAC settings from the DB, and if we override the VcThr setting with zero we do not change what is in the cache, so no new DB accesses are needed when we want to enable the ROCs. The PixelDACSettings::generateConfiguration()

method has an argument for the calling function in PixelFECSupervisor to indicate the status of the HV. To enable the ROCs, the nominal VcThr and CCR settings are programmed.

15.6 Start and Resume transitions

The status of the HV is checked and the DACs are programmed if needed, as explained above.

15.7 Pause, Stop, and Halt transitions

The ROCs are disabled as if the HV was off.

15.8 Expected behavior

As currently implemented, the following is the expected sequence of events when turning on the FPix:

- User turns FPix to ON in DCS
- FPix is powered to LV_ON_REDUCED (called DIGITAL_ON_REDUCED in DCS)
- User begins the Configure transition in PixelSupervisor
- FPix DACs are programmed
- POS notifies DCS that the DACs are programmed; POS then waits in the “Configuring” state
- DCS raises the digital voltage to the nominal value
- POS is notified that we have reached the LV_ON state
- DCS begins to ramp the high voltage.
- Configuration in POS proceeds with the reprogramming of the DACs and the programming of the trims and masks
- Configuration in POS finishes. At some point (probably after the end of configuration), the high voltage reaches its nominal value and DCS notifies POS that we have reached the HV_ON state (called simply ON in DCS).
- Because the HV was not on at the time of configuration, the DACs are set to the special settings used for HV_OFF
- User presses Start to begin a calibration in POS

- The FEC Supervisor notices that the state of the HV has changed since the configuration. Therefore the VcThr and chip control register DACs are set to their nominal values. Because of this settings change, the user will see a large change in the amount of digital current drawn by the FPix.
- The calibration proceeds as normal.

15.9 Notes on hardcodeding of information

Because the XDAQ software (POS) and DCS use different naming schemes and different granularity, several translations must be made in order for POS to communicate with DCS. Some of these translations are encoded in the `interface.xml` file, while others are hardcoded in the logic of the PixelDCSFSMInterface.

Specifically, the xml file defines which ROGs¹¹ exist, and which PVSS datapoint name belongs to which ROG¹². The xml file also defines the POS applications that need to be notified in case the state of a ROG changes.

The hardcoded logic in PixelDCSFSMInterface defines which ROG correponds to which pieces of the detector (which ROCs are supplied by which power supply).

In principle, this information could be instead contained in a database table with roughly 5 fields:

| Field in database | Example value |
|-------------------------|--|
| POS name | FPix_BmI_D2_BLD1 |
| DCS name | CMS_TRACKER:PixelEndCap:PixelEndCap_BmI:PixelEndCap_BmI_D2:PixelEndCap_BmI_D2_ROG1 |
| datapoint name | cms_trk_dcs_10:tkPg_CAEN\CMS_TRACKER_SY1527_5\branchController05\easyCrate3 ... |
| SOAP connection (TKFEC) | PixelTKFECSupervisor instance 2 |
| SOAP connection (FEC) | PixelFECSupervisor instance 2 |

With a table such as this, we could eliminate both the xml file and the hardcodeding.

16 Procedures to follow when hardware components are replaced

This section documents the procedures that are to be followed when off detector hardware components are replaced that affects the online software and calibrations.

16.1 FED board

As the FED decodes the analog levels it will be required to redo several calibrations to check the optical connections after replacing a FED. The FED firmware version has to be validated. The following calibrations should be run

¹¹readout groups

¹²These datapoints are used to indicate when POS has configured the ROCs on a ROG.

- FED baseline
- AOH gain
- AOH bias
- FED baseline (it is possible that the previous three steps have to be repeated.)
- FEC clock and phase
- TBM UB
- ROC UB
- Address level

16.2 Pixel FEC including mFECs

No calibrations should be affected by replacing a FEC board or a mFEC — in principle. However, we have seen that the size of the window for good communication changes a little between different mFECs. This seems mostly to affect the return data and this check is now disabled for both FPix and BPix. It is probably prudent to rerun a Delay25 calibration to check that after reconnecting the fibers we have the expected performance.

Need to verify that the right firmware is installed. We assume that we manage to connect the fibers in the right slot.

16.3 TKFEC

No calibrations should be affected by replacing a TKFEC board.

Need to verify that the right firmware is installed. We assume that we manage to connect the fibers in the right slot.

17 Low-level commands

In addition to the complex sequences of commands issued by the high-level configuration procedure, a number of low-level commands are available to the user through the GUIs of the various Supervisors. These can be useful for expert-level problem solving. These are described in this section.

17.1 Resets

- ResetROC: sent via the TTC. Available as a button in the PixelSupervisor GUI. Implemented and verified to work.

- ResetTBM: sent via the TTC. Available as a button in the PixelSupervisor GUI. Implemented and verified to work.
- ResetCCU: causes the TKFECSupervisors to issue the command `resetPlxFec()`. Available as a button in the PixelSupervisor GUI. Implemented. Sends `resetPlxFec` via the FecSoftware.
- PIA resets: a menu of resets available from the PixelTKFECSupervisor GUI, implemented by passing a value to the FecSoftware command `testPIAResetfunctions`
 - roc: value = 0x1
 - aoh: value = 0x2
 - doh: value = 0x4
 - res1: value = 0x8
 - res2: value = 0x10
 - fpixroc: value = 0x2A (bits 1, 3, and 5); this resets the TBM and ROCs
 - fpixdevice: value = 0x15 (bits 0, 2, and 4); this resets the portcard devices (Delay25, DOH, AOH, DCU)

The PIA resets are hard resets using hardware lines. In the FPix they go from the CCU parallel output lines to reset lines on the portcard. One is fanned out to all the portcard devices. The other is routed to the TBM and ROCs through the adapter board¹³.

¹³<https://hypernews.cern.ch/HyperNews/CMS/get/pixelOnlineSW/1155/1/1.html>

A FED phase and delay calibration

This section will describe some aspects of determining the so called ‘phase’ and ‘delay’ of the FED. The sample plots shown in this section are taken from runs on the FPIX pilot run detector. This was done using a version 4 FED with the Aug. 22, 2007 firmware version.

In order to get the transparent data to ‘look’ OK we had to send an LRES before each trigger. We think that the final conclusion was that the FED state machine gets confused when the input data is non standard. In the scans over phases and delays there are certain settings that are invalid, i.e., that you try to read out the ADC before the data is available. At the end of this section we will discuss two issues that we noticed in looking at the transparent data.

For the above reason we also turn off the automatic baseline correction to avoid it getting confused if you send non valid data or if you don’t have correct address levels.

In these tests we take 10 triggers for each delay and phase setting. We will show plots, as in Fig. 24, where for each of the two possible phases the profile histogram shows the adc value as a function of clock+delay/16. The third plots just show the two plots on top of each other. We see that there are certain values of the delays that produce garbage as you try to read out the ADC before the data is available. To identify the invalid combinations of phase and delay we calculate the rms for the different phase and delay settings in the first 20 clock cycles, i.e. before the TBM header arrives. Based on the rms distribution it is seen that for phase=0 we got invalid data with delays of 1, 2, and 3, while for phase=1 the delays of 10, 11, and 12 generate invalid data. From now on we will reject these combinations. (In the plots we set the adc value to 0.)

After rejecting the invalid phase and delay combinations we now have a signal that looks as what is shown in Fig. 25. The remaining signal still has artifacts due to when the adc latches on to the signal such that it looks like the signal is not read out in order. This is corrected by ‘reordering’ the data within a clock. We shifted the phase=0 events backward by 3 delay steps and the phase 1 events forward by 3 delay steps. This makes the signal continuous as shown in Fig. 26.

The last step is to align the data so that the phase=0 and phase=1 data overlaps. We do this by shifting the data for the events with phase=1 by 10 delay steps. The final result is shown in Fig. 27. For the bins in the plot where there is data from both phase=0 and phase=1 we see that there is excellent agreement. We also note that the signal undershoots in a transition from black to ultrablack and that it overshoots on the reverse transition.

To determine a best delay we look at the combined data - where we average the bins that has both phase=0 and phase=1 contributions - and calculate a difference for bin i by taking the difference between bin $i-1$ and $i+1$. We calculate this difference for the 25 MHz clock cycles from clock 27 to 100. We add the magnitude of these differences for all bins that corresponds to the same delay. We pick the smallest such sum of differences for the different delay as the optimal choice. If more than one phase is allowed we pick the choice that is farthest from the invalid delay for that phase.

A new algorithm has been implemented that determines the phase to be right at the

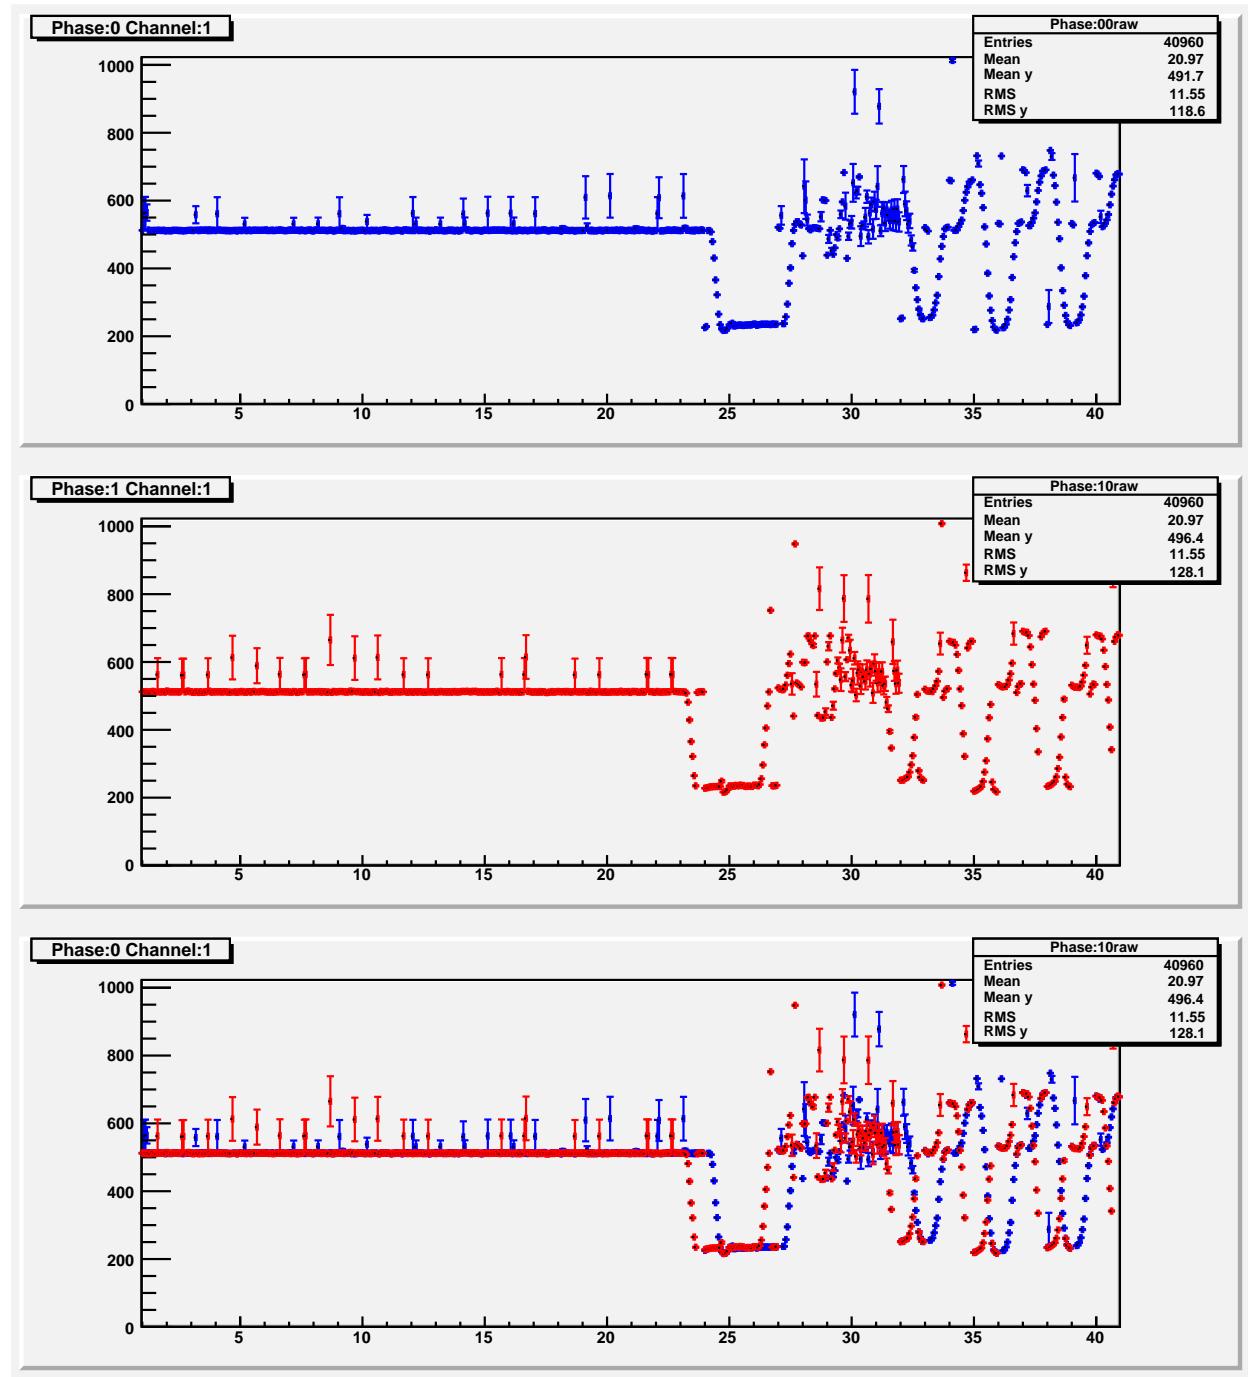


Figure 24: This figure shows on top the adc value as a function of clock+delay/16 for phase=0 and the same plot for phase=1 in the middle. The two plots are overlaid at the bottom.

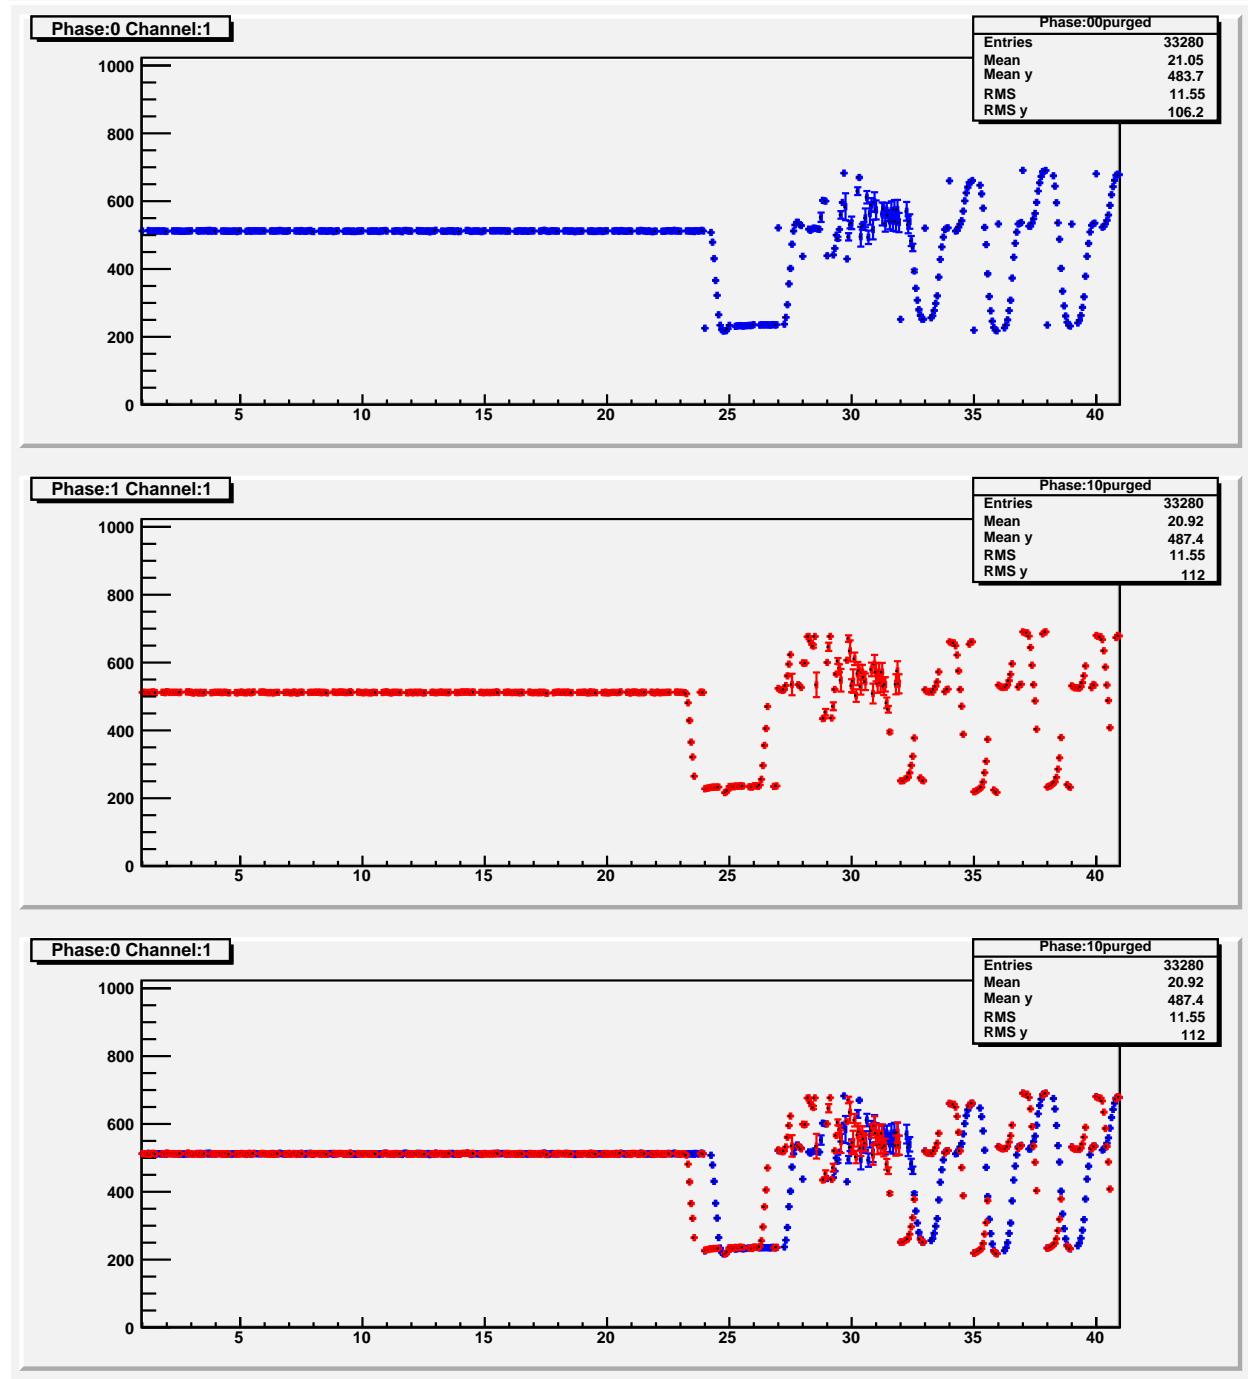


Figure 25: This figure shows on top the adc value as a function of clock+delay/16 for phase=0 and the same plot for phase=1 in the middle. The two plots are overlaid at the bottom. Here we have set the adc value to 0 for the combinations of clock and phase that are not valid.

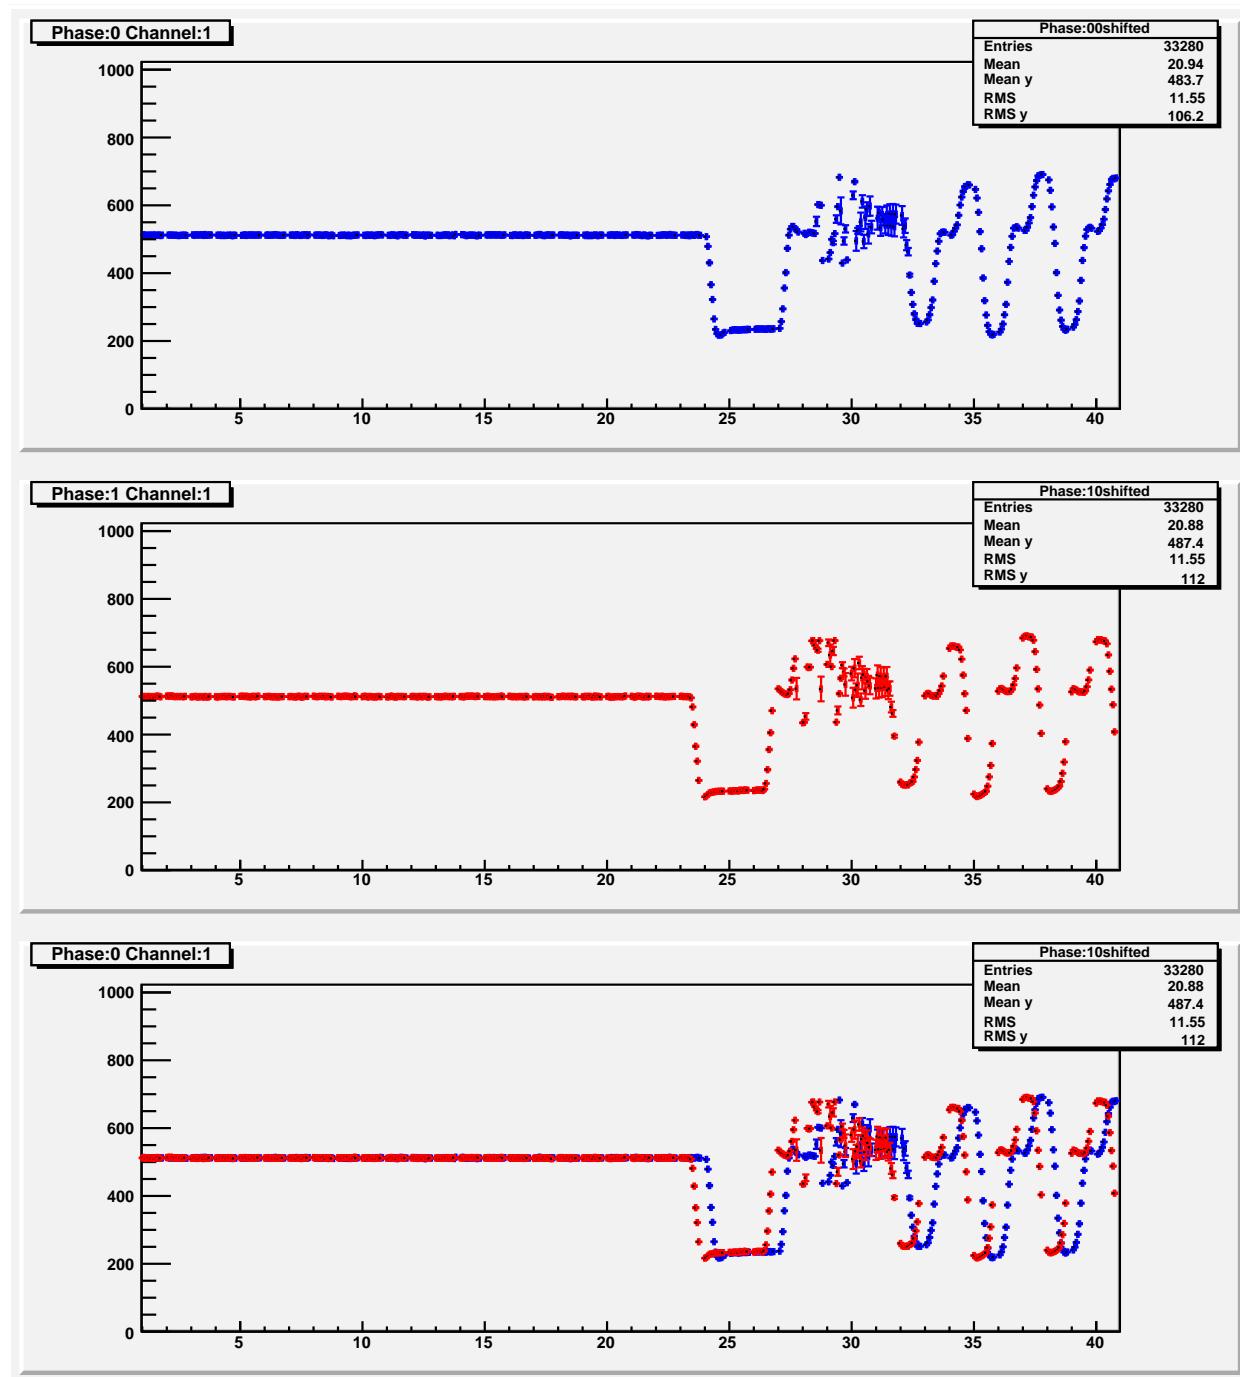


Figure 26: This figure shows on top the adc value as a function of clock+delay/16 for phase=0 and the same plot for phase=1 in the middle. The two plots are overlaid at the bottom. Here the data within a clock has been ordered to give a continuous data signal.

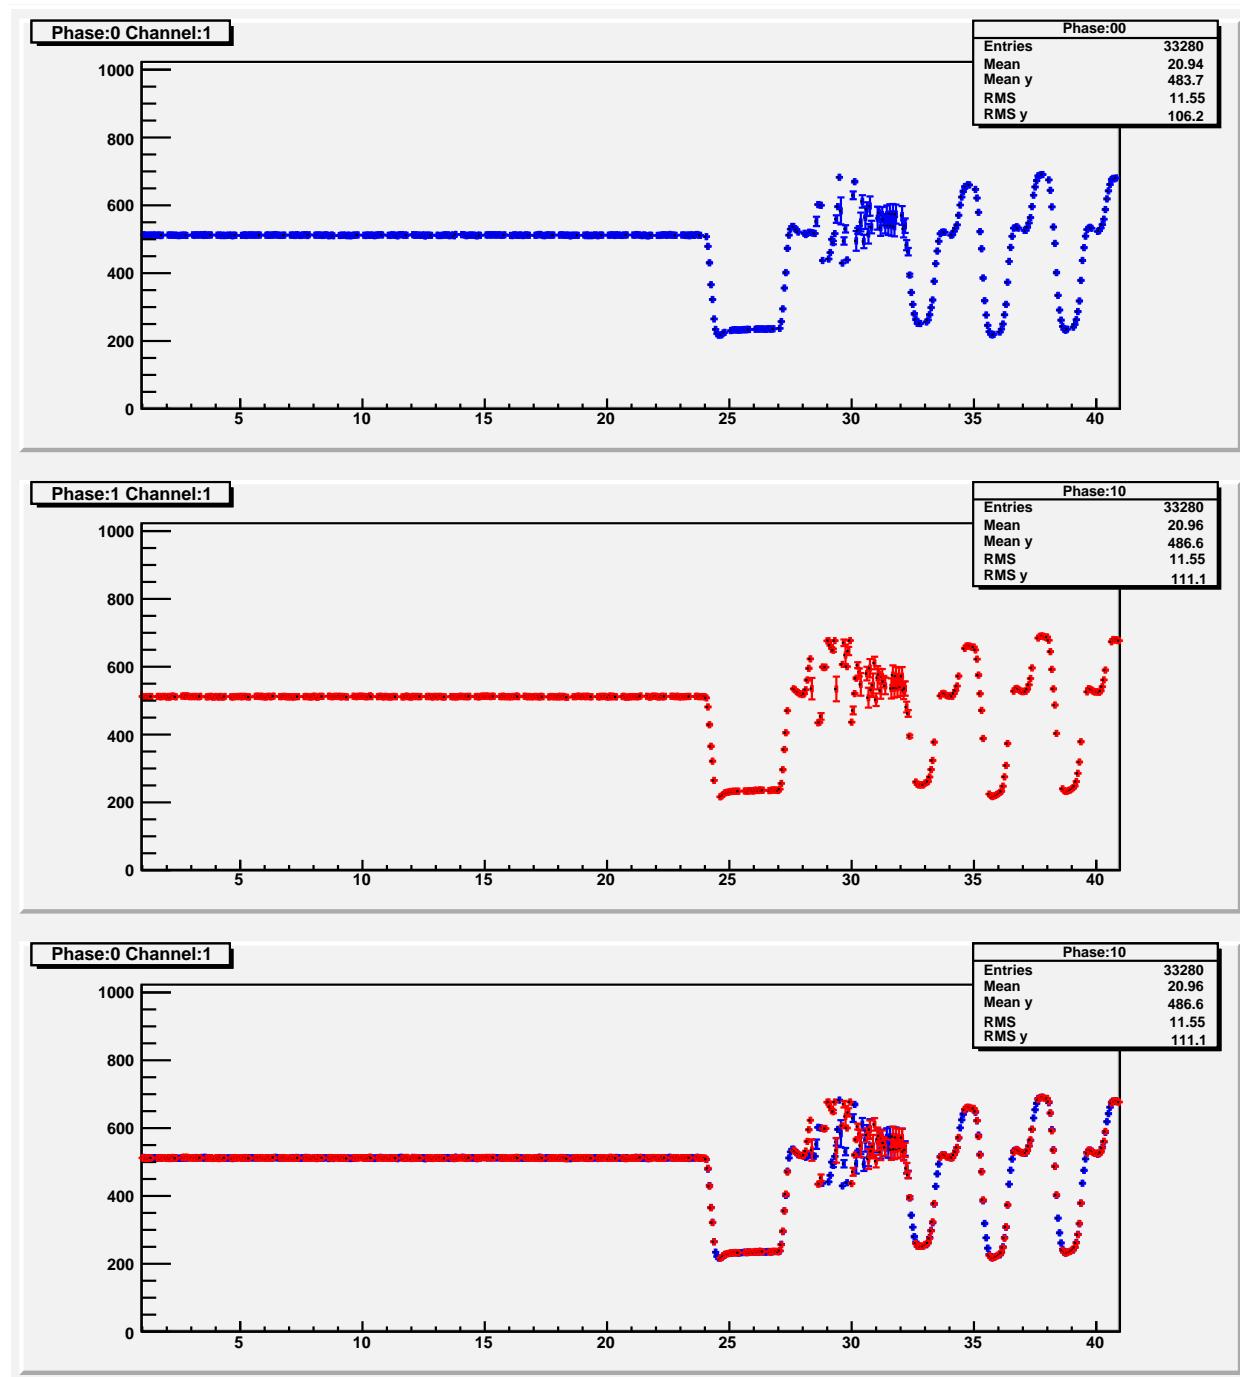


Figure 27: This figure shows on top the adc value as a function of clock+delay/16 for phase=0 and the same plot for phase=1 in the middle. The two plots are overlaid at the bottom. Here the phase=1 data has been shifted so that it overlaps with the phase=0 data.

edge of the transition from black to ultrablack in the TBM trailer. You can still run the old algorithm by setting the parameter 'oldMode' to 'Yes'.

Looking at the 8 channels that were connected to the pilot run detector We find that the results are very consistent. The phase is either 0 or 15 in the (aligned delay) for the raw delay setting this corresponds to 2 or 3 for phase=1.

The code was written as a root macro. We will 'transplant' this code into XDAQ so we can run it automatically to produce the phase and delay settings needed.

However, looking at this we observed an odd feature. This is illustrated in Fig. 28. In Fig. 29 is a close up of the region with the noise. Here you can see that there is 'wiggle' in the data that is large compared to the rms, as indicated by the error. This wiggle seems to persist for many clock cycles after the TBM trailer. However, it was not present in the black data before the data train. It seems to first start in the TBM trailer when the UB signal is generated. With the pilot run detector we are using 8 channels (1, 2, 3, 4, 7, 8, 9, and 10). We see the same feature on almost all channels, though they are not as strong on channel 1, 7, and 9.

B Calibration Organization

The calibrations all share a similar format.

- The PixelSupervisor controls a loop over 'events' and coordinates the activity of the FED, FEC, TKFEC, and TTC supervisors.
- Each calibration needs to do some initialization before processing the event data.
- During event processing data is acquired and processed.
- After all event data is taken it is analyzed and the results are saved.

This section will describe a proposal for how to formalize this process. A first step in this direction was taken when the calibration classes were broken out from the PixelSupervisor and the PixelFEDSupervisor. This first step was almost trivial as it just copied the code out to separate classes. There are a few more changes we should make that will allow us to run these jobs in workloops. Among other things, this will allow us to look at the progress of the calibrations from the web browser.

B.1 PixelSupervisor calibration code

We change the `PixelCalibrationBase` class to have the interface below.

```
class PixelCalibrationBase : public PixelSupervisorConfiguration,  
    public SOAPCommander {  
public:
```

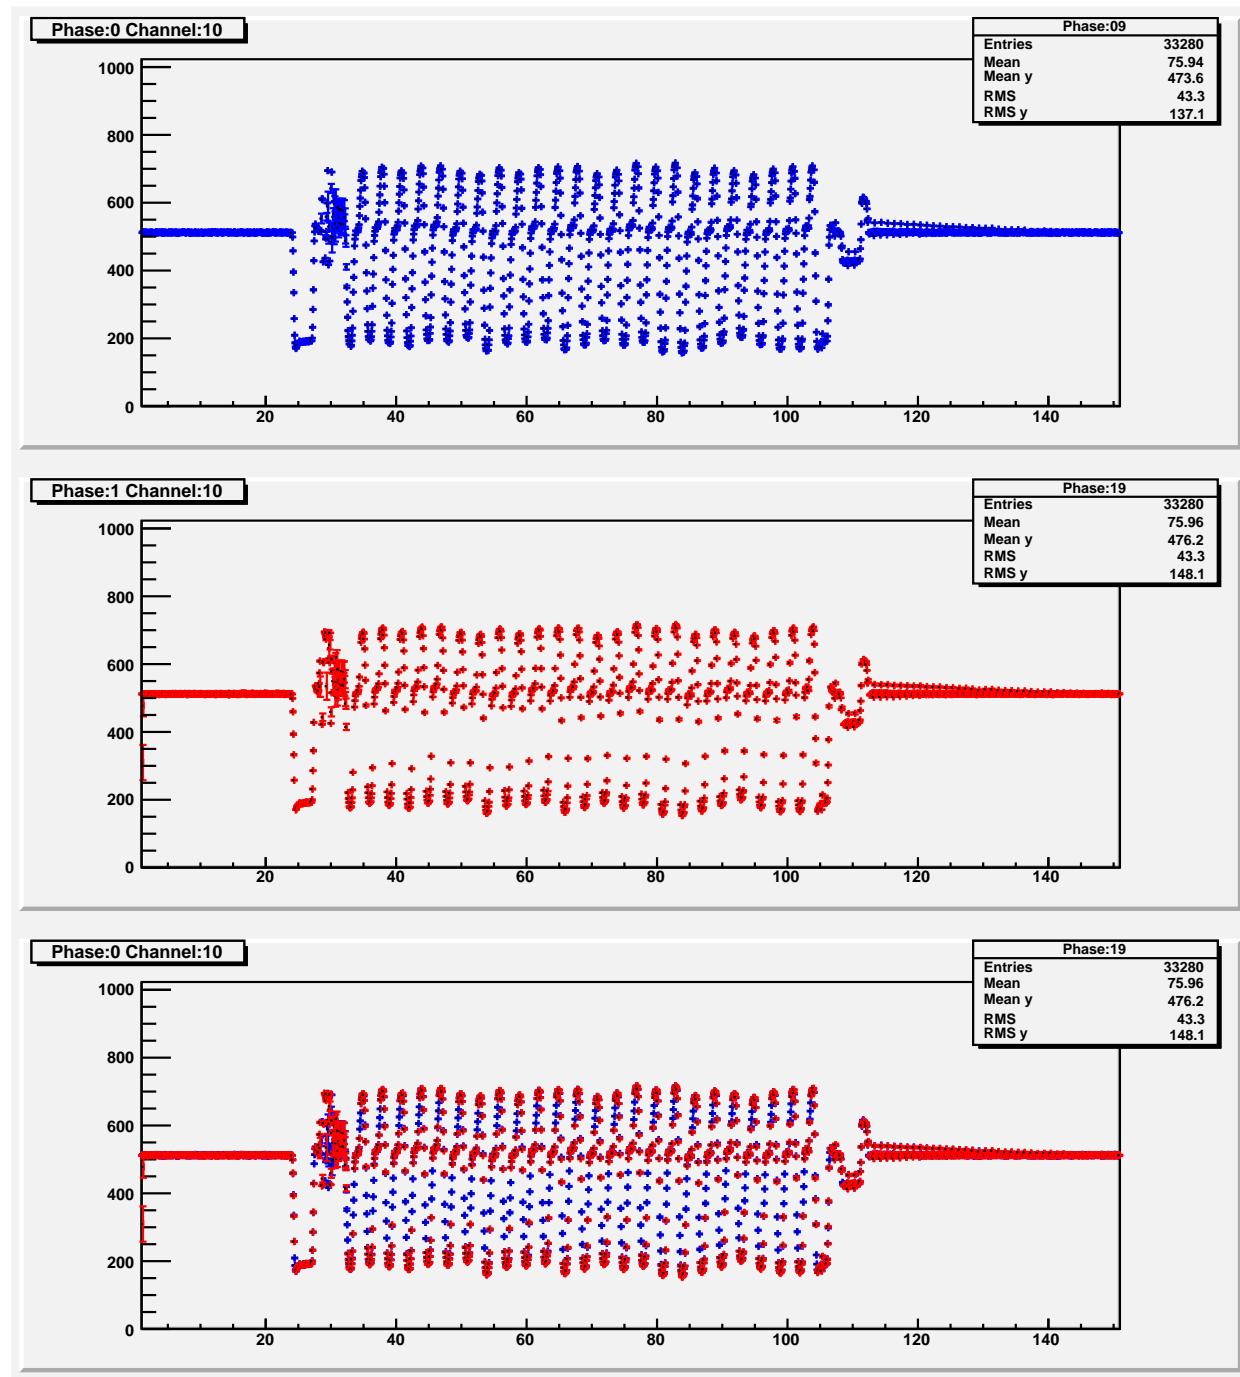


Figure 28: Here you see that for clocks in and after the TBM trailer that there is a feature in that some delays are oscillating. This 'excitation' seems to decay of after $O(20)$ clocks.

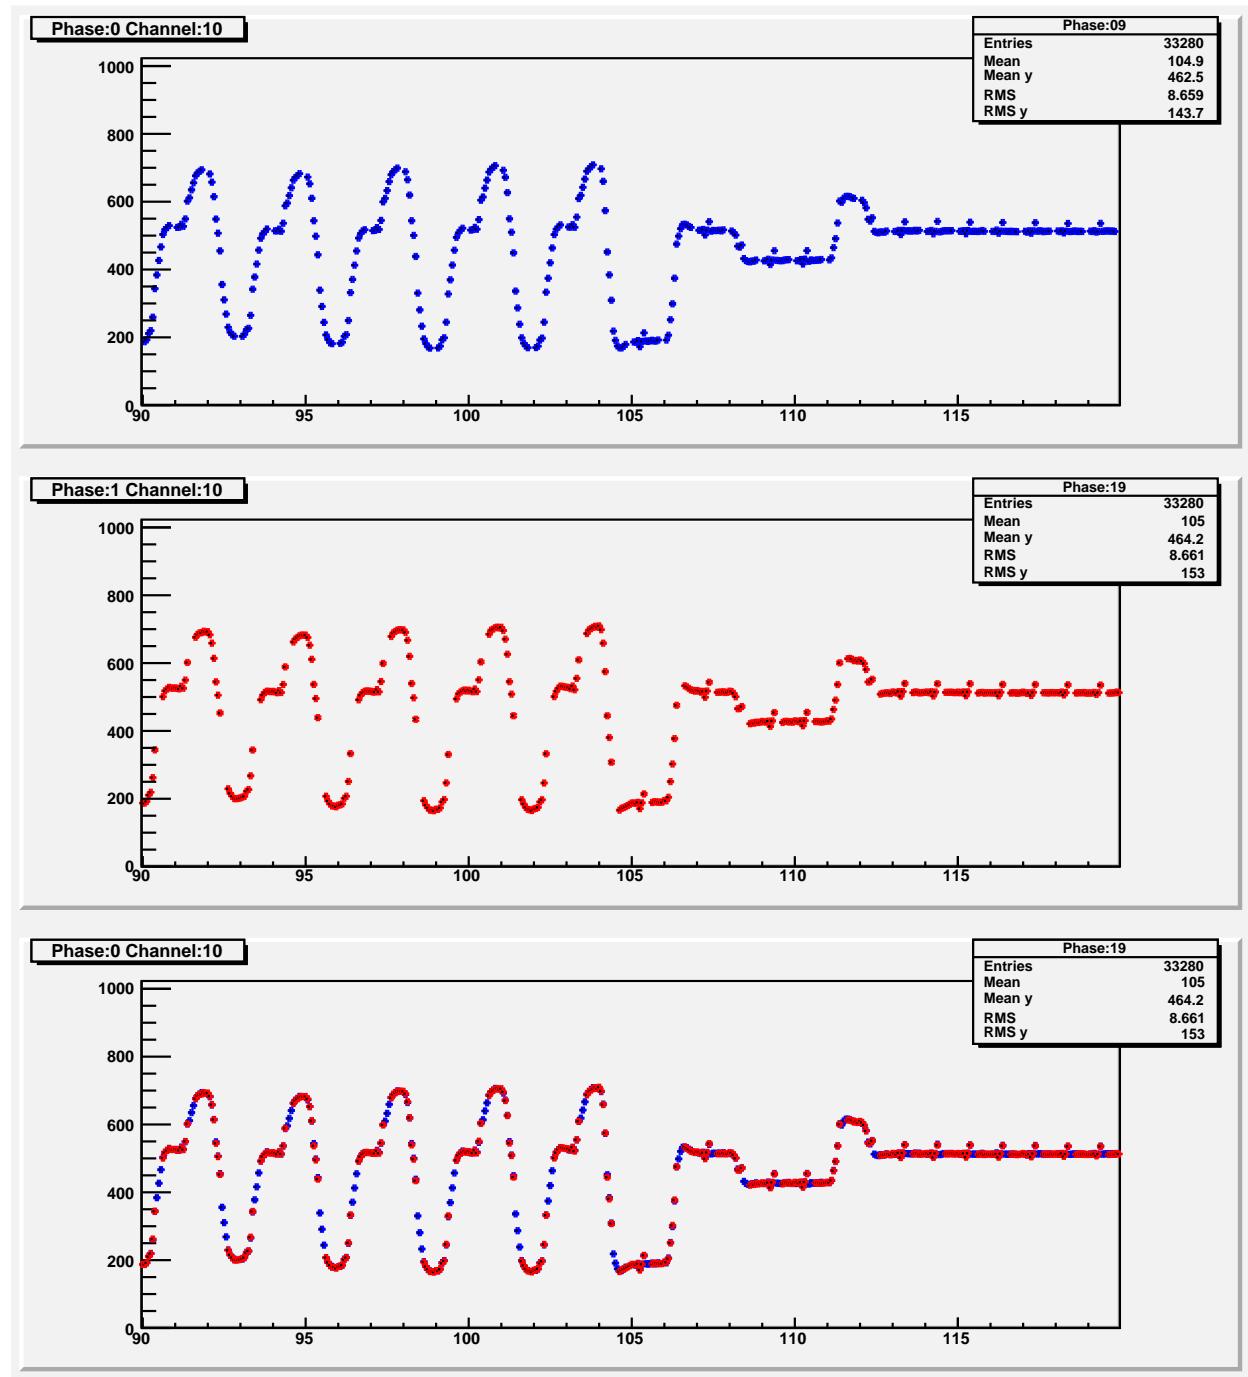


Figure 29: Zooming in you can see that data in the two phases both agree on the shape of this feature.

```

PixelCalibrationBase( const PixelSupervisorConfiguration &,
                      const SOAPCommander& soapCommander ) ;

virtual ~PixelCalibrationBase(){} ;

virtual void beginCalibration() ;

virtual bool calibrationEvent() ;

virtual void endCalibration() ;

protected:

private:

// PixelCalibrationBase Constructor
PixelCalibrationBase(const SOAPCommander& soapCommander) ;

};


```

The method `begin()` should do any initialization that is needed before processing any triggers. The method `event()` is called repeatedly until it returns true, this indicates that the calibration has come to an end. Then the `end()` method is invoked to perform any analysis of the acquired data.

More details on what these methods need to do will be discussed below.

B.2 PixelFEDSupervisor calibration code

Similarly, we suggest a more general structure for the calibration code that runs in the PixelFEDSupervisor.

```

class PixelFEDCalibrationBase : public PixelFEDSupervisorConfiguration,
                                public SOAPCommander {
public:

PixelFEDCalibrationBase( const PixelFEDSupervisorConfiguration &,
                        const SOAPCommander& soapCommander ) ;

virtual ~PixelFEDCalibrationBase(){} ;

virtual xoap::MessageReference beginCalibration(xoap::MessageReference msg) ;

virtual xoap::MessageReference calibrationEvent(xoap::MessageReference msg) ;

```

```

virtual xoap::MessageReference endCalibration(xoap::MessageReference msg);

//Should really not be here.....
unsigned int TransparentDataStart (unsigned long *buffer);

protected:

//SOAPCommander* theSOAPCommander_;

private:

// PixelCalibrationBase Constructor should never be called
PixelFEDCalibrationBase(const SOAPCommander& soapCommander);

};

```

Again the three methods `beginCalibration()`, `calibrationEvent()`, and `endCalibration()` corresponds to the pre calibration data taking, processing of trigger and the post data taking analysis respectively. We suggest that these methods are bound to soap messages that invokes the implemented methods.

In particular we should move the code that sets the mode and control registers to the `beginCalibration()` method. This will allow us to move out code from the PixelFEDSupervisor that is specific to different calibrations.

B.3 Implementation of algorithms

During the configuration step the `PixelSupervisor` and `PixelFEDSupervisor` instantiates the Calibration class for the calibration objects. When we go to Run the `PixelSupervisor` invokes the `beginCalibration()` method. This method is responsible for sending the `beginCalibration()` message to the `PixelFEDSupervisors`.

After this the `PixelSupervisor` should execute the `calibrationEvent()` in a workloop until it has completed. Again, the code executed in the `PixelFEDSupervisor` is responsible for calling the appropriate supervisors. This gives the maximal flexibility in the calibration algorithm.

Last, the `endCalibration` method in the `PixelSupervisor` is invoked to complete the analysis of the collected data and report the results.

References

- [1] The pixel online wiki is available at
<https://twiki.cern.ch/twiki/bin/view/CMS/PixelOnlineSoftware>.
- [2] *PSI46 Pixel Chip*, version 2. Available at <http://cms.web.psi.ch/roc/psi46v2.pdf>.
- [3] Token Bit Manager 05a Chip Documentation, version 1.13. Available at
<http://www.physics.rutgers.edu/~bartz/cms/tbm2/0.25u/index.html>.
- [4] M. Barbero, W. Bertl, G. Dietrich, A. Dorokhov, W. Erdmann, K. Gabathuler, St. Heising, Ch. Hörmann, R. Horisberger, H. Chr. Kästli, D. Kotlinski, B. Meier, and R. Weber. Design and test of the cms pixel readout chip. *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, 517(1):349 – 359, 2004.
- [5] RCMS Documentation (Level 1 Function Managers, updated 17 Feb 2009). Link available at <http://cmsdoc.cern.ch/cms/TRIDAS/RCMS/Docs/Manuals/manualIndex.html>.
- [6] Katarina Gromova, ETH Zurich, Semester Thesis, *Adjustment of the Readout Range of the CMS Pixel Detector* (2008). Available at <http://www.phys.ethz.ch/~ursl/home/v0/projects/p06.pdf>.